Informatica Data Ingestion and Replication October 2026 Data Ingestion and Replication Release Notes

© Copyright Informatica LLC 2019, 2026

Contents

Common. . . . 2
Fixed issues. . . . 2
Known issues. . . . 4
Connector and connectivity issues. . . . 4
Application Ingestion and Replication. . . . 9
Fixed issues. . . . 9
Known issues. . . . 9
Connector and connectivity issues. . . . 9
Database Ingestion and Replication. . . . 11
Fixed issues. . . . 11
Known issues. . . . 13
Connector and connectivity issues. . . . 14
File Ingestion and Replication. . . . 34
Fixed issues. . . . 34
Known issues. . . . 34
Connector and connectivity issues. . . . 34
Streaming Ingestion and Replication. . . . 35
Fixed issues 35
Known issues. . . . 35
Connector and connectivity issues. . . . 35
Informatica Global Customer Support. . . . 36

Read the *Data Ingestion and Replication Release Notes* to learn about fixed issues and known issues in the July 2026 release of Data Ingestion and Replication, which includes Application Ingestion and Replication, Database Ingestion and Replication, File Ingestion and Replication, and Streaming Ingestion and Replication.

For information about new features and enhancements, see *Data Ingestion and Replication What's New*.

Common

Fixed issues

The following tables describe fixed issues that apply to two or more ingestion types. Not all monthly releases include fixed issues.

October 2026 fixed issues

Issue	Description
DBMI-31134	When you run an application ingestion and replication or database ingestion and replication combined load job, the job might fail with an error indicating that a duplicate backlog storage queue already exists.
DBMI-30866	When you run a combined load job using Parquet staging and audit columns are enabled, the initial load task fails intermittently during backlog processing with a data type mismatch error. This issue occurs during the transition from the initial load phase to the incremental load phase when the job does not detect schema changes between phases.
DBMI-30398	When you run application ingestion and replication tasks or database ingestion and replication tasks that write data in Avro format, the Secure Agent process might experience extremely high CPU utilization, causing tasks to run slowly or fail. This issue is more likely to occur with wide tables, large data volumes, or column names with non-ASCII characters, especially when multiple tasks run concurrently.
DBMI-29251	When you create an application ingestion and replication task or a database ingestion and replication task, a warning message appears stating that the selected task location does not exist. This occurs when the default task location points to a deleted project or folder.
DBMI-28633	When you try to stop a database or application ingestion and replication job that is already in a stopped state because of the API endpoint POST request `/dbmi/public/api/v2/job/stop`, the API incorrectly returns a 404 status code with the error message "JOB_IS_NOT_ACTIVE". Instead, the API should return a 409 status code to indicate a conflict.

July 2026 fixed issues

Issue	Description
DBMI-30090	An application ingestion and replication or database ingestion and replication job's status displayed in the monitoring interface might show a transitional state such as "Stopping" or "Running" even after the underlying run has finished, instead of reflecting the correct final status such as "Stopped" or "Completed." The discrepancy is caused by a race condition in which a concurrent metrics update overwrites the final status written by the status update path.
DBMI-29730	Calls to the DBMI REST API v2 endpoints /dbmi/public/api/v2/task/create and /dbmi/public/api/v2/task/details might return HTTP 500 errors with the message "`Agent with id ... does not exist`," even though the runtime environment is present and running.
DBMI-29278	When using the Data Ingestion and Replication Command-Line Interface (CLI) to retrieve per-table metrics for a table whose name is a prefix substring of another table name in the same job, and when rename rules cause the alphabetical sort order of source table names to differ from that of target table names, the CLI might return "Subtask name does not match either source table or target table name" even though the specified table exists and has valid metrics.

April 2026 fixed issues

Issue	Description
DBMI-28703	You do not have the option to configure task-level Java Virtual Memory (JVM) for tasks running in serverless runtime environments.
DBMI-28508	For application ingestion and replication and database ingestion and replication combined load jobs, markers intended for initial load and backlog processing might incorrectly trigger the incremental load processing with no recovery checkpoints written, which might lead to duplicate data on restart.
DBMI-28447	Data Ingestion and Replication components contain outdated jose4j release 0.9.4 inherited from a third party driver.
DBMI-28238	When an application ingestion and replication job or a database ingestion and replication job restarts after a failure, data for a table that was already sent to the target might be resent while recovering data for other tables not yet sent to the target.
DBMI-27833	For organizations that use the shared secure agent group with multiple agents from parent organization, same unload tasks within a database ingestion and replication or application ingestion and replication might be are assigned to more than one agent, resulting in duplicate unloads for the same table.
DBMI-27520	Application ingestion and replication combined load jobs and database ingestion and replication combined load jobs that use custom properties related to the load transition processing might ignore those custom properties and instead apply only the default values. The following error might be issued: `Remote task submission failed due to invalid ReservationId`
DBMI-27223	Adding the customPrimaryKeys source custom property to an existing incremental load or combined initial and incremental load task might result in a primary key change not allowed error. Following a resync, new DML changes are not replicated to the target, and a restart of the task shows the primary key change not allowed error for the resynced table.
DBMI-26372	When you view details of an application ingestion and replication job or a database ingestion and replication job from Monitor, the overall job status might be missing from the job details page.
DBMI-26176	Data Ingestion and Replication components contain vulnerable Simba JDBC driver.

Known issues

The following table describes known issues that apply to two or more ingestion and replication types. Note that the (*month*, *year*) value in each issue description represents the release in which the issue was found.

Issue	Description
DBMI-31701	When you undeploy an application ingestion and replication job or a database ingestion and replication job from the Actions menu in the upper-right corner of the task configuration page, the undeploy action completes immediately without a confirmation prompt. (October 2026)
DBMI-31669	When an application ingestion and replication job or a database ingestion and replication job fails, the **Resume with Options** command is missing from the Actions menu in the upper-right corner of the task configuration page. Workaround: The **Resume with Options** command is available in any of the following monitoring interfaces: - **My Jobs** page that's accessed from the navigation bar of the Home page - **All Jobs** page in Monitor - **All Jobs** tab on the **Data Ingestion and Replication** page in Operational Insights (October 2026)
DBMI-28100	When you click the **Refresh** button next to the **New Connection** button in the task wizard, the connection list does not update. Connections created in Administrator while the task is open do not appear in the list. Workaround: Close and reopen the task to select the new connection. (October 2026)

Connector and connectivity issues

Read the following pages to learn about fixed issues, known limitations, and third-party limitations that apply to connectors that you can use in your service. Consider the following guidelines for the release notes:

- If a connector has known or fixed issues, you will find it in the release notes even if the connector is not updated in the current release.
- If a connector does not have any notable known or fixed issues, you will not find it in the release notes.
- The (month, year) value in each issue description indicates the release in which the issue was fixed or reported.

Amazon S3 V2 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-26804	Application ingestion and replication jobs and database ingestion and replication jobs that have an Amazon S3 target might fail if the AWS region of the S3 bucket is AWS GovCloud (US-West). The following error message is issued: `Cannot create enum from region description AWS GovCloud (US-West) value` (April 2026)

Databricks Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29670	For application ingestion and replication jobs and database ingestion and replication jobs that have a Databricks target and use Parquet file staging, the infa_operation_time audit column might be written with only 3 decimal places of sub-second precision (milliseconds) instead of the 6 decimal places (microseconds) available in the source journal timestamps. (October 2026)
DBMI-29669	For application ingestion and replication jobs and database ingestion and replication jobs that have a Databricks target with the CSV file staging, the infa_operation_time audit column might be written with only 3 decimal places of sub-second precision (milliseconds) instead of the 6 decimal places (microseconds) available in the source timestamps. (October 2026)
DBMI-26101,DBMI-26373	For application ingestion and replication or a database ingestion and replication tasks that use the Databricks target, the volume generated automatically as the staging environment is not deleted along with the "temp_jobid" directory when the task is in stopped and failed state. (October 2026)
DBMI-28330	For application ingestion and replication jobs and database ingestion and replication jobs that have a Databricks target and use volume staging, when the initial file upload attempt fails because of network connectivity issues and leaves partially written files, the subsequent retry attempts might fail. The failure occurs because the retry logic does not overwrite existing partial files by default, resulting in an error. The error indicates that the file already exists in staging and overwriting it is not enabled. (April 2026)
DBMI-27152	For application ingestion and replication tasks and database ingestion and replication tasks that have a Databricks target, empty strings ("") in a source string column are incorrectly replicated as NULL to the Databricks target. (April 2026)

Issue	Description
DBMI-26154	Application ingestion and replication jobs and database ingestion and replication jobs that have a Databricks target might fail during the COPY INTO operation if a source table begins with a special character. (April 2026)
DBMI-22028	Application ingestion and replication jobs and database ingestion and replication jobs that have a Databricks Unity Catalog target might fail with an error related to invalid Min Workers or Max Workers values. The job directly parses the values without handling invalid or malformed input. (April 2026)

Google BigQuery V2 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29723	When configuring an application ingestion and replication task or a database ingestion and replication task that has a Google BigQuery target, an option to specify a separate billing project ID at the task level is missing in the task definition. As a result, all BigQuery query costs might be billed to the project defined in the connection, making it impossible to separate FinOps costs across multiple BigQuery projects. (July 2026)

Kafka Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30420	If you configure a case transformation rule for Kafka topic names but do not configure a table renaming rule, the case transformation is not applied. The topic name retains the original source table casing instead of using the requested transformation. This issue affects application ingestion and replication jobs and database ingestion and replication jobs that have a Kafka target and do not have a table renaming rule defined. (October 2026)

Microsoft Azure Synapse Analytics Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29281	For application ingestion and replication tasks and database ingestion and replication tasks that have a Microsoft Azure Synapse Analytics target, when an incremental load job encounters a connection timeout while executing a CREATE EXTERNAL FILE FORMAT statement, the job might fail immediately instead of retrying the timed-out query. Connection timeout errors might not be included in the retry handling logic for Microsoft Azure Synapse Analytics target operations. (July 2026)
DBMI-28958	For application ingestion and replication tasks and database ingestion and replication tasks, when a connection reset occurs immediately after a DROP TABLE statement is executed against the Microsoft Azure Synapse Analytics target, the job might fail with the following error upon reconnection and retry: `2714 "There is already an object named '...' in the database"` The retry logic might only retry the CREATE TABLE statement rather than the full DROP-then-CREATE sequence, so the create attempt might collide with the object that the earlier DROP has not yet fully propagated across Synapse compute nodes. (July 2026)
DBMI-19405	Tests of a Microsoft Azure Synapse Analytics Database Ingestion connection for an application ingestion and replication task or database ingestion and replication task fail with the following error if the user name specified in the connection properties includes a special character such as @ and the target database uses AAD authentication: `Cannot instantiate datasource because of error: Failed to initialize pool: Cannot open server 'noconline.onmicrosoft.com' requested by the login. The login failed. ClientConnectionId: <identifier>` This problem occurs when the Synapse Analytics connector package is missing some required .jar files. These .jar files have now been added in the connector package at `package-AzureDWGen2MI.xxx\package\dw\thirdparty\informatica.azuredwgen2mi`: - activation-1.1 - adal4j-1.4.0 - jcip-annotations-1.0-1 - json-smart-1.3.1 - mail-1.4.7 - nimbus-jose-jwt-9.40 - oauth2-oidc-sdk-5.24.1 (April 2026)

Oracle Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29443	For application ingestion and replication jobs and database ingestion and replication jobs that have an Oracle target, when a combined load job fails or completes the initial load phase, Oracle database connections might not be released, which might lead to connection pool exhaustion and `ORA-00020 "maximum number of processes exceeded"` errors in the Oracle alert log. Connections might persist until the Oracle database is restarted. (July 2026)

Snowflake Data Cloud Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29970	When an application ingestion and replication job or a database ingestion and replication job that has a Snowflake target encounters HTTP 503 errors during merge apply operations or while executing LIST stage queries, the job might fail instead of retrying, even when retry codes are configured via the writerCustomRetryCodes custom property. The retry logic does not cover the autoCommitOff path or the stage initialization path, leaving those failure points unprotected. (July 2026)
DBMI-29745	Application ingestion and replication jobs and database ingestion and replication jobs that have a Snowflake target and use the Superpipe option might intermittently fail with `"JDBC driver encountered communication error. Message: HTTP status=503"` errors during CDC merge operations, even when writerCustomRetryCodes=503,200015 is configured. Certain internal code paths might bypass the retry mechanism, causing job failures on transient Snowflake service unavailability. (July 2026)
DBMI-29594	For application ingestion and replication jobs and database ingestion and replication jobs that have a Snowflake target and use the Superpipe option, when a Snowflake Streaming channel is invalidated due to insert throttling, a deadlock in the InputHelper shutdown sequence might prevent the CDC job from stopping or restarting promptly. The job might continue reporting a "Running" status for several hours after the failure while no data is being processed. (July 2026)
DBMI-29563	For application ingestion and replication jobs and database ingestion and replication jobs that have a Snowflake target and use the Superpipe option, when a table-specific resync is triggered, the operation might incorrectly close all open streaming channels instead of only the channel corresponding to the target table being resynced. This might cause unnecessary channel invalidations and disrupt replication for all other tables in the job. (July 2026)

Issue	Description
DBMI-29448	For application ingestion and replication and database ingestion and replication combined load jobs that have a Snowflake target and use the Superpipe option, when the Soft Delete apply mode is enabled, the initial load phase might fail with "`java.lang.IndexOutOfBoundsException: Index 128 out of bounds for length 128`" for tables that have exactly 128 source columns. The error might not occur when the job runs in Standard Apply mode. (July 2026)
DBMI-28191	For application ingestion and replication and database ingestion and replication jobs that have a Snowflake target and use the Superpipe option, when the job is restarted, the generated recovery table might contain duplicate entries which are not detected at startup. (April 2026)

Application Ingestion and Replication

Fixed issues

There were no fixed issues in this release of Application Ingestion and Replication. Not all monthly releases include fixed issues.

Known issues

There were no known issues in this release of Application Ingestion and Replication. Not all monthly releases include known issues.

Connector and connectivity issues

Read the following pages to learn about fixed issues, known limitations, and third-party limitations that apply to connectors that you can use in your service. Consider the following guidelines for the release notes:

- If a connector has known or fixed issues, you will find it in the release notes even if the connector is not updated in the current release.
- If a connector does not have any notable known or fixed issues, you will not find it in the release notes.
- The (month, year) value in each issue description indicates the release in which the issue was fixed or reported.

Amazon Redshift Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29163	When running an application ingestion and replication incremental load job with a SAP ECC or SAP S4/HANA source using the SAP ODP Extractor connector and an Amazon Redshift target, multiple updates to the same source record within a single CDC cycle might result in duplicate rows in the target table instead of a single row reflecting only the final state. The target apply logic might not consolidate multiple change events for the same primary key within one apply cycle. (July 2026)

Microsoft Dynamics 365 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-31118	For application ingestion and replication tasks that have a Microsoft Dynamics 365 source, when an incremental load job terminates abnormally, the schema cache becomes corrupted and prevents you from restarting the job. All subsequent restart attempts fail immediately with an error indicating that the ignoredChangeHistory attribute is not set. (October 2026)

SAP OData V2 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
AIN-8403	If you restart an aborted application ingestion and replication task that contains an SAP OData source with rules that rename target tables for the selected source objects, the task might cause data duplication when writing to the target. (April 2026)

SAP ODP Extractor Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29163	When running an application ingestion and replication incremental load job with a SAP ECC or SAP S4/HANA source using the SAP ODP Extractor connector and an Amazon Redshift target, multiple updates to the same source record within a single CDC cycle might result in duplicate rows in the target table instead of a single row reflecting only the final state. The target apply logic might not consolidate multiple change events for the same primary key within one apply cycle. (July 2026)

ServiceNow Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
AIN-8624	When you configure an application ingestion and replication task with a ServiceNow source that contains more than 10,000 tables, some of the available tables don't appear for selection under **Source Tables** in the Task Details page. (July 2026)

Database Ingestion and Replication

Fixed issues

The following tables describe recent Database Ingestion and Replication fixed issues. Not all monthly releases include fixed issues.

October 2026 fixed issues

Issue	Description
DBMI-31462	A CDC staging job ignores the **Flush Interval** minutes value you defined in the staging group dialog box and applies a 30-second default flush period instead.

July 2026 fixed issues

Issue	Description
DBMI-29847	When configuring the Access Management transformation settings for a database ingestion and replication task, users who use SAML authentication do not appear in the consumer selection list. The identity service API excludes them because roles are assigned by means of group mappings rather than directly based on the user object.
DBMI-29559	For database ingestion and replication jobs, the pwx_agent_service log files might not be deleted according to the configured log retention period.
DBMI-29381	Database ingestion and replication initial load jobs with partitioning disabled fail with the following message: `Configuration setting <unload.source.partition.count> for <UNLOAD-HR-HR_ALL_ORGANIZATION_UNITS Subtask:0> cannot be parsed into an integer.` The failure occurs because the **Unload Source Partition Count** custom property is internally set to false instead of a positive integer of 1 or more, as expected, even though partitioning is disabled. This fix changes the default for this property to 5, which avoids the failure. This positive partition count is not used as long as parititioning is disabled.

April 2026 fixed issues

Issue	Description
DBMI-28261	Database ingestion and replication combined load jobs that have persistent storage enabled might encounter unexpected errors related to leaving an orphaned InputPersist helper thread running after a task fails.
DBMI-28035	If you add multiple new apply jobs to a CDC staging group and then deploy and run them while the group's CDC Staging Task is running, the CDC Staging Task fails with an error and all of the other apply jobs in the group that were previously running stop.
DBMI-26672	For database ingestion and replication jobs that use the customPrimaryKeys source custom property and go through schema drift changes or multiple restarts, the table information from the cache might contain duplicate primary key entries, causing unsupported primary key changes to be detected and resulting in the job failure.
DBMI-25748	A CDC staging task or a database ingestion and replication task might stall as a result of lock contention between synchronized methods during redeployment. Long-running operations, such as purging, might cause timeouts between Cloud Services and the Agent, preventing the job from redeploying or resuming properly.
DBMI-27577	A database ingestion and replication combined load job that has schema drift enabled might fail after it's restarted if the first schema received in the CDC change stream contains a DDL change such as an added table. In this case, the system might not have the previous schema in cache to process older pending data in the stream.
DBMI-27480	Starting an apply job with a large number of tables in a CDC staging group can cause failures in that job as well as other apply jobs within the group, resulting from the method of submitting the task definition as a long string argument passed directly via the command line to the task processor JVM. The following error might occur: `CreateProcess error=206 ("The filename or extension is too long")`

Issue	Description
DBMI-27381	Database ingestion and replication combined load jobs might fail during the initial unload phase. When the backlog process detects that the unload failed, it clears the name of the table that was being processed. Then, if another unload event occurs, the backlog process checks if it's for the same table. At this point, the backlog process issues the following error because that table name was cleared `Internal logic error - getTransitionBacklogReaderTableNameValue() has no transitionTableNameValue!`
DBMI-27223	Adding the source customPrimaryKeys custom property to an existing incremental load or combined initial and incremental load task might result in the following error: `primary key change not allowed` After you perform a resync operation on a pair of source and target tables, new DML changes are not replicated to the target. If you restart of the job, the same primary key change error occurs for the resynced table.
DBMI-24943, DBMI-27151	When communication between the source CDC capture process and Data Ingestion and Replication fails, the monitoring interface still shows the job status as Up and Running although no data is being processed.
DBMI-22977	The intermediateStorageDirectory and storageBackupDirectory properties of the Database Ingestion service that runs under the Secure Agent must specify unique directory locations. However, no checking for their uniqueness occurs. With this fix, an exception will occur if the properties point to the same location.

Known issues

The following table describes Database Ingestion and Replication known issues. Note that the (*month year*) value in each issue description represents the release in which the issue was found.

Issue	Description
DBMI-31729	When a combined initial and incremental load task in a CDC staging group starts with no CDC activity on the source, the associated apply jobs remain in Backlog stage for up to one hour. Workaround: Set the readerMinutesRestartAdvanceMarkers source custom property to a lower value, such as 15, to reduce the wait time from the default 60 minutes. (October 2026)
CLDDI-51190	If you try to create a database ingestion and replication task with a Microsoft Fabric Data Warehouse, Open Table, or Salesforce Data 360 target in CLAIRE Copilot for Data Integration, Copilot responds that the target type is not valid and doesn't allow you to proceed, even though these target types are supported in manual task creation outside of Copilot. Copilot does not yet support these target types. (April 2026)
DBMI-29328	Database ingestion and replication jobs might fail if you use the Generalize Date data protection technique for Access Management. (April 2026)

Connector and connectivity issues

Read the following pages to learn about fixed issues, known limitations, and third-party limitations that apply to connectors that you can use in your service. Consider the following guidelines for the release notes:

- If a connector has known or fixed issues, you will find it in the release notes even if the connector is not updated in the current release.
- If a connector does not have any notable known or fixed issues, you will not find it in the release notes.
- The (month, year) value in each issue description indicates the release in which the issue was fixed or reported.

Amazon S3 V2 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30152	When replicating data from the Microsoft SQL Server source to Amazon S3 cloud storage as Apache Iceberg tables, numeric(p,s) source columns with scale 0 might be incorrectly mapped to integer type at the target instead of the expected decimal(p,s) type. As a result, CDC write operations might fail with a ClassCastException when the data arrives at the target. (July 2026)
DBMI-28771	Database ingestion and replication incremental load or combined load jobs that have an Amazon S3 target might create a very large cache file that consumes excessive on-disk space if a mass update occurs on the source. With this fix, on-disk space for the caching data has been improved to accommodate this situation. (April 2026)
DBMI-26870	Database ingestion and replication apply jobs in a CDC staging group that have an Amazon S3 target fail with the following error if the **Region Name** connection property is set to a value such as AWS GovCloud (US-West) that is not supported in the older connector package version that is still being used by existing jobs: `Caused by: java.lang.IllegalArgumentException: Cannot create enum from region description AWS GovCloud (US-West) value` (April 2026)
DBMI-26170	When you run a database ingestion and replication job with CDC Staging hosted on an Amazon S3 target, the task might fail if the accessKey and secretKey fields in the S3 connection properties are left empty. (April 2026)

Known issues

The following table describes known issues:

Issue	Description
DBMI-11732	If database ingestion and replication incremental load or combined initial and incremental load jobs replicate LOB source data to an Amazon S3 target and use the CSV format for the target output file, the LOB data appears as empty strings in the target file. (November 2022)
DBMI-2297	Although the Amazon S3 connection properties allow users to specify an IAM role, you cannot use temporary security credentials generated by the AssumeRole method of the AWS Security Token Service API to authorize user access to AWS Amazon S3 resources. (April 2020)

Databricks Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29353	For database ingestion and replication jobs that have a Databricks target and have the auto-tuner option enabled, when the Databricks connection is configured for the SQL Warehouse mode, where the host name field is optional, the job might fail because the auto-tuner attempts to read the null host value without a null check. (July 2026)
DBMI-29114	If you enable supplemental logging on only primary key columns for database ingestion and replication incremental load jobs that have an Oracle source and Databricks target, nulls are replicated to the target for the after images of unmodified source columns. During Update processing, the nulls might overwirte target values, which can cause the source and target to become out of sync. Furthermore, enabling supplemental logging on all columns as a workaround might not be an acceptable if a high volume of changes is expected. (July 2026)

Db2 for i Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29134	Database ingestion and replication combined load jobs that have a Db2 for i source might encounter the following exception and fail when processing a journal entry with Code 'R' and Type 'DR' for a record that was deleted as part of a ROLLBACK operation: `[CDCPUB_10066] TRACE: [IBMiClient setErrorCondition() called. *Error<IBMiUOWManager error, Index 1325 out of bounds for length 325>].*` (July 2026)
DBMI-28584	Database ingestion and replication combined load jobs that have a Db2 for i source and Snowflake target load duplicate record to the target after an insert followed by a rollback operation occurs on the source. In this case, the insert is processed but the rollback operation is not. (April 2026)
DBMI-28399	When starting a database ingestion and replication combined load job that has a Db2 for i source, the job might fail during initial schema drift collection with the following error: `Cannot read the array length because 'esdData' is null.` (July 2026)
DBMI-25956	A database ingestion and replication incremental load or combined load job that has a Db2 for i source and Snowflake target fails if you add or drop a source table while the job and journaling are stopped and then try to resume the job. A way to resume the job without losing data or re-reading a large amount of source data is needed. (April 2026)
DBMI-21838	A database ingestion and replication incremental load job that has a Db2 for i source fails when you first attempt to start it if you set the **Initial Start Point for Incremental Load** property to a specific position or date and time and if the log collector is using only a sequence number and not a log timestamp to access the data. (April 2026)

Db2 for LUW Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-31194	For database ingestion and replication jobs that have a Db2 for LUW source, if you restart an incremental or combined load job after you stop and restart the Database Ingestion service, the job takes hours before data replication resumes. The job might also fail with an inactivity timeout error after 120 minutes of no CDC activity.
DBMI-29270	Database ingestion and replication jobs with a Db2 for LUW source fail after a restart attempt if the source database name is less than eight characters in length. (October 2026)

Issue	Description
DBMI-28869	Database ingestion and replication incremental load or combined load jobs that have a Db2 for LUW source and use the Log-based CDC method might acquire and hold a persistent lock on the capture catalog table until the job ends. (April 2026)
DBMI-28857	Database ingestion and replication incremental load or combined load jobs that have a Db2 for LUW source and use the Log-based CDC method might end abnormally when a large number of DML events are waiting to be published and a timer triggers a request for a capture progress marker. (April 2026)
DBMI-28605	Database ingestion and replication jobs that have a Db2 for LUW source might remove trailing spaces from fixed-length data type columns, such as CHAR and GRAPHIC, resulting in duplicate values and violations of unique constrains. (April 2026)
DBMI-28382	For database ingestion and replication jobs that have a Db2 LUW source and an Oracle target, the deployment process might not incorporate the source character encoding, resulting in Oracle target columns being created with an incorrect VARCHAR2 size, which causes issues with multibyte characters and the job fails. (April 2026)
DBMI-26383	For database ingestion and replication jobs that have a Db2 for LUW source and use the Query-based CDC method, values in scientific notation (for example, 0E-14) are not converted into plain string format before normalization, resulting in a DmlBoundException error. (April 2026)

Known issues

The following table describes known issues:

Issue	Description
DBMI-26309	Database ingestion and replication incremental load and combined load jobs that have a Db2 for LUW source on Linux or Windows and use the Log-based CDC method fail if transient ODBC connection errors, such as network errors or database restarts, occur during communication with Db2. Workaround: Resume the job. (October 2025)

Db2 for z/OS Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-31281	When you run a database ingestion and replication combined load job that has a DB2 for z/OS source and an additional table is added to an existing job, the job fails with a null pointer exception. (October 2026)
DBMI-30168	When the CDC Reader processes a Db2 for z/OS source table that has columns added with ALTER TABLE ADD COLUMN but not yet reorganized, it might misinterpret change records and produce incorrect target data or a runtime error. When a Db2 for z/OS incremental load job restarts with a token that contains low nanosecond values, the job resumes from the wrong log position. (October 2026)
DBMI-29555, DBMI-29233	Database ingestion and replication incremental load jobs that have a Db2 for z/OS source and are in a CDC staging group might fail with a log parsing out-of-bounds error when trying to process an Add Column schema change operation. (July 2026)
DBMI-29285	A database ingestion and replication combined initial and incremental load job that has a Db2 for z/OS source might become permanently stuck in a Queued state. The initial load portion of the job might not start because the process waits indefinitely for a marker event that the log reader already passed. (July 2026)
DBMI-29270	Database ingestion and replication jobs with a Db2 for z/OS source fail after a restart attempt if the source database name is less than eight characters in length. (October 2026)
DBMI-29237	After a CDC staging job with a Db2 for z/OS source is restarted, CAPTURE_PROGRESS markers might be sent with a sequence value that resets to an earlier position rather than continuing from the last recorded position. This might cause apply jobs downstream of the staging group to report "previous sequence" errors and stall. (July 2026)
DBMI-28872	Database ingestion and replication incremental load jobs that have a Db2 for z/OS source might stop capturing change data and hang indefinitely because the log parser consumer thread can't acquire a lock on a parser queue. (April 2026)
DBMI-28561, DBMI-28989	Database ingestion and replication tasks that have a Db2 for z/OS source generate log timestamps from the Db2 WLM stored procedure, which might contain errors in calculations involving sub-second values below 0.1. (April 2026)
DBMI-28506	When a database ingestion and replication incremental load job that has a Db2 for z/OS source encounters certain SYSIBM.SYSINDEXES catalog events that contain fewer columns than expected, the schema drift parser might raise a ClassCastException and cause the job to fail. The error might occur because the code attempts to read index-name data from an incomplete catalog record without first verifying that the required data is present. (July 2026)

Issue	Description
DBMI-28143	For database ingestion and replication incremental load or combined load jobs that have a Db2 for z/OS source, the latency of writing changes to the target can't be calculated as the difference between the values in the Operation Time and Last Replicated Time metadata columns because the Operation Time records the log time of the DML operation. Instead, for latency reporting, the Operation Time should report of the UOW commit time. (April 2026)
DBMI-28092	Performance statistics for database ingestion and replication incremental load and combined load jobs might indicate a delay in processing in an environment with many sub-second requests for Db2 log data that return few records. Each request creates a record queue. After processing a queue, the Db2 reader waits 1 second by default for the next queue. With this fix, the the wait time is reduced to 100 milliseconds to get the log data from the queues faster and reduce latency. (April 2026)
DBMI-27718	Database ingestion and replication incremental load and combined load jobs that have a Db2 for z/OS need to write enhanced statistics messages to the job log to help with diagnosing CDC performance issues. (April 2026)
DBMI-27717	If database ingestion and replication incremental load and combined load jobs that have Db2 for z/OS source encounter Db2 compensation records in the log, a slowdown in CDC processing might occur. The processing of compensation records needs to be improved. (April 2026)
DBMI-27554	Database ingestion and replication incremental load or combined load jobs might generate multiple inactive connections on a Db2 for z/OS source system if the generation of capture progress markers is enabled and the job disconnects from the subsystem and then retries the connection multiple times. While the job is retrying the connection, each capture progress marker request creates an inactive connection thread on the Db2 subsystem. (April 2026)

Google BigQuery V2 Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-28558	Database ingestion and replication jobs that have a Google BigQuery target run temporary table creation and merge queries that have significantly higher compute costs when using Soft Deletes apply mode, as compared to the costs when using Standard apply mode. This fix adds the CLUSTER BY clause in the DDL for new tables to optimize the queries run on the target when using Soft Deletes mode. This solution requires the source to have a primary key or unique index. (April 2026)

Google Cloud Storage V2 Connector

Known issues

The following table describes known issues:

Issue	Description
DBMI-11732	If database ingestion and replication incremental load or combined initial and incremental load jobs replicate LOB source data to a Google Cloud Storage target and use the CSV format for the target output file, the LOB data appears as empty strings in the target file. (November 2022)

Kafka Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29543	Database ingestion and replication incremental load and combined load jobs that have a PostgreSQL source and Kafka target fail with the following error shortly after startup when comparing the source and target checkpoints: `ClassCastException: object <SimpleSequenceToken> cannot be used to determine relative stream positon because it in not an instance of PayloadSequenceAsToken` (July 2026)

Microsoft Azure Data Lake Storage Gen2 Connector

Known issues

The following table describes known issues:

Issue	Description
DBMI-11732	If Database ingestion and replication incremental load or combined initial and incremental load jobs replicate LOB source data to a Microsoft Azure Data Lake Storage Gen2 target and use the CSV format for the target output file, the LOB data appears as empty strings in the target file. (November 2022)

Microsoft Fabric Data Warehouse Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30118	Database ingestion and replication jobs that have a Microsoft Fabric Data Warehouse target might fail when the source table contains decimal columns with precision greater than 28. (October 2026)
DBMI-28927	In a combined load database ingestion and replication task enabled with the **Audit** apply mode, the INFA_OPERATION_TYPE columns are not populated in the initial unload phase, which results in NULL values in operational columns. (October 2026)

Known issues

The following table describes known issues:

Issue	Description
DBMI-31696	For a Microsoft Fabric Data Warehouse target, when a task processes multiple tables and DML operations for different tables arrive simultaneously, the task fails with the following error: `Failed to execute target queries: Failed to execute transactional statements.` (October 2026)
DBMI-31507	For database ingestion and replication tasks that have Microsoft Fabric Data Warehouse and Open Table targets, the CYCLE_ID and LAST_REPLICATED_TIME columns are not populated, which results in NULL values in these columns. This issue affects the following job types: initial load jobs, the initial load portion of the combined load jobs, and incremental load and combined load jobs that use the Soft Deletes apply mode. (October 2026)
DBMI-27827	If you enable the **Add Operation Sequence** option for a combined load database ingestion and replication task with the **Audit** apply mode to load data from Oracle to Microsoft Fabric Data Warehouse, the task fails. (April 2026)

Microsoft SQL Server Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30743	Database ingestion and replication jobs that have a SQL Server source might fail after the July 2026 upgrade when the source database is read-only. When the CDC management attempts to re-enable a capture instance on the read-only database, SQL Server rejects the operation, causing the reader process to crash and drop the session with the following error: `[CDCPUB_10066] PwxSessionException: SessionGet failed / End of socket stream reached` (October 2026)
DBMI-30152	When replicating data from the Microsoft SQL Server source to Amazon S3 cloud storage as Apache Iceberg tables, numeric(p,s) source columns with scale 0 might be incorrectly mapped to integer type at the target instead of the expected decimal(p,s) type. As a result, CDC write operations might fail with a ClassCastException when the data arrives at the target. (July 2026)
DBMI-29777	Database ingestion and replication initial load, incremental load, or combined load jobs that have a SQL Server source and PostgreSQL target might fail if a source column contains the \0 (NUL) character, which is not supported on the target. With this fix, you can avoid the failure by using the formatEncoderReplaceUnsupportedChars and formatEncoderReplacementChar target custom properties to replace the unsupported charcter with another character on the target. (July 2026)
DBMI-29545, DBMI-29547	Database ingestion and replication jobs that have a SQL Server source and contain source tables with a large number of columns might fail to update the CDC tables after a DDL change. For tables where combined column names exceed 4000 characters, this operation might fail with a TDS RPC protocol stream error (8016). The job continues to run but with an outdated CDC schema. (July 2026)
DBMI-29128	Database ingestion and replication jobs that have a SQL Server source and use the CDC Tables capture method might intermittently fail with a fatal ODBC driver error (HY000) when the job attempts to query the cdc.lsn_time_mapping system table but encounters a transient row versioning conflict. (July 2026)
DBMI-28957	For database ingestion and replication jobs that have a SQL Server source and use the Query-based CDC capture method, the source BIT column might be incorrectly replicated as the strings "false"/"true" instead of "0"/"1", This behavior can cause jobs to fail with the `ORA-12899: value too large for column` error when the target column is sized to hold only one character. (July 2026)
DBMI-28492	Database ingestion and replication incremental and combined load jobs that have a SQL Server source might assign duplicate sequence values to change records, potentially causing inconsistencies or data corruption in the target tables. (April 2026)

Issue	Description
DBMI-27721	Database ingestion and replication combined load jobs with a SQL Server source that uses Query-based CDC and a Snowflake target might fail with the following type of error after running for a few hours: `Writer <CDC-UNLOAD-data_store.dbo.employee Async Object Distributor 1> failed to process <END_OF_TABLE_DATA> marker event received for table <data_store.dbo.fx_rate> with stream id <1> as it is not applicable to flush strategy <TimeoutOrRowCount>` (April 2026)
DBMI-27160	During CDC processing, database ingestion and replication jobs that have a SQL Server source might hang and stop sending change data to the target because the SQL Server DataDirect ODBC driver fails to recognize the timeout set for SQL query execution. (April 2026)
DBMI-26134	Database ingestion and replication incremental load or combined loads with a Snowflake target that uses Snowpipe Steaming (Superpipe option) might fail to restart after a channel invalidation error. (April 2026)
DBMI-26127	When a deadlock occurs during the FETCH call for a a database ingestion and replication job that has a SQL Server source, the job might fail without attempting any retries. (April 2026)
DBMI-25390	Database ingestion and replication incremental load and combined load jobs that have a SQL Server target might fail to drop the temporary xxx_VER table that's used in the merge apply cycle at job restart, which can cause the job to fail. This table should have been deleted as part of task cleanup processing. (April 2026)

MongoDB Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29396	When running a database ingestion and replication initial load that has a MongoDB sharded cluster source, the countDocuments() call issued before iterating the collection cursor might take several hours to return on large collections, causing the previously opened cursor to expire. As a result, the job might fail with the following error: `Command failed with error 43 (CursorNotFound).` (July 2026)
DBMI-27212	For database ingestion and replication combined load jobs that have a MongoDB source and use the Audit apply mode with the Oracle target, the order of audit and non-audit columns in the source schema might not match the order of values on the target side, causing values to be inserted into incorrect columns. (April 2026)
DBMI-26342	Database ingestion and replication combined load tasks that have a MongoDB source and use the MongoDB Atlas instance might fail. (July 2025)

MySQL Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30126, DBMI-29816	Database ingestion and replication incremental load and combined load jobs might drop the fractional-seconds part of DATETIME, TIMESTAMP, and TIME column values during the incremental load phase, causing the target values not to match the source. The initial load phase is not affected. (October 2026)
DBMI-29725	Database ingestion and replication jobs running on the Task Processor might produce incorrect output when the target format is Avro or Parquet. The Task Processor reuses a shared Avro schema cache that can retain a stale schema from a previous job, causing new columns from the current job to be missing from the output. (July 2026)
DBMI-28896	The initialization of database ingestion and replication initial load jobs that have a MySQL Community Edition source might be slower than expected. The problem occurs because the jobs try to use the Progress DataDirect for MySQL driver, which is not supported with Community Edition. After three unsuccessful connection attempts, the jobs switch to the recommended MySQL JDBC driver that you downloaded from the MySQL Community Downloads website. By default, the MySQL JDBC driver should be used for the first and subsequent connection attempts to avoid these delays. (July 2026)

Issue	Description
DBMI-28599	For database ingestion and replication jobs that have a MySQL source, the SSL Attributes (For MySQL 8.X driver) defined in the connection properties are not being honored. (April 2026)
DBMI-28074	Database ingestion and replication jobs that have a MySQL source fail with following error if the MySQL sql_mode system variable includes the ANSI_QUOTES mode: `Caused by: java.sql.SQLSyntaxErrorException: Unknown column 'trust_reg_reporting' in 'where clause'` (April 2026)
DBMI-27622	Database ingestion and replication jobs that have a MySQL source and Snowflake target might fail with the following error after insert or update operations occur on the Sequence column in the INFORMATICA_CDC_RECOVERY table causing unexpectedly long restart token strings: `[Failed to execute target queries. Error: DML operation to table xxx.INFORMATICA_CDC_RECOVERY failed on column SEQUENCE with error: String '<long_string>' is too long and would be truncated].` (April 2026)

Open Table Connector

Known issues

The following table describes known issues:

Issue	Description
DBMI-31507	For database ingestion and replication tasks that have Microsoft Fabric Data Warehouse and Open Table targets, the CYCLE_ID and LAST_REPLICATED_TIME columns are not populated, which results in NULL values in these columns. This issue affects the following job types: initial load jobs, the initial load portion of the combined load jobs, and incremental load and combined load jobs that use the Soft Deletes apply mode. (October 2026)

Oracle Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-29671	Database ingestion and replication jobs that have an Oracle source might regularly fail with the following error, particularly when processing a large number of high-volume tables: `PWX-36000 ORAD: Internal error Set failed for unwanted list for obj` The error might be triggered when the CDC reader encounters DML events for object IDs that fall outside the range tracked in the internal unwanted-object filter list. (July 2026)
DBMI-29209	When an Oracle source connection is configured to use the Oracle JDBC Thin driver for the log-based CDC method, the driver selection might not be propagated to the metadata manager component, which might always default to the DataDirect driver instead. This might cause metadata operations to fail in environments where only the Oracle JDBC Thin driver is permitted or reachable. (July 2026)
DBMI-29114	If you enable supplemental logging on only primary key columns for database ingestion and replication incremental load jobs that have an Oracle source and Databricks target, nulls are replicated to the target for the after images of unmodified source columns. During Update processing, the nulls might overwirte target values, which can cause the source and target to become out of sync. Furthermore, enabling supplemental logging on all columns as a workaround might not be an acceptable if a high volume of changes is expected. (July 2026)
DBMI-28969	Database ingestion and replication incremental load and combined load jobs with an Oracle source and Oracle target might result in data loss on the target if a source table schema change, such as a partition add or drop, occurs during CDC. The schema change triggers a flush of data to the target only for the table with that schema and then resets the job restart point. If the job is restarted before the normal flush processing occurs for the other source tables, the data captured for those tables is not applied to the target. (July 2026)
DBMI-28939	Database ingestion and replication combined initial and incremental load jobs with an Oracle source that contains CLOB columns with the NOT NULL constraint might fail when processing an insert operation that writes an empty LOB to the table followed by an update operation that populates the LOB column with data. In this situation, the following error is issued: `PWX-36200 ORAD: Column conversion error: [36202] NULL value found for column that is not nullable for Table name=<table_name>, column[25] segment[25] RO_TERMS_AND_CONDITIONS. Detail Type = CLOB, Oracle Type = CLOB.` (April 2026)
DBMI-28770	If process ID memory usage on Linux reaches 60%, Secure Agent services might stop, causing database ingestion and replication jobs to fail. This problem occurs because of a memory leak in the Oracle Client. (April 2026)

Issue	Description
DBMI-28382	Deployment of database ingestion and replication jobs that have a Db2 for LUW source and an Oracle target might not incorporate the source character encoding. In this case, Oracle target columns might be created with an incorrect VARCHAR2 size, which can cause multibyte-character issues and the job to fail. (April 2026)
DBMI-27796	If database ingestion and replication jobs capture change data from Oracle CLOB columns that use BASICFILE storage, unreadable or garbled characters might be replicated to the corresponding target columns. (April 2026)
DBMI-27222	Database ingestion and replication incremental load and combined load jobs that have an Oracle source and Oracle target might write corrupted data to the target if an update to an existing record and an insert of a new record occurs in the same transaction. (April 2026)
DBMI-27212	For database ingestion and replication combined load jobs that have a MongoDB source and use the Audit apply mode with the Oracle target, the order of audit and non-audit columns in the source schema might not match the order of values on the target side, causing values to be inserted into incorrect columns. (April 2026)
DBMI-27146	Database ingestion and replication combined load jobs with an Oracle source and Oracle target fail with the following error if you specify a lowercase prefix such as “audit” in the **Prefix for Metadata Columns** property for the task: `java.lang.NumberFormatException: Character I is neither a decimal digit number, decimal point, nor "e" notation exponential mark.` (April 2026)
DBMI-27063	After Oracle performs a QMI-array insert operation on an Oracle index-organized table (IOT) that contains only key data, database ingestion and replication jobs that capture change data from that table might fail with a segmentation fault (SIGSEGV) in the PWXORL code. (April 2026)
DBMI-27055	Database ingestion and replication jobs that have an Oracle physical standby source might time out with the following message: `PWX-36564 ORAD: LRdr for member 2 has been unresponsive for more than the timeout threshold of 600. Last response time was 2025/09/30 06:54:09.` The problem occurs because an incomplete LWN was written to the standby logs. Log transport abandons the standby log in favor of an archive log, which causes a code loop. (April 2026)

Issue	Description
DBMI-25346	Database ingestion and replication incremental load and combined load jobs that have an Oracle source and use the BFILE log access method might report stale buffer files when the files are in an Oracle server file system on Linux. If the default settings for the status check interval and archive wait time are too high, this cache storage problem might not be detected or detection might be delayed. (April 2026)
DBMI-23048, DBMI-21130	If you configure an Oracle Database Ingestion connection to use both the Oracle JDBC Thin driver and SSL encryption, when you try to test the connection, the test fails. (April 2026)

Known issues

The following table describes known issues:

Issue	Description
DBMI-31127	Row-level filtering fails for a database ingestion and replication CDC task with an Oracle source configured to capture only modified columns when supplemental logging is not enabled for all the source columns. The filter reads the wrong column value and the job fails. Workaround: Enable supplemental logging for all columns in the Oracle source tables when you use row-level filtering in CDC jobs. (October 2026)
DBMI-28993	Database ingestion and replication tasks that have an Oracle source and a Microsoft Fabric Data Warehouse or Salesforce Data 360 target do not replicate Oracle TIMESTAMP_WITH_LOCAL_TIMEZONE and TIMESTAMP_WITH_TIMEZONE data types to the target. (July 2026)
DBMI-24444	For database ingestion and replication tasks that have an Oracle source and use the Audit apply mode, LOB data might not be correctly replicated to the target. When a non-LOB data is updated, the remaining columns are replicated as NULLs. (May 2025)
DBMI-23360	Database ingestion and replication incremental load or combined load jobs that have an Oracle source with XML columns and an Oracle target might fail with the following Oracle error when processing a DML operation: `[Oracle JDBC Driver][Oracle]ORA-00932: inconsistent datatypes: expected - got CLOB. Error code: 932` (April 2025)

Issue	Description

DBMI-19145	Database ingestion and replication jobs that use the Log-based CDC method and have an Oracle source with only LOB columns selected and no primary key do not create subtasks on the tables. As a result, the jobs cannot capture change data from the tables. Messages such as the following are written to the log:
	`[DBMIP_23026] The process [CDC-LOBS] with the thread ID [9708] encountered new table [lob_table_name]. The new table is excluded from processing.`
	`[CDCPUB_10066] TRACE: [PwxCDCRequestProcessor.askSchema() returned: Don't capture.].`
	Workaround: In the task, select some non-LOB source columns, in addition to the LOB columns, for replication if you want to continue using the table without a primary key.
	(July 2024)

DBMI-14767	If a database ingestion and replication job uses an Oracle Database Ingetion source connection that's configured to use the Oracle JDBC Thin driver, the job can replicate up to 39 digits from source columns that have a numeric data type to the target. If a source numeric value has 40 or more digits, the fortieth digit and all additional digits are replaced by zeroes (0) on the target. (September 2023)
DBMI-13605	The Oracle Database Ingestion connection properties page includes no property for entering JDBC connection properties, such as EncryptionLevel, when they're needed. Workaround: In the **Service Name** field, you can add JDBC connection properties after the Oracle SID value, using a semicolon (;) as the separator. (April 2023)
DBMI-10794	Oracle source columns with the TIMESTAMP WITH TIME ZONE data type are supported only for initial load jobs. Workaround: To enable database ingestion and replication incremental load and combined load jobs to process change data from TIMESTAMP WITH TIME ZONE columns, set the source custom property pwx.cdcreader.oracle.option.additional ENABLETIMSTAMPWITHTZ to Y. (July 2022)

PostgreSQL Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30153	When you run a database ingestion and replication combined load job that has a PostgreSQL source and select individual columns for data replication, after stopping the job and resuming it with options, Resync (refresh) and Resync (retain) might be missing from the list of available options. (October 2026)
DBMI-30864	Database ingestion and replication incremental load jobs that have a PostgreSQL source might intermittently stop capturing data and become unresponsive without any error, until the built-in inactivity watchdog automatically restarts the job after 120 minutes, causing a 2-hour data loading delay. (October 2026)

Issue	Description
DBMI-30773, DBMI-31235	When running a database ingestion and replication incremental load job that has a PostgreSQL source, the replication slot read query might be cancelled after 60 seconds with the following message: `canceling statement due to user request` The replication slot remains in an active state and causes the job to loop indefinitely until a task retry exception forces a restart after approximately 120 minutes. (October 2026)
DBMI-29777	Database ingestion and replication initial load, incremental load, or combined load jobs that have a SQL Server source and PostgreSQL target might fail if a source column contains the \0 (NUL) character, which is not supported on the target. With this fix, you can avoid the failure by using the formatEncoderReplaceUnsupportedChars and formatEncoderReplacementChar target custom properties to replace the unsupported charcter with another character on the target. (July 2026)
DBMI-29543	Database ingestion and replication incremental load and combined load jobs that have a PostgreSQL source and Kafka target fail with the following error shortly after startup when comparing the source and target checkpoints: `ClassCastException: object <SimpleSequenceToken> cannot be used to determine relative stream positon because it in not an instance of PayloadSequenceAsToken` (July 2026)
DBMI-29480	A database ingestion and replication incremental load or combined initial and incremental load job that has a PostgreSQL source might fail if an apply cycle is interrupted before completion, causing the system to mistakenly use a table-level checkpoint to prematurely prune replication slot data. On restart, the job crashes with a `Failed to read persisted records for LSN` error because the required restart location no longer exists in the pruned persistence table. (July 2026)
DBMI-29410	When a database ingestion and replication combined load job with a PostgreSQL source using local table-based persistence is restarted, the job might not resume and might fail with the following error: `Persistence restart failed - Error encountered while locating restart LSN=... Xid=... in persistence table` The failure is caused by a type mismatch in the persistence table query with extremely large transaction IDs. This prevents the restart information from being located correctly. (July 2026)
DBMI-27474	Database ingestion and replication incremental load jobs that have a PostgreSQL source and Snowflake target tries to re-create some target tables during redeployment, even if no DDL changes occurred on the corresponding source tables. In this case, HTTP 403 errors appear in the Metadata Manager logs, but the job does not fail as expected. Instead, it continues and drops the target tables. (April 2026)

Issue	Description
DBMI-26959	Database ingestion and replication tasks that have a PostgreSQL source might display error messages if not all sub-partitions are included in the publication. (April 2026)
DBMI-26936	Database ingestion and replication combined load jobs that have a PostgreSQL source and Google BigQuery target might fail with the following error: `java.lang.RuntimeException: java.lang.NullPointerException: Cannot invoke "java.util.Map.get(Object)" because "this.desriptorsByDataType" is null` The problem occurs when multiple tasks are running simultaneously after the Secure Agent starts. This situation can occasionally lead to a condition where the internal configuration related to data types is in an inconsistent state. If you resume the failed job, it runs without an issue. (April 2026)

Known issues

The following table describes known issues:

Issue	Description
DBMI-26256	For database ingestion and replication incremental load and combined initial and incremental load tasks that have a PostgreSQL source and use the pgoutput replication plug-in, the generated CDC script does not include partitioned tables in the publication along with the primary tables. Workaround: Edit the CDC script manually to add the partitions to the publication for any desired partitioned tables and execute the script. The CDC script might need to be altered if adding a table to the publication results in an error. (October 2025)

SAP HANA Database Ingestion Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-30521	When you use log-based CDC with a SAP HANA source and the backup server is temporarily unreachable, SAP HANA might generate duplicate redo-log files with identical start and end block IDs. The duplicate files cause incremental load and combined load jobs to get stuck in the BACKLOG stage. (October 2026)
DBMI-26633	When one or more database ingestion and replication combined load jobs that have an SAP HANA Trigger-based CDC source are running and writing metadata about DML changes to the same PKLOG table, a subtask on a source table might remain in the Queued state. This situation occurs when the rate of publishing metadata to the PKLOG table is greater than the rate of consuming metadata from the PKLOG table. If no record is found for the queued table, the task waits either to EOL or until it receives a record. (April 2026)

Known issues

The following table describes known issues:

Issue	Description

DBMI-24971	If a CDC Staging Task is running with associated apply jobs to perform SAP HANA Log-based CDC and a schema drift DDL change occurs on a source table, the schema drift change is ignored with a warning. The apply job that processes the table might fail with the following error, even if the job is restarted:
	`[DBMIP_23002] The process [CDC-CDC_STAGING InputEndpoint Helper] with the thread ID [HanaLogCollector pwxHanaCDCApi.readNextData() failure. null] received an unexpected error from the PowerExchange Capture Service. Error:`
	`…`
	`Caused by: java.nio.BufferUnderflowException`
	Workaround: None.
	(July 2025)

Snowflake Data Cloud Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
DBMI-28162	Database ingestion and replication incremental load or combined load jobs that have a Snowflake target that uses the Superpipe (Snowpipe Streaming) option might replicate duplicate records to the target. The problem occurs if the job deletes a record and then re-inserts the record into a source table that has no primary key while the job is stopped. When the job restarts, if it tries to merge data from columns that contain null values and the merge query is not null aware, deletes fail and inserts can then result in duplicate rows on the target. (April 2026)
DBMI-27721	Database ingestion and replication combined load jobs with a SQL Server source that uses Query-based CDC and a Snowflake target might fail with the following type of error after running for a few hours: `Writer <CDC-UNLOAD-data_store.dbo.employee Async Object Distributor 1> failed to process <END_OF_TABLE_DATA> marker event received for table <data_store.dbo.fx_rate> with stream id <1> as it is not applicable to flush strategy <TimeoutOrRowCount>.` (April 2026)
DBMI-27622	Database ingestion and replication jobs that have a MySQL source and Snowflake target might fail with the following error after insert or update operations occur on the Sequence column in the INFORMATICA_CDC_RECOVERY table causing unexpectedly long restart token strings: `[Failed to execute target queries. Error: DML operation to table xxx.INFORMATICA_CDC_RECOVERY failed on column SEQUENCE with error: String '<long_string>' is too long and would be truncated].` (April 2026)

Issue	Description
DBMI-27484	In a database ingestion and replication combined load job that has a Snowflake target, if a Resync(Refresh) operation is requested on a table that is already in the Normal state, the job might stop with an invalid channel error. The problem occurs because the resync triggers an unload operation, which re-creates the LOG table. This action causes the channel in the CDC stream to become invalid. (April 2026)
DBMI-27474	Database ingestion and replication incremental load jobs that have a PostgreSQL source and Snowflake target tries to re-create some target tables during redeployment, even if no DDL changes occurred on the corresponding source tables. In this case, HTTP 403 errors appear in the Metadata Manager logs, but the job does not fail as expected. Instead, it continues and drops the target tables. (April 2026)
DBMI-27355	Database ingestion and replication jobs that have a Snowflake target that uses the Superpipe (Snowpipe Streaming) option might fail with the following error after running for several hours: `Cannot invoke "com.infa.rat.dbmi.writer.impl.recovery.snowflake_streaming.SnowflakeStreamingRecoveryManager.getCurrentCycleNumber()" because "this.recoveryManager" is null` While the job was running, the numbers of processed rows and written rows indicated job progress, but the target objects remained in the Unloading state with a much lower record count than the source. Also, the target sometimes had duplicate records or missing records. (April 2026)
DBMI-27118	You might encounter out-of-memory issues, particularly when using a Snowpipe Streaming (Superpipe) for a Snowflake target. Now you can use Shenandoah GC to get rid of OOM issues and promptly release memory back to the operating system if a job fails or is aborted. (April 2026)
DBMI-26871	During a database ingestion and replication combined load job that has a Snowflake target using the Superpipe (Snowpipe Streaming) option, initial unload processing might trigger a flush of data to the target every hour, which unintentionally causes the stream channel to close and then to reopen in the middle of the unload. With this fix, a single flush is performed at the end of the job. (April 2026)
DBMI-26829	A database ingestion and replication job that has a Snowflake target using the Superpipe (Snowpipe Streaming) option fails while validating a target table because of the following column index error: `[CDCPUB_10066] TRACE: [Process <CDC-dbo OutputEndpoint Helper 7> with thread ID <147> failed to fetch the target table format for table <ADDRESS_BASE_LOG> . Error was <Invalid column index: 22>.].` (April 2026)
DBMI-26134	Database ingestion and replication incremental load or combined loads with a Snowflake target that uses Snowpipe Steaming (Superpipe option) might fail to restart after a channel invalidation error. (April 2026)

File Ingestion and Replication

Fixed issues

There were no fixed issues in this release of File Ingestion and Replication. Not all monthly releases include fixed issues.

Known issues

The following table describes File Ingestion and Replication known issues. Note that the (*month*, *year*) value in each issue description represents the release in which the issue was reported.

Issue	Description
FMI-4132	If you are assigned the Monitor role and try to view the File Watch Task step details of a taskflow on the Monitor page, the Current Run tab fails to load. This issue doesn't occur if you are assigned the Admin role or have execute permissions for the taskflow. Workaround: To view the File Watch Task step details, use an account that is assigned an Admin role or has execute permissions for the taskflow. (October 2026)

Connector and connectivity issues

Read the following pages to learn about fixed issues, known limitations, and third-party limitations that apply to connectors that you can use in your service. Consider the following guidelines for the release notes:

- If a connector has known or fixed issues, you will find it in the release notes even if the connector is not updated in the current release.
- If a connector does not have any notable known or fixed issues, you will not find it in the release notes.
- The (month, year) value in each issue description indicates the release in which the issue was fixed or reported.

Oracle Cloud Object Storage Connector

Known issues

The following table describes known issues:

Issue	Description
CCON-113412	When you use GZIP compression in a file ingestion and replication task to load files to an Oracle Cloud Storage target, the task uploads files from the first batch multiple times. This issue occurs because gzipped files are not deleted from the workspace after upload, so the task includes them again in subsequent uploads. Workaround: To avoid this issue, perform one of the following steps: - Disable file compression by setting the **File Compression** option to **None** before the transfer. - If you want to use GZIP, either increase the batch size to a value higher than the total number of files when there are 20 or fewer files, or apply GZIP compression through the file processing actions instead of in the target properties. (July 2026)

Streaming Ingestion and Replication

Fixed issues

There were no fixed issues in this release of Streaming Ingestion and Replication. Not all monthly releases include fixed issues.

Known issues

There were no known issues in this release of Streaming Ingestion and Replication. Not all monthly releases include known issues.

Connector and connectivity issues

Read the following pages to learn about fixed issues, known limitations, and third-party limitations that apply to connectors that you can use in your service. Consider the following guidelines for the release notes:

- If a connector has known or fixed issues, you will find it in the release notes even if the connector is not updated in the current release.
- If a connector does not have any notable known or fixed issues, you will not find it in the release notes.
- The (month, year) value in each issue description indicates the release in which the issue was fixed or reported.

Google PubSub Connector

Fixed issues

The following table describes fixed issues:

Issue	Description
SI-5473	The Google PubSub connection test fails even if you provide valid connection parameters. (October 2026)

Microsoft Azure Event Hubs Connector

Known issues

The following table describes known issues:

Issue	Description
SI-6005	The Microsoft Azure Event Hubs test connection fails with a `RxJavaCallAdapterFactory` error. Workaround: You can ignore the failure message, save the connection, and successfully use it in streaming ingestion and replication tasks. (April 2026)

Informatica Global Customer Support

You can contact a Global Support Center through the Informatica Network or by telephone.

To find online support resources on the Informatica Network, click **Contact Support** in the IDMC Help menu to go to the **Cloud Support** page. The **Cloud Support** page includes system status information and community discussions. Log in to Informatica Network and click **Need Help** to find additional resources and to contact Informatica Global Customer Support through email.

The telephone numbers for Informatica Global Customer Support are available from the Informatica web site at https://www.informatica.com/services-and-training/support-services/contact-us.html.