Informatica® Data Ingestion and Replication
October 2026

Data Ingestion and Replication What's New

Informatica Data Ingestion and Replication Data Ingestion and Replication What's New
October 2026

© Copyright Informatica LLC 2019, 2026

This software and documentation are provided only under a separate license agreement containing restrictions on use and disclosure. No part of this document may be reproduced or transmitted in any form, by any means (electronic, photocopying, recording or otherwise) without prior consent of Informatica LLC.

U.S. GOVERNMENT RIGHTS Programs, software, databases, and related documentation and technical data delivered to U.S. Government customers are "commercial computer software" or "commercial technical data" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, the use, duplication, disclosure, modification, and adaptation is subject to the restrictions and license terms set forth in the applicable Government contract, and, to the extent applicable by the terms of the Government contract, the additional rights set forth in FAR 52.227-19, Commercial Computer Software License.

A current list of Informatica trademarks is available on the web at https://www.informatica.com/trademarks.html. Other company and product names may be trade names or trademarks of their respective owners.

Portions of this software and/or documentation are subject to copyright held by third parties. Required third party notices and AI assistant acknowledgements are included with the product.

The information in this documentation is subject to change without notice. If you find any problems in this documentation, report them to us at infa_documentation@informatica.com.

Informatica products are warranted according to the terms and conditions of the agreements under which they are provided. INFORMATICA PROVIDES THE INFORMATION IN THIS DOCUMENT "AS IS" WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING WITHOUT ANY WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND ANY WARRANTY OR CONDITION OF NON-INFRINGEMENT.

Publication Date: 2026-10-16

Table of Contents

Preface 5

Chapter 1: October 2026 6

Changelog October 2026 6

Important notices and upcoming changes 7

- Preview lifted 8
- Reinstated support 8
- Deprecated support 8

New features and enhancements 9

- Common 9
- Application Ingestion and Replication 12
- Database Ingestion and Replication 12
- File Ingestion and Replication 14

Changed behavior 15

Chapter 2: July 2026 17

Changelog July 2026 17

Important notices 17

New features and enhancements 18

- Database Ingestion and Replication 18
- File Ingestion and Replication 18

Changed behavior 19

Chapter 3: April 2026 20

Changelog April 2026 20

Important notices 21

New features and enhancements 22

- Common 22
- Application Ingestion and Replication 25
- Database Ingestion and Replication 25
- File Ingestion and Replication 28
- Streaming Ingestion and Replication 28

Changed behavior 29

Chapter 4: November 2025 31

Changelog November 2025 31

Important notices 31

New features and enhancements 32

- Common 32

Chapter 5: October 2025 . . . . . **33**
Important notices 33
New features and enhancements 34
Common 34
Database Ingestion and Replication 37
Changed behavior 39

Index **42**

Preface

Read *Data Ingestion and Replication What's New* for a brief overview of new features in the Data Ingestion and Replication service, which includes Application Ingestion and Replication, Database Ingestion and Replication, File Ingestion and Replication, and Streaming Ingestion and Replication.

CHAPTER 1

October 2026

This section provides information about new features, enhancements, and behavior changes in the October 2026 release of Data Ingestion and Replication.

Changelog October 2026

Read the changelog to learn about significant changes to this publication for this release.

The following table describes changes in this publication:

Publication Date	Type of Change	Section Title
October 16, 2026	Added	"SAP HANA log-based CDC archive log discovery" on page 15
October 16, 2026	Modified	"Apache NiFi 2.5.0 runtime engine for Streaming Ingestion and Replication" on page 8
October 16, 2026	Added	"Table-level error isolation in incremental load tasks" on page 12
October 16, 2026	Modified	"Microsoft Fabric Lakehouse target" on page 9
October 16, 2026	Added	"Path validation for local directories in Google Cloud Storage V2 tasks" on page 16
October 16, 2026	Modified	"Audit metadata columns for Microsoft Fabric OneLake targets" on page 9
October 16, 2026	Added	"Billing project override for Google BigQuery targets" on page 11
October 16, 2026	Modified	"Multiple object distributors for Cloud Data Lake targets" on page 11
October 16, 2026	Added	"Auto-tuning enhancements" on page 12

Publication Date	Type of Change	Section Title
October 16, 2026	Modified	"Number of Rows in Output File field location" on page 15
September 24, 2026	Modified	"Granular privileges for deploy, redeploy, and undeploy actions " on page 13
September 24, 2026	Added	"Java Runtime Environment 17 requirement for Data Ingestion and Replication Command-Line Interface" on page 9
September 24, 2026	Added	"Job management from the task configuration page" on page 11
September 24, 2026	Added	"Serverless runtime environments extended to Cloud Data Lake targets" on page 11
September 24, 2026	Added	"Automatic agent switchover extended to Cloud Data Lake targets" on page 11
September 24, 2026	Added	"Improved security for Mass Ingestion service ports" on page 14
September 24, 2026	Modified	"Enhanced Query-based CDC logic to release task units faster and reduce queuing time" on page 15
September 24, 2026	Added	"Additional targets for Application Ingestion and Replication" on page 12
September 24, 2026	Added	"New Data Ingestion and Replication CLI options to override connection and runtime environment IDs" on page 10
September 10, 2026	Initial publication	

Important notices and upcoming changes

The October 2026 release includes the following important notices and upcoming changes.

Preview lifted

Effective in the October 2026 release, the following functionalities are generally available and production-ready. Previously, these functionalities were only available in preview.

Auto-tuning for database ingestion and replication tasks

You can enable auto-tuning on database ingestion and replication initial load tasks and the initial load phase of combined load tasks. Auto-tuning automatically sets optimal values for pipeline parameters, including source partition count, helper thread counts, distributor counts, and the number of rows per output file. To enable auto-tuning for a task, select the **Auto-Tune** option on the last page of the task configuration wizard.

SAP HANA Log-based CDC processing of encrypted logs

Database ingestion and replication jobs that use Log-based CDC can now capture changes from encrypted SAP HANA online redo logs and archive logs. When you configure the SAP HANA Database Ingestion connection, provide the Key Backup Password and Key Backup Data Path to allow the connector to decrypt AES-256 CBC encrypted online redo and archive log files. The connector checks each log file to detect whether encryption is active. This action ensures jobs continue uninterrupted even if you turn encryption on or off on the source database while a job is running.

Apache NiFi 2.5.0 runtime engine for Streaming Ingestion and Replication

Streaming Ingestion and Replication uses the Apache NiFi version 2.5.0 runtime engine. No user action is required. The upgrade includes key updates to custom processors and AMQP NAR components. Streaming Ingestion and Replication is now enhanced to ensure compatibility with JDK 21, as Apache NiFi 2.5 requires a minimum of JDK 21 version to run effectively.

Reinstated support

Effective in the October 2026 release, historical statistics for job monitoring in Streaming Ingestion and Replication is reinstated and is available.

You can view the historical performance statistics for jobs on the **Performance** tab in My Jobs and Operational Insights. The performance charts track data movement trends, such as the total number of messages and the data volume transferred between the source and the target. You can use the charts to compare a job's current performance against its history, spot unusual changes, and troubleshoot issues.

Deprecated support

Effective in the October 2026 release, OAuth 2.0 Username-Password Flow authentication in Salesforce Mass Ingestion Connector is deprecated because Salesforce is retiring the SOAP API login() call in API versions 31.0 through 64.0 to align with enhanced cybersecurity standards. For more information, see Platform SOAP API login() Retirement.

Deprecation Notice: Deprecated functionality is supported, but Informatica intends to drop support in a future release. Informatica requests that you transition to different functionality before the functionality is dropped.

Informatica intends to drop support for OAuth 2.0 Username-Password Flow on May 1, 2027. In the interim, transition to OAuth 2.0 JWT Bearer Flow authentication.

Additionally, you can no longer create new connected apps in Salesforce. Existing connected apps continue to work, but you cannot create new ones. To create a new Salesforce Mass Ingestion connection, configure an External Client App in Salesforce. For more information, see Create an External Client App. To migrate an existing connected app for your Salesforce Mass Ingestion connection to an external client app, see Connected App to External Client App Migration.

New features and enhancements

The October 2026 release of Data Ingestion and Replication includes the following new features and enhancements.

Common

The October 2026 release of Data Ingestion and Replication includes the following new features that are common to multiple types of ingestion and replication tasks.

Java Runtime Environment 17 requirement for Data Ingestion and Replication Command-Line Interface

The Data Ingestion and Replication Command-Line Interface (CLI) requires Java Runtime Environment (JRE) 17 or later version. If you run the CLI on an earlier JRE version, the CLI stops with a message indicating that JRE 17 or later is required.

REST API enhancements for creating tasks

The REST API request `POST /dbmi/public/api/v2/task/create` that creates an application ingestion and replication or database ingestion and replication task includes additional configuration options. The REST APIs now handle errors better and return clearer messages when a request fails.

The following list shows some of the new options you can configure through this API:

- Source properties such as change data capture method, query-based CDC columns and intervals, large object replication, source partitioning for parallel initial loads, and reader performance tuning.
- Endpoint-specific source properties for connectors such as Workday, Oracle Fusion Cloud, and Salesforce Bulk API 2.0.
- Target properties such as apply mode, Databricks unmanaged tables, Snowflake ingestion options including Superpipe, writer performance tuning, JSON output format, and audit columns.

You can also enable automatic performance tuning in your request. The task automatically tunes initial load jobs and the initial phase of combined load jobs, so you don't need to adjust performance settings manually.

For special processing needs, you can set custom properties at the direction of Informatica Global Customer Support or a technical staff member.

For the full list of attributes that you can specify in your request, see *Database Ingestion and Replication REST API Reference*.

Microsoft Fabric Lakehouse target

You can now use Microsoft Fabric Lakehouse as a target for application ingestion and replication tasks and database ingestion and replication tasks of all load types.

To connect to a Microsoft Fabric Lakehouse target, use the Microsoft Fabric Lakehouse connector.

Confirmation prompt when undeploying tasks

A confirmation dialog appears before you undeploy an application ingestion and replication, database ingestion and replication, or streaming ingestion and replication task.

You must confirm the action before the task is undeployed. This prompt prevents you from accidentally undeploying a running task when you click **Undeploy**.

Audit metadata columns for Microsoft Fabric OneLake targets

You can now configure audit metadata columns for application ingestion and replication tasks and database ingestion and replication tasks that have a Microsoft Fabric OneLake target.

In **Advanced Options** under **Advanced Target Properties**, you can enable columns that track change-data metadata, such as operation type, operation sequence, and last replicated timestamp. To identify the source table or source schema for each replicated row, you can also enable the source table name and source schema name columns using target custom properties.

Schema drift options enabled for Microsoft Fabric OneLake targets

Automatic schema drift detection and handling is now available for Microsoft Fabric OneLake targets in incremental load and combined initial and incremental load application ingestion and replication jobs and database ingestion and replication jobs.

Schema drift occurs when you add, modify, drop, or rename columns in a source table.When defining a task, you can configure the types of source schema changes to propagate and how to handle them on the target. For example, you can configure schema drift options to ignore the changes, replicate the changes, or stop the job or a subtask when a schema change occurs. If you choose to stop the job, you can use the **Resume With Options** command to resume the job with an override schema drift option.

IAM Role ARN property available for Amazon S3 targets

You can now use the **IAM Role ARN** property in the Amazon S3 V2 target connection properties for application ingestion and replication tasks and database ingestion and replication tasks. Previously, this property did not apply to Data Ingestion and Replication tasks.

The **IAM Role ARN** property is used in the AssumeRole with EC2 role and AssumeRole with IAM user authentication types. In both types, the Secure Agent assumes the specified role through AWS Security Token Service and obtains temporary credentials automatically, so you don't need to store a long-lived access key or secret key in the connection. Cross-account Amazon S3 access is now possible without modifying the Secure Agent's own instance role.

Amazon Redshift JDBC driver upgrade

The Amazon Redshift JDBC driver has been upgraded from version 2.1.0.30 to 2.2.7 to address a security vulnerability that could allow unauthorized code execution under certain conditions. No configuration changes are required. Existing Amazon Redshift connections continue to work as before the upgrade.

Microsoft SQL Server JDBC driver upgrade

Data Ingestion and Replication now includes Microsoft SQL Server JDBC driver version 13.2.1 when you connect to Microsoft SQL Server sources and targets. This version replaces driver version 12.10.2, which reached end of life.

Secure Agent on Red Hat Enterprise Linux (RHEL) 10

You can now install the Secure Agent on RHEL 10 to handle your Data Ingestion and Replication tasks.

New Data Ingestion and Replication CLI options to override connection and runtime environment IDs

In the Data Ingestion and Replication Command-Line Interface (CLI) `task create` and `task replace` commands, you can now specify the source connection ID, target connection ID, and runtime environment ID as command-line arguments to override the values in the configuration file. This lets you reuse the same configuration file across organizations, such as when you promote a task configuration from a development environment to production. In the `task replace` command, you can also specify an argument to retain the connection and runtime environment IDs already stored in the repository.

The syntax for `task create` is:

```
task create --pathToConfig path_to_configuration_file [--taskName task_name][--
taskLocation task_location][--taskType task_type][--deploy][--sourceConnectionId
source_connection_id][--targetConnectionId target_connection_id][--runtimeEnvId
runtime_environment_id]
```

The syntax for `task replace` is:

```
task replace --pathToConfig path_to_configuration_file [--taskName task_name][--
taskLocation task_location][--deploy][--noprompt][--sourceConnectionId
source_connection_id][--targetConnectionId target_connection_id][--runtimeEnvId
runtime_environment_id][--retainExistingIds]
```

Job management from the task configuration page

You can now run, stop, abort, resume, redeploy, and undeploy application ingestion and replication jobs and database ingestion and replication jobs directly from the task configuration page. The Actions menu in the upper-right corner of the page displays options based on the current job status. You no longer need to navigate to the **My Jobs** page, Monitor, or Operational Insights to manage your jobs.

Serverless runtime environments extended to Cloud Data Lake targets

You can run application ingestion and replication and database ingestion and replication jobs on a serverless runtime environment hosted on Microsoft Azure with Cloud Data Lake targets such as Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage, Microsoft Fabric OneLake, and Oracle Cloud Object Storage targets.

Database Ingestion and Replication can now run incremental load and combined initial and incremental load jobs that use the Query-based CDC method to capture changes from an Oracle or SQL Server source and load to Cloud Data Lake targets using a serverless runtime environment.

Automatic agent switchover extended to Cloud Data Lake targets

Automatic switchover to another Secure Agent in a Secure Agent group is enabled by default for new application ingestion and replication jobs and database ingestion and replication jobs with Cloud Data Lake targets, such as Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage, Microsoft Fabric OneLake, and Oracle Cloud Object Storage.

Automatic agent switchover now also applies to database ingestion and replication jobs that use the Query-based CDC method to capture changes from a Db2 for LUW, Oracle, or SQL Server source to a Cloud Data Lake target.

Multiple object distributors for Cloud Data Lake targets

You can configure multiple writer object distributors to run in parallel during initial load jobs for Cloud Data Lake targets, including Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage Gen2, Microsoft Fabric OneLake, and Oracle Cloud Object Storage. Multiple distributors can reduce the initial load job run time on high-volume datasets. This feature is available for application ingestion and replication tasks and database ingestion and replication tasks. For combined initial and incremental load tasks, multiple object distributors apply to the initial load phase in synchronous mode only.

To enable the feature and enter the number of parallel distributors, use the **Writer Unload Multiple Distributors** and **Writer Distributor Count** target custom properties to run multiple writer instances simultaneously to help the writer throughput keep pace with parallel source reads. Alternatively, for database ingestion and replication tasks, enable the **Auto-Tune** option to calculate the optimal values automatically.

Billing project override for Google BigQuery targets

For application ingestion and replication tasks and database ingestion and replication tasks that use Google BigQuery targets, you can specify a separate billing project for query execution costs. To use this feature, configure the writerGbqBillingProjectId custom property with the billing project ID under **Advanced Target Properties** on the **Target Details** page of the task wizard. Google Cloud then bills query costs to that project while you continue to access data from the dataset in the storage project defined in the connection.

Application Ingestion and Replication

The October 2026 release of Application Ingestion and Replication includes the following new features and enhancements.

Additional targets for Application Ingestion and Replication

You can now use Microsoft Fabric Data Warehouse and Salesforce Data 360 targets with Oracle Fusion Cloud, Salesforce, SAP OData, and SAP ABAP sources, including SAP ECC and SAP S4/HANA with SAP Mass Ingestion Connector, for all load types in application ingestion and replication tasks.

All load types added for Cloud Data Lake targets with SAP OData

You can configure an application ingestion and replication incremental or combined load task from SAP OData to cloud data lake targets such as Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage Gen2, Microsoft Azure Synapse Analytics, Microsoft Fabric OneLake, and Oracle Cloud Infrastructure Object Storage. Previously, you could use these targets with SAP OData source for initial loads only.

Readable display values from Microsoft Dynamics 365 sources

When you create a task that uses Microsoft Dynamics 365 as the source, the task now writes the readable value for reference and coded fields to the target, next to the original column.

For OWNERTYPE and Lookup fields, the task adds a column named `<ColumnName>_display_name` that holds the display value, and the original column retains the GUID. For Choice, Yes-No, State, and Status fields, the task writes the label column `<field>name`, such as statuscodename, as a string column next to the numeric code. The target now holds both the reference and its resolved value. Existing tasks keep their current target columns. To add the display value columns, you need to redeploy your tasks.

Database Ingestion and Replication

The October 2026 release of Database Ingestion and Replication includes the following new features and enhancements.

Auto-tuning enhancements

Auto-tuning is now available for the following cloud data lake targets: Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage Gen2, Microsoft Fabric OneLake, and Oracle Cloud Object Storage. This release also introduces the following performance and resource optimization enhancements:

- Adaptive ring buffer sizing. Auto-tuning now sizes each table's internal pipeline ring buffer by selecting a slot count sized to each table's profile, reducing the risk of failures on wide tables and avoiding unnecessary memory use on narrow ones. The unloadEventBatchSize value is automatically tuned together with ringBufferSize so that peak in-flight payload stays near the configured size target.
- Dynamic writer flush size. Auto-tuning now adjusts the maximum number of rows that the writer batches before flushing to the target based on observed target latency to improve throughput.
- Evaluate-and-stop mode. You can run auto-tuning in evaluate-and-stop mode, to preview the auto-tuning recommendations without moving any data. In this mode, auto-tuning collects metrics and runs a full tuning evaluation, logs the recommended values for all components, and then stops the task gracefully before the unload and load execution begins. The task exits with a SUCCESS status. Use this mode to validate the auto-tuning decisions against a real task's metrics, compare recommended values across different source and target configurations, and identify regressions in the tuning algorithm without incurring the cost of a full data-movement run. This mode applies to initial load jobs only.

Table-level error isolation in incremental load tasks

You can enable table-level error isolation for database ingestion and replication tasks that replicate data from Oracle to Snowflake and use the log-based CDC method. When you enable error isolation and a table

encounters an irrecoverable error, that table stops while all other tables in the task continue replicating. The task transitions to a Running with Warning state and raises an alert that identifies the stopped table and the reason for the error.

To enable this feature, contact Informatica Global Customer Support.

Granular privileges for deploy, redeploy, and undeploy actions

As an administrator, you now have granular control over the deploy, redeploy, and undeploy actions on database ingestion and replication tasks through three separate feature privileges. You can restrict any of these actions for a role.

To let a user perform these actions, assign the Deploy Ingestion Task, Redeploy Ingestion Task, and Undeploy Ingestion Task feature privileges to the role.

For example, you can let a user run ingestion tasks but withhold the undeploy privilege. Previously, a Run permission combined the deploy, redeploy, start, stop, abort, and undeploy actions.

Note: Effective in the October 2026 release, an organization-level flag controls these feature privileges. To enable this functionality for your organization, create a request with Informatica Global Customer Support. If you select and save these feature privileges, they won't take effect until Informatica Global Customer Support activates the functionality for your organization.

When enabled, the Admin, Deployer, and Designer system-defined roles include all three feature privileges by default, but custom roles don't. When a user in a custom role lacks these privileges, the Deploy, Redploy, and Undeploy options appear disabled in the ingestion and replication task setup flows, and don't appear in the Monitor and My Jobs menus or in Operational Insights. To grant these actions, an administrator must add the feature privileges to the custom role.

Oracle Autonomous Database target

You can now use Oracle Autonomous Database as a target in database ingestion and replication initial load tasks. To connect to a Oracle Autonomous Database target on Oracle Cloud Infrastructure, use the Oracle Autonomous Database connector.

Access Management for Oracle sources

You can now apply Cloud Data Access Management (CDAM) governance rules to Oracle source tables in database ingestion and replication jobs. In the **Transform** page of the task configuration, select one or more tables and add the Access Management transformation to assign de-identification and filtering policies defined in Cloud Data Governance and Catalog.

Connections available for Access Management in database ingestion and replication tasks

When you add access management to a database ingestion and replication task, the connection list on the Access Management page includes all the connections that match the source vendor type, including connections that are not Database Ingestion connections. Previously, the Access Management page listed only Database Ingestion connections, which meant that source vendors that required a non-ingestion connection, such as Oracle, had no available connection for access management configuration.

Truncate event handling for Oracle sources with Snowflake targets

Database ingestion and replication jobs that have an Oracle source can now detect TRUNCATE TABLE operations on an Oracle source table during log-based CDC replications to a Snowflake target. The incremental load task truncates the target table in standard apply mode or inserts a truncate record with an operation type of "T" in audit apply mode. You can configure this behavior with custom properties in the task configuration. You can also enable an alert that notifies you when the task detects a truncate event at the source.

Encrypted Oracle source connection passwords

Data Ingestion and Replication extends the CLI password encryption to support Oracle-specific connection credentials. You can now provide encrypted values for the following fields:

- Oracle database password
- Reader ASM Password
- TDE Wallet Password

The encryptText command introduces a new --method parameter with a selfcontained option. Unlike the existing keystore method, selfcontained embeds the encryption key directly in the encrypted output, and does no require a key-store.txt file on the machine that runs the CLI commands.

Soft Deletes for Microsoft Fabric Data Warehouse and Open Table targets

Database ingestion and replication incremental load jobs and combined initial and incremental load jobs can now process hard delete operations on the source as soft deletes on Microsoft Fabric Data Warehouse and Open Table targets. Database Ingestion and Replication marks the soft-deleted records with a "D" in the INFA_OPERATION_TYPE column on the target without actually deleting the records.

Schema drift options enabled for Microsoft Fabric Data Warehouse, Open Table, and Salesforce Data 360 targets

When you configure database ingestion and replication tasks that have Microsoft Fabric Data Warehouse, Open Table, or Salesforce Data 360 targets, you can now set schema drift options in the task configuration. Select REPLICATE, IGNORE, STOP TABLE, or STOP JOB options to control how the task responds to source column additions and modifications. The task does not replicate dropped or renamed columns to the target. For schema drift events that involve dropped or renamed columns, you can select IGNORE, STOP TABLE, or STOP JOB options.

File Ingestion and Replication

The October 2026 release of File Ingestion and Replication includes the following new features.

GCM cipher and RSA SHA-2 support for the Advanced SFTP V2 connector

You can use the Advanced SFTP V2 connector to connect to SFTP servers that require the aes128-gcm@openssh.com and aes256-gcm@openssh.com GCM ciphers. You can also use the connector to authenticate to SFTP servers that require the RSA SHA-2 host key algorithms `rsa-sha2-512` or `rsa-sha2-256`. The connector also works with servers that still accept the older ssh-rsa algorithm, with no change to your connection configuration.

Improved security for Mass Ingestion service ports

The Mass Ingestion service uses internal service ports. These ports now bind to the local host (127.0.0.1), so processes running on the same machine as the Secure Agent can connect to them.

Changed behavior

The October 2026 release of Data Ingestion and Replication includes the following changed behavior.

Enhanced Query-based CDC logic to release task units faster and reduce queuing time

For database ingestion and replication combined initial and incremental load tasks that have a Db2 for LUW, Oracle, or SQL Server source and use the query-based CDC method, a new dynamic algorithm triggers a CDC cycle as soon as one or more tables complete their initial load and releases the task unit for those tables sooner. During this initial period, CDC cycles might run at shorter, dynamic intervals. After all tables reach a normal state, the task returns to the CDC interval that you configured.

Previously, the task didn't release the task unit for a table until a full CDC cycle completed. This behavior caused tables that had finished unloading to remain in backlog status for hours if you configured a long CDC interval with a low Max Task Units value.

New parameter for encryptText command in Data Ingestion and Replication Command-Line Interface

When you encrypt a text value with the `encryptText` command in the Data Ingestion and Replication Command-Line Interface, you must now specify an encryption method with the `--method` parameter: Available options are:

- `keystore`. Legacy value. The encryption key is stored in a key-store.txt file
- `selfcontained`. New value. The encryption key is embedded in the encrypted output, making the encrypted value portable without requiring a separate key file.

Previously, the `keystore` encryption method was the only available option and the `encryptText` command did not include a method parameter.

Number of Rows in Output File field location

For application ingestion and replication tasks and database ingestion and replication tasks, the **Number of Rows in Output File** field has moved from the **Advanced Options** section to the **Custom Properties** section on the **Let's Go!** page of the task configuration wizard. If you don't select a value, a default of 100000 rows is applied. For tasks that are already deployed, the field appears in the **Custom Properties** section and is pre-populated with the value saved at deploy time.

SAP HANA log-based CDC archive log discovery

For SAP HANA log-based CDC, Database Ingestion and Replication now reads the list of archive log files from the SAP HANA backup catalog instead of scanning the mount point directory. Because only catalog-registered files are processed, duplicate archive log files on disk are excluded, preventing incremental and combined load jobs from getting stuck in the BACKLOG stage.

The database user configured for the SAP HANA source connection must have read access to the M_CATALOG_BACKUP and M_CATALOG_BACKUP_FILES monitoring views.

Previously, archive log discovery scanned the mount point directory directly, which caused CDC jobs to stall when duplicate redo-log files were generated during a BACKINT backup server outage.

Updated names for Google Cloud Storage V2 connection properties

The Google Cloud Storage V2 connection properties now use the names Service Account Private Key ID and Service Account Client ID. Previously, these connection properties had the names Private Key ID and Client ID. The name change doesn't affect existing connections or upgrades.

Path validation for local directories in Google Cloud Storage V2 tasks

File ingestion and replication tasks that move files between a local directory and Google Cloud Storage now validate the local directory. A directory that contains relative or shortcut segments such as /.. or ~, for example /root/FMI/f1/../f2/, fail with the following error message:

```
[GoogleStorageV2_AFTER_FILE_PICKUP_ACTION_DIR_NOT_VALID_7013]
The specified local directory is not valid.
```

Previously, tasks resolved these directories. Existing tasks that use such directories fail after the upgrade. To fix a task, replace the directory with an absolute directory that has no relative or shortcut segments.

CHAPTER 2

July 2026

This section provides information about new features, enhancements, and behavior changes in the July 2026 release of Data Ingestion and Replication.

Changelog July 2026

Read the changelog to learn about significant changes to this publication for this release.

The following table describes changes in this publication:

Publication Date	Type of Change	Section Title
July 10, 2026	Added	Support for Oracle 23ai and Oracle AI Database 26ai sourceson page 18
July 10, 2026	Added	"Preview initiated" on page 17
June 18, 2026	Added	Column selection for PostgreSQL sourceson page 18
June 11, 2026	Initial publication	

Important notices

The July 2026 release includes the following important notice:

Preview initiated

Effective in the July 2026 release, Apache NiFi version 2.5.0 for Streaming Ingestion and Replication is available for preview. You can now upgrade the Apache NiFi runtime used by Streaming Ingestion and Replication from version 1.14.0 to 2.5.0. The preview package includes key updates to custom processors and AMQP NAR components. Informatica recommends that you upgrade and test your existing streaming ingestion and replication tasks only in a non-production environment to verify compatibility before the general availability rollout, planned for the October 2026 release. To access the preview package and review the migration details, contact Informatica Global Customer Support.

Preview functionality is supported for evaluation purposes but is unwarranted and is not supported in production environments or any environment that you plan to push to production. Informatica intends to

include the preview functionality in an upcoming release for production use, but might choose not to in accordance with changing market or technical circumstances. For more information, contact Informatica Global Customer Support.

Dropped support

Effective in the July 2026 release, Informatica drops support for the following connectors in Application Ingestion and Replication:

- Adobe Analytics Mass Ingestion Connector
- Zendesk Mass Ingestion Connector

New features and enhancements

The July 2026 release of Data Ingestion and Replication includes the following new features and enhancements.

Database Ingestion and Replication

The July 2026 release of Database Ingestion and Replication includes the following new features and enhancements:

Column selection for PostgreSQL sources

If you use rules to select PostgreSQL source tables when you define a database ingestion and replication task, you can individually select or deselect the columns in each of the selected tables from which to replicate data.

Support for Oracle 23ai and Oracle AI Database 26ai sources

You can now use Oracle 23ai and Oracle AI Database 26ai as sources in database ingestion and replication tasks.

File Ingestion and Replication

The July 2026 release of File Ingestion and Replication includes the following new feature:

Support for Oracle Cloud Object Storage sources and targets

You can use Oracle Cloud Object Storage as a source and target in file ingestion and replication tasks. You can also use Oracle Cloud Object Storage as a source in file listeners. For more information, see the *Connectors and Connections* and *File Ingestion and Replication* documentation.

Changed behavior

The July 2026 release of Data Ingestion and Replication includes the following changed behavior.

Task configuration wizard enhancement

After you select a destination connection or source connection for an application or database ingestion and replication task in the task configuration wizard, the connection name appears after the page name in the row of navigation buttons at the top for easy reference, for example, Destination: <connection_name> or Source: <connection_name>.

Note: If you add a new connection in the wizard, the connection name doesn't appear at the top until you click the newly added box for it on the Destination or Source page.

Nulls for unmodified Oracle columns can be skipped during change data replication to Databricks

Database ingestion and replication incremental load jobs that have an Oracle source and Databricks target and that have Oracle supplemental logging enabled on primary key columns only can now be configured to skip replicating null values for the before images and after images of any unmodified source columns to the target. This new behavior is available for **Standard** and **Soft Deletes** apply modes only. To enable this behavior, set the target custom property writeSkipUpdatingUnmodifiedColumns to true. By default, the property value is false to preserve pre-existing behavior. However, with the pre-existing behavior, null values for the unmodified columns might not be properly handled during Update processing, leading to the source and target becoming out of sync.

Note: If you enable skipping nulls for unmodified columns and a primary key column in a source table changes during CDC processing, the table is placed into an error state to prevent the job from failing on restart. To restart the job and recover the table, use the **Resume with Options > Resync** command.

CHAPTER 3

April 2026

This section provides information about new features, enhancements, and behavior changes in the April 2026 release of Data Ingestion and Replication.

Changelog April 2026

Read the changelog to learn about significant changes to this publication for this release.

The following table describes changes in this publication:

Publication Date	Type of Change	Section Title
April 8, 2026	Added	"Additional target types for Oracle Fusion Cloud sources" on page 25 "Deprecated support for session-based authentication in Data Ingestion and Replication" on page 22
April 2, 2026	Modified	"Access Management policies for filtering and protecting data" on page 28 - Noted the Access Management policies are for SQL Server sources only.
April 2, 2026	Added	"Replication slot data persistence moved to PostgreSQL database table " on page 29 "Change to default data-type mappings for SQL Server sources and targets that use Parquet format" on page 30
March 26, 2026	Modified	"New target types for database ingestion and replication tasks" on page 26 - Noted the new target types are for Database Ingestion and Replication only.

Publication Date	Type of Change	Section Title
March 26, 2026	Added	"Providing JDBC connection details in the tnsnames.ora file for Oracle unload processing " on page 27 "Support for PostgreSQL 17.x" on page 24 "New Data Ingestion and Replication CLI options in the task replace command" on page 24
March 19, 2026	Modified	MySQL source columns with LOB data typeson page 26 - Noted limitation for JSON columns
March 19, 2026	Added	Preview initiated for auto-tuningon page 21 Dropped support for email and password authentication in Zendesk Mass Ingestion Connectoron page 22
March 11, 2026	Initial publication	

Important notices

The April 2026 release includes the following important notice:

Preview initiated for auto-tuning

Effective in the April 2026 release, auto-tuning is available for preview.

Auto-tuning dynamically adjusts key parameters for database ingestion and replication initial load jobs and the unload phase of combined load jobs to optimize performance and resource use. Auto-tuning is based on performance and system metrics automatically collected from your environment and on application-specific metrics such as JVM heap allocation and task capacity. To enable auto-tuning for a task, select the **Auto-Tune** option on the last page of the task configuration wizard. The following source types are supported with any target type: Db2 for i, Db2 for LUW, Db2 for z/OS, MySQL, Oracle, PostgreSQL, SAP HANA, SQL Server, and Teradata.

Preview functionality is supported for evaluation purposes but is unwarranted and is not supported in production environments or any environment that you plan to push to production. Informatica intends to include the preview functionality in an upcoming release for production use, but might choose not to in accordance with changing market or technical circumstances. Note that if you are working on a preview POD, all data is excluded from SOC 2 compliance coverage. For more information, contact Informatica Global Customer Support.

Dropped support for the older task configuration wizard

Effective in the April 2026 release, Data Ingestion and Replication drops support for the older task configuration wizard for application ingestion and replication tasks and database ingestion and replication tasks. The newer task configuration wizard, which was first introduced in October 2024, is now the only supported configuration wizard for these task types. The newer wizard will be automatically enabled for any users who haven't previously migrated to it.

Dropped support for Google Analytics Mass Ingestion Connector

Effective in the April 2026 release, Informatica drops support for the Google Analytics Mass Ingestion Connector following Google Analytics' deprecation of the UA4 interface. While the connector might still appear in Application Ingestion and Replication tasks, it no longer functions. Informatica will remove the connector and its documentation in a future release.

Dropped support for email and password authentication in Zendesk Mass Ingestion Connector

Effective in the April 2026 release, Informatica drops support for email and password authentication to Zendesk following Zendesk's deprecation of email and password access for API calls.

Deprecated support for session-based authentication in Data Ingestion and Replication

Effective in the October 2025 release, session-based authentication is deprecated. Informatica intends to drop support for session-based authentication in the October 2026 release.

Deprecated functionality is supported, but Informatica intends to drop support in a future release. Informatica requests that you transition to different functionality before the functionality is dropped.

Informatica plans to replace session IDs used for user authentication with JSON Web Tokens (JWT). The organization administrator can change the authentication type used for REST API requests on the **Settings** page in Administrator. If the administrator enables the JSON Web Token (JWT) authentication type, browser-based user interface communications seamlessly use JWT behind the scenes.

For REST API calls, session IDs must still be included in request headers. However, with JWT authentication, the session value carries a JWT with a fixed expiry time configured in Administrator. As a result, REST API scripts must be updated to handle automatic re-authentication before or after the token expires. For more information, see the JWT Support Knowledge article.

Session-based authentication is the default authentication type. Sometime after the July 2026 release, JWT will be the default authentication method. By October 2026, session-based authentication will no longer be available.

New features and enhancements

The April 2026 release of Data Ingestion and Replication includes the following new features and enhancements.

Watch the What's New video to learn about new features and enhancements in the April 2026 release.

Common

The April 2026 release of Data Ingestion and Replication includes the following new features that are common to multiple types of ingestion and replication tasks.

REST APIs for creating and monitoring tasks

Publicly accessible REST APIs are now available for application ingestion and replication, database ingestion and replication, file ingestion and replication, and streaming ingestion and replication tasks.

This release introduces the following resources to help you manage your application ingestion and replication and database ingestion and replication tasks:

- Create a task. `POST /dbmi/public/api/v2/task/create`

- Deploy a task. `POST /dbmi/public/api/v2/task/deploy/{taskId}`
- Run a task. `POST /dbmi/public/api/v2/job/start`
- Stop a job. `POST /dbmi/public/api/v2/job/stop`
- Resume a job. `POST /dbmi/public/api/v2/job/resume`
- Undeploy a job. `POST /dbmi/public/api/v2/job/undeploy`
- Get job status. `GET /dbmi/public/api/v2/job/status`
- Get job metrics. `POST /dbmi/public/api/v2/job/metrics`
- Get task details. `GET /dbmi/public/api/v2/task/details?taskId={taskId}&projectId={projectId}&folderId={folderId}`

For more information about each endpoint including request and response formats and usage examples, see the *REST API Reference* documentation.

CLAIRE Copilot enhancement for creating tasks

When you create an application or database ingestion and replication task in CLAIRE Copilot for Data Integration, you can now include source table or object selection critieria in the first prompt or enter the criteria later during the conversation. You can enter specific table or object names, the initial or ending part of the names, or name masks that include the * and ? wildcards, for example, `include tables with names starting with ca*, exclude tables with names starting with ca5*`. Enter include and exclude criteria in the order they should be processed. If you don't want to define a subet of tables, you must enter "select all" to select all tables. Once entered, you can't change your selection criteria later in the Copilot conversation. The summary at the end of the conversation will include the table selection criteria you entered. When you open the new task in the task configuration wizard, the criteria will appear as include and exclude rules. You can make changes there if needed.

For more information, see the *CLAIRE Copilot for Data Integration* documentation.

Open mirroring available for Microsoft Fabric OneLake targets

You can now enable Open Mirroring for Microsoft Fabric OneLake targets in application and database ingestion and replication tasks for all load types. Mirroring creates a synchronized copy of your data in Parquet format within the mirrored database in Fabric OneLake, ensuring your data stays current and readily available for timely analysis.

Last Replicated Time metadata column for Databricks and Amazon Redshift targets

When you configure application and database ingestion and replication tasks that have a Databricks or Amazon Redshft target, you can now select the new **Add Last Replicated Time** advanced option to add the INFA_LAST_REPLICATED metadata column in the target tables. This column records the timestamp at which a record was inserted or last updated in the target table.

For more information, see the *Application Ingestion and Replication* and *Database Ingestion and Replication* documentation.

Audit and Soft Deletes apply modes for Amazon Redshift targets

You can now use the **Audit** and **Soft Deletes** apply modes in database ingestion and replication incremental load and combined initial and incremental load tasks that have an Amazon Redshift target and any source type.

The **Audit** and **Soft Deletes** apply modes are also supported for application ingestion and replication incremental load and combined initial and incremental load jobs that have an Amazon Redshift target and a SAP Mass Ingestion source.

Use the **Audit** apply mode to write a row for each DML operation on a source table to a generated audit table in the target. This feature is useful when you need an audit trail of changes to perform downstream processing on the data before writing it to the target or when you need to examine the metadata for changes.

Use the **Soft Deletes** mode to process DML delete operations on the source as soft deletes in the target. Data Ingestion and Replication marks the soft-deleted records with a "D" in the INFA_OPERATION_TYPE column in the target without actually deleting the records.

For more information, see "Configure an Amazon Redshift target" in the *Application Ingestion and Replication* and *Database Ingestion and Replication* documentation.

New Programmatic Access Token (PAT) authentication to access Snowflake

You can configure your Snowflake Data Cloud connection in application, database, and file ingestion and replication tasks to use PAT to authenticate to Snowflake.

For more information, see "Snowflake Data Cloud connection properties" in the *Connectors and Connections* documentation.

Serverless runtime environments

Data Ingestion and Replication extends support for serverless runtime environments hosted on Microsoft Azure. Now you can run the following types of jobs on a serverless runtime environment:

Application ingestion and replication jobs for all load types:

- SAP ECC sources on Oracle database using the SAP Mass Ingestion connector

Database ingestion and replication jobs for all load types:

- MongoDB sources
- PostgreSQL sources

You cannot use a serverless runtime environment if you are using a read-only replica for a PostgreSQL connection.

For more information about serverless runtime environments, see *Administrator* > Runtime Environments > Serverless Runtime Environments and the Data Ingestion and Replication *Connectors and Connections* documentation for your connector.

Support for SQL Server 2022

Database Ingestion and Replication and Application Ingestion and Replication add support for SQL Server 2022. You can use version 2022 of SQL Server for RDS or on-premises Developer Edition as a source or target in database ingestion and replication jobs and as a target in application ingestion and replication jobs. In the SQL Server connection properties, the **SQL Server Version** field does not currently include 2022. You can select any option in this field. It will be ignored.

Support for PostgreSQL 17.*x*

Database Ingestion and Replication and Application Ingestion and Replication add support for PostgreSQL 17.*x*. You can use PostgreSQL Standard Edition 17.*x* or RDS for PostgreSQL 17.*x* for sources or targets in database ingestion and replication tasks and for targets in application ingestion and replication tasks.

New Data Ingestion and Replication CLI options in the task replace command

In Data Ingestion and Replication Command-Line Interface (CLI) task replace command, you can now specify the task name and task location as command line arguments. The syntax is:

```
task replace --pathToConfig path_to_configuration_file [--taskName task_name][--
taskLocation task_location]
```

Application Ingestion and Replication

The April 2026 release of Application Ingestion and Replication includes the following new features and enhancements:

HTTPS available for SAP Mass Ingestion Connector

You can enable HTTPS for SAP Mass Ingestion Connector to securely connect to SAP for your application ingestion and replication initial load jobs including the initial load phase in combined load jobs.

For more information, see the *Connectors and Connections* and *Application Ingestion and Replication* documentation.

Automatic switchover to another Secure Agent in a Secure Agent group

Automatic switchover to another Secure Agent in a Secure Agent group is now enabled by default for new application ingestion and replication jobs without persistent storage enabled for source log records.

If an active Secure Agent on which a job is running goes down unexpectedly, the job can automatically switch over to another available Secure Agent in the group after a 15-minute heartbeat interval. However, if the database ingestion service is in a stopped or error state, you need to manually stop and then resume the job.

For more information about automatic switchover and the sources and targets that jobs can include, see the *Getting Started with Data Ingestion and Replication* and the *Application Ingestion and Replication* documentation.

Additional target types for Oracle Fusion Cloud sources

Application ingestion and replication tasks of any load type with an Oracle Fusion Cloud source can now load data to Microsoft SQL Server and Microsoft Azure SQL Database targets.

For more information, see the *Application Ingestion and Replication* documentation.

Database Ingestion and Replication

The April 2026 release of Database Ingestion and Replication includes the following new features and enhancements:

Extended support for automatic switchovers to another agent in a Secure Agent group

Automatic switchovers to another Secure Agent in a Secure Agent group is now enabled for database ingestion and replication jobs that use the Log-based CDC method and have a MongoDB, MySQL, PostgreSQL, or SAP HANA source. This feature extends prior automatic switchover support for Log-based CDC with Db2 for i, Db2 for z/OS, Oracle, and SQL Server sources.

Also, automatic switchovers are now supported for tasks that use the Query-based CDC method and have an Oracle, Db2 for LUW, or SQL Server source.

If the active Secure Agent on which the job is running goes down unexpectedly, the job can automatically switch over to another available agent in the group after the 15-minute heartbeat interval elapses. The following limitations apply:

- Jobs cannot have persistent storage enabled.
- Jobs that have Kafka targets must store checkpoint information in the Kafka header. For any jobs that existed before the July 2025 release, automatic switchovers can't occur because checkpoint information is stored in the checkpoint file in the Secure Agent.
- Jobs that use the Query-based CDC method must have an Amazon Redshift, Databricks, Google BigQuery, Microsoft Azure Synapse Analytics, Oracle, PostgreSQL, Snowflake, or SQL Server target.
- Jobs that have an SAP HANA source cannot use the trigger-based CDC method.

- If you use a read-only replica for a PostgreSQL connection, you must set the readerPostgreSQLShouldDisableLocalPersistent custom property to false on the **Task Details** source page of the task configuration wizard .

The automatic switchovers are enabled by default for new tasks created in the current release and later.

New target types for database ingestion and replication tasks

Database Ingestion and Replication adds support for the following targets for all load types:

- Microsoft Fabric Data Warehouse
- Open Table
- Salesforce Data 360

For more information, see the Data Ingestion and Replication *Connectors and Connections* publication and *Database Ingestion and Replication* publication.

Support for Db2 for LUW 12.1 sources

Database Ingestion and Replication adds support for Db2 for LUW 12.1 sources for all load types.

Db2 for LUW Log-based CDC sources with additional target types

You can now use Db2 for LUW 11.x or 12.1 sources for Log-based CDC with these additional target types: Amazon Redshift, Databricks, Kafka, Oracle, PostgreSQL, or Azure SQL Server. Previously, only Snowflake targets were supported.

For more information, see the *Database Ingestion and Replication* documentation.

Db2 for LUW Log-based CDC sources on AIX

Database ingestion and replication tasks that use the Log-based CDC method can now process Db2 for LUW sources on AIX systems. You can use any supported Db2 for LUW version, the incremental load or combined load type, and the same CDC features that are available for Db2 sources on Linux or Windows, such as CDC staging groups.

For more information, see the *Database Ingestion and Replication* documentation.

Improved statistics messages for Db2 for z/OS sources

Database Ingestion and Replication introduces enhanced statistics messages than include additional details for jobs with Db2 for z/OS sources to help in diagnosing latency issues. The messages are written to the job log.

MySQL driver upgrade

Database ingestion and replication jobs that have MySQL sources now require the Progress DataDirect MySQL JDBC driver version 5.1.4.000364. This driver upgrade is delivered with April 2026 release of the product.

MySQL source columns with LOB data types

Database ingestion and replication jobs that have a MySQL source can now replicate data from source columns that have one of the following LOB data types: BLOB, LONGBLOB, MEDIUMBLOB, TINYBLOB, JSON, TEXT, LONGTEXT, MEDIUMTEXT, or TINYTEXT. The jobs can use any load type and any supported target type.

Note: JSON columns are not fully supported. The DataDirect MySQL JDBC driver does not parse Unicode characters in JSON columns correctly,which might lead to illegible characters on the target.

When you create an ingestion and replication task, select the new **Include LOBs** advanced option on the Task Details -Source Details page of the configuration wizard.

For more information, see the *Database Ingestion and Replication* documentation.

BFILE access to Oracle TDE wallet files

Database Ingestion and Replication can use BFILE directory objects for remote access to a file-based Transparent Data Encryption (TDE) wallet.

If you configured a file-based TDE wallet, Database Ingestion and Replication first checks if the ewallet.p12 file exists in the TDE wallet directory on the local machine. If the file exists, processing continues. If the file does not exist there, Database Ingestion and Replication queries the database for a directory object that has a path matching the TDE wallet directory or tries to create a directory object for BFILE access. After an appropriate directory object exists, the wallet can be read remotely. You do not need to copy the TDE wallet to the local machine.

For more information, see the *Database Ingestion and Replication* documentation.

Providing JDBC connection details in the tnsnames.ora file for Oracle unload processing

For database ingestion and replication unload processing as well as Log-based CDC, you can optionally enter JDBC connection details for the Oracle source database in an Oracle tnsnames.ora file. If you need to make changes to the connection later, you can then update the details in the tnsnames.ora file only instead of in both the Oracle Database Ingestion connection properties and in the tnsnames.ora file. To enable this feature for initial load jobs and the unload phase of combined jobs, set the source custom property useTnsNamesInJDBCUrl to true for the task. This feature does not apply to jobs that use the Query-based CDC method or to Oracle targets.

Note: When you use the tnsnames.ora file only, you can't configure SSL encryption and Kerberos authentication.

For more information, see "Providing JDBC connection details in the tnsnames.ora file in the *Database Ingestion and Replication* publication.

SAP HANA Log-based CDC processing of encrypted logs

Database ingestion and replication jobs that use Log-based CDC can now capture changes from encrypted SAP HANA online redo logs and archive logs. Currently, only the AES-256 CBC (Cipher Block Chaining) encryption algorithm is supported.

Note: To use SAP HANA Log-based CDC, with or without encrypted logs, a flag needs to be set for your organization. If you do not have access to the SAP HANA Log-based CDC option, contact Informatica Global Customer Support.

First, ensure that your database administrator created an encryption root keys file that contains the keys needed to decrypt the source logs. This file must reside on the Secure Agent machine. In a Secure Agent group, the file must reside on every Secure Agent machine in the group. Also, check that the administrator created a backup password for accessing the key file.

To enable database ingestion and replication processing of encrypted logs, you must set both the **Key Backup Password** and **Key Backup Data Path** properties in the SAP HANA Database Ingestion connection properties. If you do not set one of these properties, deployment of any job using the connection for SAP HANA Log-based CDC fails. If both properties are left empty, Database Ingestion and Replication assumes the logs are not encrypted.

For more information, see the Data Ingestion and Replication *Connectors and Connections* publication and *Database Ingestion and Replication* publication.

Row-level filtering for SAP HANA sources

You can now use row-level filtering for database ingestion and replication tasks that have a SAP HANA source. Configure row-level filtering rules to filter out data rows from selected source tables based on column conditions before the data is applied to the target. On the **Transform Data** page of the task wizard, you can create **Basic** or **Advanced** filtering rules for source tables.

For more information, see the *Database Ingestion and Replication* documentation.

Access Management policies for filtering and protecting data

Note: The Access Management feature availability is controlled by means of an organization-level feature flag. If this functionality is not available for your organization but you want to use it, create a request for Informatica Global Customer Support.

For database ingestion and replication jobs that have a SQL Server source, you can now use Access Management transformations to protect sensitive data by applying filtering and data protections before your data reaches the target.

The Access Management transformation applies data access policies that you create on the Data Access Management page in Data Governance and Catalog according to the data requirements of your organization.

You can apply data access policies for a selected set of tables on the **Transform** page of the task configuration wizard.

You can enforce the following types of data access policies:

- Data filter policies
- Data de-identification policies

Data filter policies control access to specific rows within data assets by filtering the data which users can interact with.

Data de-identification policies remove or mask sensitive information from data elements based on their classifications, applying protections to make the data anonymous or less identifiable for specific use cases.

File Ingestion and Replication

The April 2026 release of File Ingestion and Replication includes the following new feature:

Checksum algorithms for Amazon S3

You can select a checksum algorithm when loading files to Amazon S3. The checksum is stored with each file to efficiently detect changes without the need to recompute checksums, improving performance during ingestion and replication. Available algorithms include CRC64NVME, CRC32, CRC32C, SHA1, and SHA256, giving you the flexibilty to pick the best option for your data integrity and performance requirements.

Streaming Ingestion and Replication

The April 2026 release of Streaming Ingestion and Replication includes the following new feature:

CreateEntity API

You can programmatically create streaming ingestion and replication tasks using a new REST API endpoint, `CreateEntity`.

Use a POST request to the following URI:

```
/sisvc/restapi/v1/CreateEntity/Documents
```

For more information, see the *REST API Reference* documentation.

Changed behavior

The April 2026 release of Data Ingestion and Replication includes the following changed behavior.

Empty strings no longer replicated as NULL to Databricks targets

For application ingestion and replication tasks and database ingestion and replication tasks that have a Databricks target, empty strings ("") in a source string column now remain as empty strings in the target. Previously, any empty string in the source data was incorrectly replicated as NULL to the Databricks target.

With the new logic, a new default null marker is used. If a source string column contains NULL data, the data is replicated as `infa_null` in the Databricks staging file in CSV format. Data Ingestion and Replication interprets any occurrence of `infa_null` as an actual NULL in the target table.

To enable this behavior for new tasks created in the current release and later, set the writerEmptyStringsAsNulls custom property to false on the **Task Details - Target Details** page of the task configuration wizard.

For existing tasks deployed in a version earlier than April 2026, set both the writerEmptyStringsAsNulls and disableWriterDatabricksEmptyHandling custom properties to false on the **Task Details - Target Details** page of the task configuration wizard. Then redeploy the task.

SAP HANA Log-based CDC support for update SPS08

Database ingestion and replication jobs that use Log-based CDC can now process SAP HANA 2.*x* SPS08 sources. To use SPS08, the database ingestion and replication user must have read access to the SYS.M_SYSTEM_OVERVIEW system monitoring view. This view is required to fetch system overview information, such as the SAP HANA database version, that's needed to detect some changes in the structure of the binary files that SPS08 introduced. Without this information, jobs might encounter file-parsing failures or data discrepancies due to the incorrect parsing of data.

Improved processing of Db2 for z/OS compensation records

Previously, if database ingestion and replication incremental load or combined load jobs with a Db2 for z/OS source encountered unneeded Db2 compensation records in the log, a slowdown in CDC processing might occur. Now, CDC handling of compensation records is more efficient. Instead of removing the compensation log records at the end of the UOW, Database Ingestion and Replication stores the compensation records as key:value pairs in a hash map, which allows these records to be ignored when the log data is released to the log parser.

Improved handling of unload tasks across upgrade boundaries

If any unload tasks in database or application ingestion and replication jobs that use the initial load type are active when an IDMC upgrade occurs, the Database Ingestion agent service now waits for the running unload tasks to complete, up to a maximum of 24 hours, before it shuts down. This behavior helps ensure that the unload tasks have no unexpected shutdown during the upgrade and don’t need to re-read data after the upgrade. Previously, the Database Ingestion agent service stopped the unload tasks during the upgrade and restarted them from the beginning after the upgrade completed.

Replication slot data persistence moved to PostgreSQL database table

For database ingestion and replication incremental load and combined initial and incremental load jobs that have a PostgreSQL source, the persistence layer has been changed to use a PostgreSQL database table in the public schema for storing replication slot data immediately. Database Ingestion and Replication will remove data rows from the persistence table over time as records are written to the target.

For more information, see " PostgreSQL sources" in the *Database Ingestion and Replication* documentation.

Change to default data-type mappings for SQL Server sources and targets that use Parquet format

For database ingestion and replication tasks that have a SQL Server source and a target that uses the Parquet output format, the default data-type mappings now use the Avro logical data type of date, decimal, or timestamp as the target type by default, instead of using string. This change applies to mappings with source date, datetime, datetime2, decimal, money, and numeric data types. If you want to revert to the prior behavior of using string, contact Informatica Global Customer Support.

CHAPTER 4

November 2025

This section provides information about new features, enhancements, and behavior changes in the November 2025 release of Data Ingestion and Replication.

Changelog November 2025

Read the changelog to learn about significant changes to this publication for this release.

The following table describes changes in this publication:

Publication Date	Type of Change	Section Title
November 4, 2025	Initial publication	

Important notices

The November 2025 release includes the following important notice:

Upcoming wizard change

The new task configuration wizard for application ingestion and database tasks and database ingestion and replication tasks will begin rolling out to all remaining organizations. If you're still using the older wizard, you'll need to migrate to the new wizard before the April 2026 release. At that time, the new wizard will become the sole wizard for configuring application and database ingestion and replication tasks. If you haven't tried it yet, you can ask Informatica Global Customer Support or your Customer Success Manager to activate it for your organization.

Deprecation initiated

Effective in the November 2025 release, the following connectors used in Streaming Ingestion and Replicationare deprecated and have been moved to maintenance mode. Informatica intends to drop support in a future release. Your streaming ingestion and replication tasks that use these connectors will continue to function until Informatica officially drops support.

- Amazon Kinesis Firehose Connector
- Amazon Kinesis Streams Connector
- MQTT Connector

- OPC UA Connector

New features and enhancements

The November 2025 release of Data Ingestion and Replication includes the following new features and enhancements.

Common

The November 2025 release of Data Ingestion and Replication includes the following new features that are common to multiple types of ingestion and replication tasks.

Include source table selection criteria in CLAIRE Copilot prompts

In your initial prompt for creating an application or database ingestion and replication task in CLAIRE Copilot for Data Integration, you can now include criteria in plain language for selecting source tables or objects. You can specify table names or enter table-name masks that include the * or ? wildccard. Copilot translates the table selection criteria in the prompt into include and exclude rules, such as include sales* or exclude temp*. If you request a task summary or view the task in the wizard, the selection criteria is shown as include and exclude rules. For more information, see the *CLAIRE Copilot for Data Integration* documentation.

CHAPTER 5

October 2025

This section provides information about new features, enhancements, and behavior changes in the October 2025 release of Data Ingestion and Replication.

Important notices

The October 2025 release includes the following important notice:

Upcoming wizard change

Effective in the October 2025 release, the new task configuration wizard for application ingestion and database tasks and database ingestion and replication tasks is available to all organizations. Also, the Application Ingestion and Replication and Database Ingestion and Replication documentation now covers using only the new wizard to configure tasks. The July 2025 documentation for the pre-existing wizard is still available as a How-To article on the Documentation Portal.

If you're still using the older wizard, you'll need to migrate to the new wizard before the April 2026 release. At that time, the new wizard will become the sole wizard for configuring application and database ingestion and replication tasks. If you haven't tried it yet, ask Informatica Global Customer Support or your Customer Success Manager to activate it for your organization.

Deprecation of the Simba JDBC driver for Databricks

Effective in the October 2025 release for Application Ingestion and Replication and Database Ingestion and Replication, the Simba JDBC driver for Databricks targets is deprecated. Databricks recommends that you transition to the Databricks JDBC driver. Informatica intends to drop support for the Simba JDBC driver in a future release.

Upgrade task for CLI users

Effective in the October 2025 release, the Database Ingestion and Replication Command-Line Interface (CLI) files are installed in the `<Cloud Secure Agent installation>\apps\dbmicli` directory. If you have been using the CLI to create or manage application ingestion and replication tasks or database ingestion and replication tasks, before the upgrade, back up your files in the old path `<Cloud Secure Agent installation>\apps\Database_Ingestion\<latest_version>\dbmicli`, including the key-store, dbmiclienv.properties, and YAML files. After the upgrade, copy the files to the new dbmicli directory path.

Deprecation of Oracle versions

Effective in the October 2025 release, Oracle versions 11*g* R2 and 12*c* R1 and R2 are deprecated for sources and targets in database ingestion and replication tasks. Informatica intends to drop support for these versions in a future release.

New features and enhancements

The October 2025 release of Data Ingestion and Replication includes the following new features and enhancements.

Watch the What's New video to learn about the new features and enhancements in the October 2025 release.

Common

The October 2025 release of Data Ingestion and Replication includes the following new features that are common to multiple types of ingestion and replication tasks.

CLAIRE Copilot summaries of ingestion and replication tasks

CLAIRE Copilot for Data Integration can now generate brief and detailed summaries of database ingestion and replication tasks and application ingestion and replication tasks to help you understand how they're configured. To summarize a task, the task must be saved and valid. To produce the summary of a task, enter "Summarize ths asset" or "Generate detailed summary." For more information, see the *CLAIRE Copilot for Data Integration* documentation.

Serverless runtime environments

Data Ingestion and Replication extends support for serverless runtime environments hosted on Microsoft Azure. You can run the following types of application ingestion and replication jobs and database ingestion and replication jobs on a serverless runtime environment:

Application ingestion and replication sources:

- For all load types:
  - Marketo
  - Microsoft Dynamics 365
  - NetSuite
  - Oracle Fusion Cloud
  - Salesforce
  - ServiceNow
  - Workday
  - Zendesk
- For initial loads:
  - Salesforce Marketing Cloud
  - SAP OData V2

These application ingestion and replication jobs can have any of the following target types: Amazon Redshift, Databricks, Google BigQuery, Microsoft Azure Synapse Analytics, Oracle, PostgreSQL, Snowflake, and SQL Server.

Database ingestion and replication sources:

- For all load types with any target type:
  - Db2 for i
  - Db2 for z/OS
  - MySQL

  - Oracle (log-based CDC method only)
  - SQL Server (CDC tables and log-based CDC methods only)
- For initial loads with any target type:
  - PostgreSQL
  - SAP HANA
- For incremental loads and combined initial and incremental loads with Amazon Redshift, Databricks, Google BigQuery, Microsoft Azure Synapse Analytics, Oracle, PostgreSQL, Snowflake, or SQL Server targets:
  - Oracle (query-based CDC method only)
  - SQL Server (query-based CDC method only)

Previously, you could use serverless runtime environments only for application ingestion and replication initial load jobs that had a Marketo, Microsoft Dynamics 365, Salesforce, ServiceNow, or Zendesk source, and database ingestion and replication jobs that had a MySQL, PostgreSQL (initial loads only), Oracle, or SQL Server source.

You still cannot use serverless runtime environments for file ingestion and replication jobs and streaming ingestion and replication jobs.

For more information about serverless runtime environments, see *Administrator* > Runtime Environments > Serverless Runtime Environments and the Data Ingestion and Replication *Connectors and Connections* documentation for your connector.

Performance-related custom properties added in the task configuration wizard

In the latest task configuration wizard, the **Custom Properties** section of the **Task Details** source and target pages now includes some commonly used custom properties that can help improve performance, for your convenience. The specific custom properties added vary by source or target type and load type. When you configure an application or database ingestion and replication task, you can easily set one or more of these properties. The **Custom** option is also available if you need to manually enter any custom property that a technical support representative has provided for your unique requirement. For more information, see the *Application Ingestion and Replication* and *Database Ingestion and Replication* documentation.

Inline help panel added in the latest task configuration wizard

The new task configuration wizard for application and database ingestion and replication tasks now includes an inline help panel on the right side of each page for ease of access to help information about the page. To display or hide the inline help, use the small arrow icons (<, >) on the right border.

Databricks unmanaged target tables supported for volume staging

Application ingestion and replication tasks and database ingestion and replication tasks that have Databricks targets can now use Databricks unmanaged tables for volume staging.

To stage data in volumes, in the Databricks connection properties, set **Staging Environment** to **Volume**.

To use unmanaged tables, select the **Create Unmanaged Tables** check box on the target step of the task details. For volume staging, you must also provide the complete path to a parent directory that exists in Amazon S3 or Microsoft Azure Data Lake Storage to hold the files that are generated for each target table when captured DML records are processed.

Deletion of volumes generated for volume staging in Databricks targets

For application ingestion and replication tasks and database ingestion and replication tasks that have a Databricks target and use volumes as the staging environment, if you do not provide a path to the files within a volume in the **Volume Path** field of the connection properties, a new volume is generated automatically. The volume is deleted when an initial load job is successfully completed. Only a volume that is empty is

automatically deleted. The volume is not deleted when the initial load phase of a combined load job is completed.

Improved rule validation for column selection rules

For sources in application or database ingestion and replication tasks where you can select or deselect columns from the source tables to replicate data, primary key columns remain selected by default and you cannot manually deselect them. If you add rules to exclude columns, you can now run a validation to check of the rules. The validation warns you if a rule excludes a primary key column.

Less caching of information for data lake targets

The amount of information that application or database ingestion and replication jobs cache during data replication to data lake targets has been reduced to improve performance. Data lake targets include Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage, Microsoft Azure Data Lake Storage Gen2, Microsoft Fabric OneLake, and Oracle Cloud Object Storage.

New CLI commands to override schema drift settings when resuming jobs

The Data Ingestion and Replication CLI now includes commands to override the schema drift settings configured in the user interface when resuming jobs that are in a Stopped, Aborted, or Failed state because of schema drift errors in incremental load and combined load application or database ingestion and replication jobs. These overrides apply only to tables in an Error state caused by the **Stop Table** or **Stop Job** schema drift options.

Use the following CLI commands to override the schema drift settings when resuming a job:

- `job resumeWithOptions ignore <job name>`. Does not replicate DDL changes from the source to the target.
- `job resumeWithOptions stopTable <job name>`. Stops processing the source object affected by the DDL change.
- `job resumeWithOptions replicate <job name>`. Applies the DDL changes to the target and resumes the job.

The resumed job uses the specified schema drift override to process the schema change that caused the job to stop. Afterward, the schema drift options originally configured when creating the task take effect again.

For more information, see the *Data Ingestion and Replication Command-Line Interface*.

New Scheduled filter in Operational Insights

In Operational Insights service, you can now use the Scheduled filter to filter and view scheduled application or database ingestion and replication initial load jobs. This filter is available on both the **Overview** and **All Jobs** pages to help you easily check the schedule details and job status of the jobs. For more information, see "Monitoring all ingestion and replication jobs" in *Operational Insights*.

Enhancements for specifying Data Directory and Schema Directory expressions for data lake targets

You now have the option to edit the expressions for the **Data Directory** and **Schema Directory** fields when configuring an application or database ingestion and replication task that has a data lake target such as Amazon S3, Google Cloud Storage, Microsoft Azure Data Lake Storage, Microsoft Azure Data Lake Storage Gen2, Microsoft Fabric OneLake, and Oracle Cloud Object Storage.

Instead of typing directory expressions manually and searching documentation for supported syntax, you can directly edit these fields to modify the directory paths. You can select and insert supported variables and functions to create directory and schema paths.

Path pattern elements such as folder path, timestamp, schema name, and table name are available, along with a list of variables and functions to help you build your Data Directory or Schema Directory expressions.

Validation checks ensure your expression is valid, helping prevent errors before you save the expression.

Apache Iceberg open table format available to write to Amazon S3

You can use application or database ingestion and replication tasks to replicate data to Amazon S3 cloud storage as Apache Iceberg tables. You can then access these tables directly from Amazon S3 using the AWS Glue Catalog.

Snowflake-managed Iceberg tables for tasks with Snowflake targets

For application ingestion and replication jobs and database ingestion and replication jobs with Snowflake targets, you can store data and metadata in permanent Snowflake-managed Iceberg tables located in an external cloud location instead of in Snowflake tables. The external location can be in Amazon S3, Google Cloud Storage, or Azure Storage.

The external storage of data in an open format has the following benefits:

- Ability to use query processing engines other than Snowflake to process the data
- Interoperability across different data processing and analytics tools
- Better performance
- Robust security and governance

Data Ingestion and Replication automatically creates the Snowflake-managed Iceberg tables when you deploy ingestion and replication tasks enabled for Iceberg table use.

Snowflake manages the metadata and catalog for these tables. Snowflake connects to the storage location by using an external volume that stores identity and access management (IAM) information. You must create and configure the external volume beforehand.

To enable the feature, set the target custom property **writerTargetTableFormat** to **iceberg** when you configure an ingestion and replication task. Also set the **writerTargetTableExternalVolumeName** property to the name of the external volume you created.

For more information, see the *Database Ingestion and Replication* and *Application Ingestion and Replication* documentation for Snowflake targets.

Improved DBMI AGENT service for minimizing downtime during upgrades

A new light-weight version of the Database Ingestion agent service (DBMI Agent) that runs under the Secure Agent is available to run essential services during an upgrade. The older version of the DBMI Agent service shuts down after migrating all of the initial load and CDC jobs to the new DBMI Agent. Incremental load and combined load jobs that are performing CDC processing are stopped and restarted on the new DBMI Agent from where they last left off. Initial load jobs or combined load jobs in the unload phase are restarted from the beginning on the new DBMI Agent. Because the new DBMI Agent is available to accept tasks during the upgrade process, application ingestion and replication jobs and database ingestion and replication jobs experience minimal downtime. If the new DBMI Agent fails to start, for example, because of low memory on the Secure Agent machine, the older version of the DBMI Agent continues running for up to 24 hours to ensure no service disruption occurs while startup of the new DBMI Agent is retried.

Note: Database Ingestion agent service logs, such as Metadata Manager, Task Container, and DBMI_Agent logs, now include the application package version to allow for easy identification of the logs' application version and for the older DBMI Agent version to run during migration to the new DBMI Agent. For more information, see "New log file names for the Database Ingestion agent service" in "Changed behavior" on page 39.

Database Ingestion and Replication

The October 2025 release of Database Ingestion and Replication includes the following new features and enhancements:

Db2 for LUW Log-based CDC

Effective in the October 2025 release, this feature is available to organizations for which Informatica has enabled the feature. To request access, contact Informatica Global Customer Support.

Database ingestion and replication incremental load and combined load tasks that have a Db2 11.*x* for LUW source and a Snowflake target can now capture change data from the database logs. The Db2 source must be on a RedHat Linux or Windows system.

When you configure a task in the latest configuration wizard, on the **Task Details** page for the source, select **Log-based** in the **CDC Method** field and set the associated **Catalog Table Name**, **Capture Threading**, and **Log Read Buffer Size** fields. You can optionally include the tasks in a CDC staging group and use schema drift.

Before you use this CDC method, complete the following prerequisites:

- Ensure that archive logging is active for the Db2 source database from which change data will be captured.
- Define Db2 source tables with the DATA CAPTURE CHANGES clause. If you use the deprecated LONG VARCHAR or LONG VARGRAPHIC data types, define the tables with the DATA CAPTURE CHANGES INCLUDE LONGVAR COLUMNS clause. The CDC script that is generated when you configure a task includes SQL statements for setting this option on each source table.
- Ensure that a valid Db2 environment exists for the Database Ingestion and Replication user. On Linux, the path to the Db2 client must be specified in the PATH and LD_LIBRARY_PATH environment variables.

Tasks that use this CDC method can process source tables that use row compression, database encryption, or table partitioning. Also, the source can be in multiple-partition cluster.

This CDC method does not support the following items: sources on AIX or AWS systems, Db2 pureScale environments, LOB and XML data types, and user-defined DISTINCT and STRUCT data types.

For more information, see the *Database Ingestion and Replication* documentation.

MongoDB sources in combined initial and incremental load tasks

Database Ingestion and Replication now supports MongoDB sources in combined initial and incremental load tasks that have any supported target type.

The MongoDB source can use Community Edition version 4.*x*, 5.*x*, 6.*x*, or 7.*x*, with replica sets or sharded clusters.

For more information, see the *Database Ingestion and Replication* documentation.

SAP HANA sources in combined initial and incremental load tasks

Database Ingestion and Replication now supports SAP HANA log-based and trigger-based CDC sources in combined initial and incremental load tasks that have any supported target type. The following limitations apply:

- Job-level resynchronization of combined load jobs with SAP HANA log-based or trigger-based CDC sources is not supported. Table-level resynchronization continues to be supported.
- For table-level synchronization of SAP HANA log-based or trigger-based CDC sources, the default option is Resync (refresh). This option drops the table and re-creates the source schema at the point of resynchronization. However, after the resynchronization, if an Add Column schema drift change occurs, data for the newly added column isn't captured.

For more information, see the *Database Ingestion and Replication* documentation.

Custom data-type mappings for PostgreSQL sources and Oracle targets

You can now define rules to create custom mappings of PostgreSQL source data types to and Oracle target data types if you want to override the default data-type mappings for Database Ingestion and Replication.

You can add custom data-type mapping rules on the Task Details -Target Details page when creating a task.

Custom data-type mapping rules are applied when the target table is created during task deployment and also during schema drift processing. Custom data-type mappings are not evaluated at runtime when the job loads data to the target table, with the exception of schema drift handling.

For more information, see the *Database Ingestion and Replication* documentation.

Query-based CDC method enabled for Oracle and SQL Server sources with serverless runtime environments

You can now use the Query-based CDC method with a serverless runtime environment for tasks that have SQL Server or Oracle sources and Amazon Redshift, Databricks, Google BigQuery, Microsoft Azure Synapse Analytics, Oracle, PostgreSQL, SQL Server, or Snowflake targets.

For more information, see the *Database Ingestion and Replication* documentation.

Changed behavior

The October 2025 release of Data Ingestion and Replication includes the following changed behavior.

Oracle LOB data resolved as "Not Available" if it cannot be fetched from the database

For database ingestion and replication jobs that have an Oracle source, if LOB data is incompletely logged or cannot be interpreted and it also cannot be fetched directly from the database because the row has been deleted or moved, the LOB data will be replicated as "Not Available".

Previously, if LOB data could not be fetched from the database because the original row was no longer available, the task failed.

Simplified job details to improve clarity and usability

For streaming ingestion and replication jobs, the **Performance** tab, which previously displayed throughput graphs for source and target along with job performance metrics such as total messages and total kilobits of messages streamed per second is no longer included in the Job Details page.

You can refer to the **Overview** tab in the Job Details page to view the total number of events and related performance insights.

Change to the Data Ingestion and Replication CLI installation path

The Data Ingestion and Replication Command-Line Interface (CLI) files are now installed in the `<Cloud Secure Agent installation>\apps\dbmicli` directory. The directory contains .jar files, sample and customized YAML input files, default and customized dbmiclienv properties files, and any key-store.txt file that is generated when encryption is enabled. Previously, the CLI files were installed in the version-dependent path `<Cloud Secure Agent installation>\apps\Database_Ingestion\<latest_version>\dbmicli`.

Note: Before the release, back up your existing files in the old path, including the key-store, dbmiclienv.properties, and YAML files. After the upgrade, copy the files to the new dbmicli directory path.

The installation also now delivers the dbmiclienv_default.properties file instead of the dbmiclienv.properties file. You can use the default properties file to create a customized one. Your customized file will not be overwritten by subsequent installs.

Source NULL data gets loaded as NULL to Amazon Redshift targets

For application ingestion and replication tasks and database ingestion and replication tasks that have an Amazon Redshift target, if a source column that has a string data type contains NULL data, the data is replicated as `infa_null` to the CSV file in the staging location. Data Ingestion and Replication uses the `NULL AS 'infa_null'` clause to interpret any occurrence of 'infa_null' in the CSV file as an actual NULL in the target table.

Previously, when input data contained NULL values for string data type columns, the values were written as empty strings ("") to the Amazon Redshift target.

Output format selection for Amazon S3 targets affected by new Open Table Format options

With the Open Table Format now available to load data to Amazon S3, there's a small change in how you select output formats for Amazon S3 in application and database ingestion and replication task properties. Previously, you chose the output file format directly as CSV, Avro, or Parquet. Now, you first select either **Apache Iceberg** or **None** in the **Open Table Format** field. If you select **None**, you can then select CSV, Avro, or Parquet as before. If you select **Apache Iceberg**, the output format is set to Parquet by default.

Teradata INTEGER and SMALLINT data types are loaded as INTEGER to Amazon S3 targets

Database ingestion and replication initial load tasks that have a Teradata source now map the INTEGER and SMALLINT data types to the INTEGER data type for an Amazon S3 target when the encoding format is configured as Avro or Parquet.

Previously, the INTEGER and SMALLINT data types were loaded to an Amazon S3 target as the STRING data type.

New log file names for the Database Ingestion agent service

The names of the log files that the Database Ingestion service of the Secure Agent generates for its sub-services have been changed to include the application package version. Including the package version in the log file names enables you to easily differentiate the logs generated by different versions of the application and to run the older DBMI Agent version while migrating to the new DBMI Agent to avoid downtime.

This change applies to each sub-service that runs under the Database Ingestion service. For example:

For the DBMI Agent:

- agent_start.log is renamed to `agent_start_<package_version>.log`
- agent_stop.log is renamed to `agent_stop_<package_version>.log`
- agent_status.log is renamed to `agent_status_<package_version>.log`

For the task container: container_service.log is renamed to `container_service_<package_version>.log`.

For Metadata Manager: metadata_manager.log is renamed to `metadata_manager_<package_version>.log`.

Note: For information about the new light-weight DBMI Agent version, see "Improved DBMI AGENT service for minimizing downtime during upgrades" under "New features and enhancements" on page 34 > Common.

Snowflake Ingest SDK and JDBC driver upgrades

The Snowflake Ingest SDK that's packaged with the Database Ingestion agent has been upgraded to version 4.1 to improve the performance and stability of application and database ingestion and replication jobs with Snowflake targets that use Snowflake Streaming (Superpipe option). Also, the Snowflake JDBC driver has been upgraded to version 3.25.0.

Simplification of Oracle OCI configuration in serverless deployments

For database ingestion and replication tasks that have an Oracle source and use the serverless deployment, you can now configure the Oracle OCI driver using a standard Oracle instant client zip.

It is no longer required to rename the internal Oracle OCI driver zip folder and add the `libclntsh.sl` symbolic link to the zip file contents before copying the file to the Azure location.

For more information , see Data Ingestion and Replication *Connectors and Connections* > Oracle Database Ingestion connection properties > Use the serverless runtime environment

Using GV$TRANSACTION instead of V$TRANSACTION for Oracle sources

Database Ingestion and Replication now uses the GV$TRANSACTION global view instead of the V $TRANSACTION view by default to check for open transactions for CDC processing. Configure your Oracle privileges to grant SELECT to PUBLIC on GV$TRANSACTION. Using GV$TRANSACTION ensures that results for all nodes in a RAC cluster are returned.

GV$TRANSACTION is the default view for new jobs created from the current release. For existing jobs that use the V$TRANSACTION view, if the permission to access the GV$TRANSACTION view is missing, Data Ingestion and Replication generates an alert and the job continues to run with the V$TRANSACTION view.

INDEX

B

behavior changes
- Data Ingestion and Replication April 2026 29
- Data Ingestion and Replication July 2025 39
- Data Ingestion and Replication July 2026 19
- Data Ingestion and Replication October 2026 15

D

Data Ingestion and Replication
- changed behavior in April 2026 29
- changed behavior in July 2025 39
- changed behavior in July 2026 19
- changed behavior in October 2026 15

Data Ingestion and Replication *(continued)*
- new features in April 2026 22
- new features in July 2026 18
- new features in November 2025 32
- new features in October 2025 34
- new features in October 2026 9

N

new features
- Data Ingestion and Replication April 2026 22
- Data Ingestion and Replication July 2026 18
- Data Ingestion and Replication November 2025 32
- Data Ingestion and Replication October 2025 34
- Data Ingestion and Replication October 2026 9