How-To Library

Informatica

Deploying the Informatica Enterprise Data Catalog 10.4.0 AWS Marketplace Solution

© Copyright Informatica LLC 2019, 2022. Informatica, the Informatica logo, and Data Engineering Integration are trademarks or registered trademarks of Informatica LLC in the United States and many jurisdictions throughout the world. A current list of Informatica trademarks is available on the web at https://www.informatica.com/trademarks.html.

Abstract

This deployment reference provides step-by-step instructions for deploying Informatica Enterprise Data Catalog on the Amazon Web Services Marketplace. Automated reference deployments use AWS CloudFormation templates to launch, configure, and run the AWS compute, network, storage, and other services required to deploy a specific workload on AWS.

Supported Versions

- Enterprise Data Catalog 10.4.0

Table of Contents

Overview. 2
Intended Audience. 3
Prerequisites. 3
Specialized Knowledge. 3
Deployment Prerequisites. 3
Network Prerequisites. 3
User Permissions. 4
Costs and Licenses. 4
AWS Resources Created by the Deployment. 5
Architecture. 7
Deployment Options. 9
Deployment Steps. 9
Step 1: Prepare an AWS Account. 9
Step 2: Obtain a License for Enterprise Data Catalog. 10
Step 3: Launch the Deployment. 10
Verifying Lambda CloudWatch Logs for Failures. 16
Verifying AWS CloudWatch Logs for Failures. 16
FAQ. 17
Additional Resources. 18

Overview

This deployment reference provides step-by-step instructions for deploying Informatica Enterprise Data Catalog on the AWS Marketplace. You can deploy Enterprise Data Catalog on an new Virtual Private Cloud (VPC) or an existing VPC. Automated reference deployments use AWS CloudFormation templates to launch, configure, and run the AWS compute, network, storage, and other services required to deploy a specific workload on AWS.

In addition, the deployment reference also provides information about the required prerequisites, user permissions, associated costs and licenses, and troubleshooting information.

This guide is for users who are responsible for deploying Enterprise Data Catalog on the AWS Marketplace.

Intended Audience

As a user with administrator privileges to deploy applications on AWS, it is assumed that you are familiar with AWS resources such as Cloud formation, VPC, EC2, S3, RDS, Internet gateway, NAT Gateway, Route table, Security group, and Elastic IP.

Familiarity with the concepts of IP CIDR and public and private IP addresses is recommended. In the current deployment, you can connect to cloud-based resources only using a public IP address.

You must also be familiar with Informatica Enterprise Data Catalog.

Prerequisites

The following sections provide information about concepts that you must know, network prerequisites, and user permissions that you must configure before you begin the deployment:

Specialized Knowledge

Before you deploy Enterprise Data Catalog, we recommend that you become familiar with the following AWS services:

- Amazon VPC
- Amazon EC2
- Amazon RDS
- Amazon S3
- Elastic IP Addresses

If you are new to AWS, see Getting Started with AWS.

Deployment Prerequisites

Before you deploy Enterprise Data Catalog in the AWS Marketplace, verify that you have performed the following steps:

- Complete all the network prerequisites.
- Verify the user permissions.

Network Prerequisites

Before you deploy Enterprise Data Catalog in the AWS Marketplace, verify that you have completed the following network prerequisites for the Enterprise Data Catalog application services:

VPC and Subnet Prerequisites

- Ensure that you set the DNS resolution to **True**.
- Make sure that each subnet in the VPC is associated with a route table.
- The routes in the route table must have the route defined to any of the following gateways:
  - Internet Gateway
  - NAT Gateway
  - Virtual Private Gateway
- Ensure that the subnet associated with Internet Gateway has the Auto-assign public IPv4 property enabled.

- Make sure that the subnet is associated with any of the following gateways:
  - Internet Gateway
  - NAT Gateway
  - Virtual Private Gateway

Existing VPC Deployment

- Ensure that the VPC has the DNS resolution enabled.
- You must attach the subnet to a route table. Ensure that you use a local route to the VPC CIDR.
- Make sure that the service subnets you choose are attached to a route table under the selected VPC.
- The VPC must meet the following requirements:
  - It must be set up with public access through the Internet via an attached Internet gateway.
  - The **DNS Resolution** property of the VPC must be set to **Yes**.
  - The **Edit DNS Hostnames** property of the VPC must be set to **Yes**.

New VPC Deployment

- For the availability zone, ensure that selected availability zone has sufficient capacity to create a new subnet, route table, internet gateway, and NAT gateway.
- Ensure that the VPC creation has not exceeded the supported limit.

User Permissions

Before you deploy Enterprise Data Catalog on the AWS Marketplace, verify that you have the required user permissions to launch Enterprise Data Catalog application services on the AWS Marketplace.

Ensure that the AWS user account that you use has the administrator privileges or the following policies assigned to the user account:

- arn:aws:iam::aws:policy/AWSCloudFormationFullAccess
- arn:aws:iam::aws:policy/AmazonEC2FullAccess
- arn:aws:iam::aws:policy/CloudWatchFullAccess
- arn:aws:iam::aws:policy/IAMFullAccess
- arn:aws:iam::aws:policy/AmazonS3FullAccess
- arn:aws:iam::aws:policy/AmazonRDSFullAccess
- arn:aws:iam::aws:policy/AmazonVPCFullAccess
- arn:aws:iam::aws:policy/AWSLambdaFullAccess

Costs and Licenses

You are responsible for the cost of the AWS services used while running this deployment. There is no additional cost for using this deployment.

The AWS CloudFormation template for this deployment includes configuration parameters that you can customize. Some of these settings, such as instance type, will affect the cost of deployment. See the pricing pages for each AWS service you plan to use, for cost estimates.

This deployment requires a license for Informatica Enterprise Data Catalog. To sign up for a demo license, please contact Informatica.

The following table lists the instance types that you can choose based on sizing requirements:

Virtual Machine	Instance Type	Cluster Size
Database	M5.xlarge	Small, Medium, Large
Informatica Domain	M4.2xlarge / M5.2xlarge	Small, Medium, Large
Informatica Cluster Service	M4.4xlarge / M5.4xlarge	Small
	M4.2xlarge / M5.2xlarge	Medium, Large
Bastion Server	M5.xlarge/ M4.xlarge	Small, Medium, Large

See the Enterprise Data Catalog EDC Performance Tuning Guide to choose an SLA-based sizing recommendation.

The deployment chooses M4 or M5 type of instances based on the instance type availability for the specific AWS region.

AWS Resources Created by the Deployment

The following components are created when you deploy Enterprise Data Catalog in a new Virtual Private Cloud (VPC):

Component	Number of Components Created
VPC	1
Internet gateway	1
Subnet	2
Route table	2
Security group	4 security groups: - 1 for RDS, 1 for Informatica domain instance, and 2 for Hadoop nodes - Based on your choice, 1 security group for RDP server is created for to support remote log in. The group is created if you selected the option to create the RDP server as **Yes**.
DB subnet group	1
RDS	1 Oracle RDS instance with 200 GB of storage capacity that runs on a db.m5.xlarge DB instance type.

Component	Number of Components Created
EC2	1. Informatica domain with 1 EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge / m4.2xlarge based on availability of the instance size in the selected region. Based on your choice, one Informatica high availability node with an EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, a small-sized Informatica Hadoop Cluster with 1 EC2 instance and 150 GB of storage space is created. The instance created is of type. m5.4xlarge / m4.4xlarge based on availability of the instance size in the selected region. Based on your choice, a medium-sized Informatica Hadoop Cluster with 3 EC2 instances with each instance providing a storage space of 150 GB is created. The instances created are of type m5.2xlarge or m4.2xlarge based on the availability of the instance size in the selected region. Based on your choice, a large-sized Informatica Hadoop Cluster with 6 EC2 instances with each instance providing a storage space of 150 GB is created. The instances created are of type m5.2xlarge or m4.2xlarge based on availability of the instance size in the selected region. Based on your choice, a Windows Remote Desktop server with 1 EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large or m4.large based on availability of the instance size in the selected region.
Elastic IP	1 Elastic IP attached to the NAT gateway.
Internet Gateway	1
IAM Role and Instance Profile	1

The following components are created when you deploy Enterprise Data Catalog in an existing VPC:

Component	Number of Components Created
Security group	4 security groups: - 1 for RDS, 1 for Informatica domain instance, and 2 for Hadoop nodes - Based on your choice, 1 security group for RDP server is created for to support remote log in. The group is created if you selected the option to create the RDP server as **Yes**.
DB subnet group	1
RDS	1 Oracle RDS instance with 200 GB of storage capacity that runs on a db.m5.xlarge DB instance type.

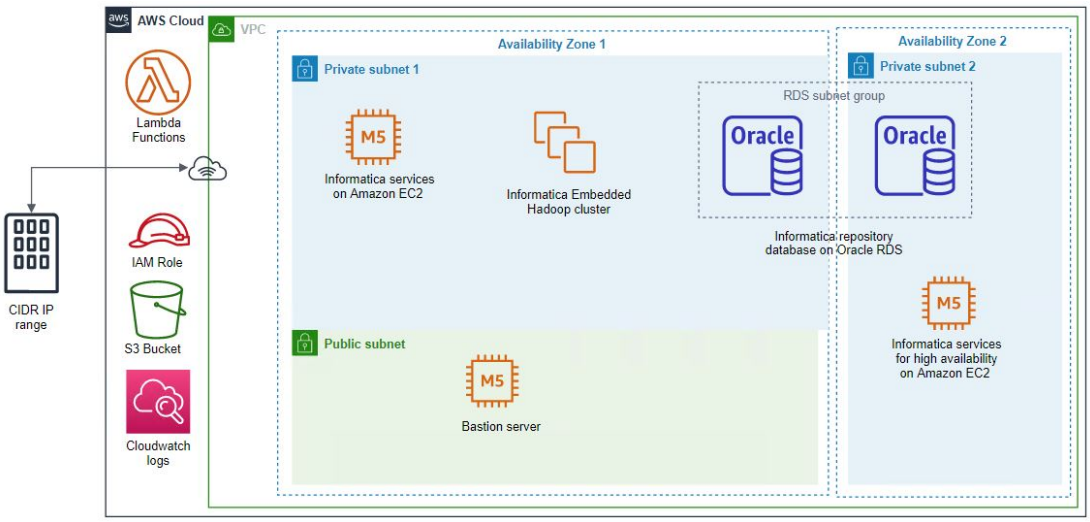
Component	Number of Components Created
EC2	1. Informatica domain with 1 EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge / m4.2xlarge, based on availability of the instance size in the selected region. Based on your choice, one Informatica high availability node with an EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, a small-sized Informatica Hadoop Cluster with 1 EC2 instance and 150 GB of storage space is created. The instance created is of type. m5.4xlarge / m4.4xlarge based on availability of the instance size in the selected region. Based on your choice, a medium-sized Informatica Hadoop Cluster with 3 EC2 instances with each instance providing a storage space of 150 GB is created. The instances created are of type m5.2xlarge / m4.2xlarge based on the availability of the instance size in the selected region. Based on your choice, a large-sized Informatica Hadoop Cluster with 6 EC2 instances with each instance providing a storage space of 150 GB is created. The instances created are of type m5.2xlarge / m4.2xlarge based on availability of the instance size in the selected region. Based on your choice, a Windows Remote Desktop server with 1 EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large / m4.large based on availability of the instance size in the selected region.
AWS Lambda Function	5 AWS Lambda functions to check the feasibility for VPC, subnets, and Informatica license.
IAM Role and Instance Profile	1

Architecture

When you deploy Enterprise Data Catalog Marketplace solution on a new Virtual Private Cloud (VPC) with the default parameters, the deployment builds the following Enterprise Data Catalog environment in the AWS Marketplace.

The following image shows the architecture of the Enterprise Data Catalog on AWS Marketplace:

Figure 1. Deployment architecture for Enterprise Data Catalog on AWS Marketplace



The architecture includes the following components:

- A VPC configured across two Availability Zones. For each Availability Zone, this deployment provisions one private and one public subnet. The Enterprise Data Catalog deployment uses both subnets.
- Managed network address translation (NAT) gateways deployed into the public subnets and configured with an Elastic IP address for outbound internet connectivity.
- An AWS Identity and Access Management (IAM) role with fine-grained permissions for access to AWS services necessary for the deployment process.
- Appropriate security groups for each instance or function to restrict access to only necessary protocols and ports.
- In the public subnets, EC2 instances for Enterprise Data Catalog, including the following:
  - A single-node or multi-node, embedded cluster based on the cluster size you configure. A node represents an EC2 machine. The Informatica domain runs on a single node with all the associated services. The cluster runs on a separate node or nodes based on your selection. The application services such as the model repository service and the data integration service run on the Informatica domain node.
  - Scanners to extract metadata from all data sources supported by Enterprise Data Catalog.
  - **Oracle database**, where the Informatica domain configuration repository, the model repository, the content management service reference data warehouse, and the profiling warehouse are configured.
  - **Informatica domain**, which is the fundamental administrative unit of the Informatica platform. The Informatica platform has a service-oriented architecture that provides the ability to scale services and share resources across multiple machines.
  - **Model Repository Service**, which is a relational database that stores all the metadata for projects created using Informatica client tools. The model repository also stores run-time and configuration information for applications that are deployed to a Data Integration Service. The information relating to the administrative module of Enterprise Data Catalog, called Catalog Administrator, is stored in the model repository.
  - **Data Integration Service**, which is a compute component within the Informatica domain that manages requests to submit big data integration, big data quality, and profiling jobs to the Hadoop cluster for processing.
  - **Informatica Cluster Service**, which runs and manages all the Hadoop services, Apache Ambari server, and Apache Ambari agents on an internal Hadoop cluster.
  - **Catalog Service**, which runs Enterprise Data Catalog and manages connections between service components and external applications.
  - **Content Management Service**, which manages reference data. It provides reference data information to the Data Integration Service and Informatica Developer.
  - **Analyst Service**, which runs the Analyst tool in the Informatica domain. The Analyst Service manages the connections between the service components and the users who log in to the Analyst tool. You can perform column and rule profiling, manage scorecards, and manage bad records and duplicate records in the Analyst tool.
  - **Profiling**, which helps you find the content, quality, and structure of data sources of an application, schema, or enterprise. A profile is a repository object that finds and analyzes all data irregularities across data sources in the enterprise and hidden data problems that put data projects at risk. The profiling results include unique values, null values, data domains, and data patterns. In this deployment, profiling can be run on the Data Integration Service (default) or Hadoop.
  - **Business Glossary**, which consists of online glossaries of business terms and policies that define important concepts within an organization. Data stewards create and publish terms that include information such as descriptions, relationships to other terms, and associated categories. Glossaries are stored in a central location for easy lookup by end-users. Glossary assets include business terms, policies, and categories that contain information that consumers might search for. A glossary is a high-level container that stores

Glossary assets. A business term defines relevant concepts within the organization, and a policy defines the business purpose that governs practices related to the term. Business terms and policies can be associated with categories, which are descriptive classifications.

- **Metadata and Catalog**, which include the metadata persistence store, search index, and graph database in an embedded Hortonworks cluster. The catalog represents an indexed inventory of all the data assets in the enterprise that you configure in Enterprise Data Catalog. Enterprise Data Catalog organizes all the enterprise metadata in the catalog and enables the users of external applications to discover and understand the data.

You can collect metadata both from resources deployed on a VPC and from resources deployed on premises. For on-premises resources, Enterprise Data Catalog collects metadata through a virtual private network (VPN), a VPC, or by using AWS Direct Connect.

After deploying Informatica Enterprise Data Catalog, you must configure the required resources to extract metadata.

If you decide to deploy Enterprise Data Catalog into your existing VPC (see Deployment Options), the deployment assumes that the infrastructure components already exist, and deploys Enterprise Data Catalog into the environment you specify during deployment.

Deployment Options

You can select one of the following deployment options:

- Deploy Enterprise Data Catalog into a new VPC (end-to-end deployment). This option builds a new AWS environment consisting of the VPC, subnets, NAT gateways, security groups, and other infrastructure components, and then deploys Enterprise Data Catalog into this new VPC.
- Deploy Enterprise Data Catalog into an existing VPC. This option provisions Enterprise Data Catalog in your existing AWS infrastructure.

The deployment also lets you configure additional settings such as CIDR blocks, instance types, and Enterprise Data Catalog settings, as discussed later in this guide.

Deployment Steps

Perform the following steps to deploy the Informatica Enterprise Data Catalog 10.4.0 AWS Marketplace solution:

Step 1: Prepare an AWS Account

Perform the following steps to prepare an AWS account:

1. If you don't already have an AWS account, create an account at http://aws.amazon.com by following the on-screen instructions.
2. Use the region selector in the navigation bar to choose the AWS Region where you want to deploy Enterprise Data Catalog on AWS.
3. Create a key pair in your preferred region.
4. When you log in to an EC2 instance, you use a password file for authentication. The file is called a private key file and has a file name extension of .pem.
5. If you do not have an existing .pem key to use, follow the instructions in the AWS documentation to create a key pair.

   Note: Your administrator might ask you to use an existing key pair.

   When you create a key pair, you save the .pem file to your desktop system. Simultaneously, AWS saves the key pair to your account. Make a note of the key pair that you want to use for the Enterprise Data Catalog instance, so that you can provide the key pair name during network configuration.

6. If necessary, request a service limit increase for the EC2 instance type you've decided to use for the Informatica domain. You might need to do this if you already have an existing deployment that uses this instance type, and you think you might exceed the default limit with this reference deployment.

Step 2: Obtain a License for Enterprise Data Catalog

1. This deployment requires a license for Informatica Enterprise Data Catalog. To sign up for a demo license, please contact Informatica.
2. In your AWS account, create an S3 bucket.

   (Optional) Create a directory under the bucket.

   Place the Enterprise Data Catalog license key file for the software in the S3 bucket or directory. You'll be prompted for its location in step 3.

Step 3: Launch the Deployment

You are responsible for the cost of the AWS services used while running this deployment. There is no additional cost for using this deployment. See the pricing pages for each AWS service you will be using for full details. Each deployment takes approximately one to two hours to complete.

Perform the following steps to launch the deployment:

1. Login to AWS Marketplace using the AWS account.
2. Type `Enterprise Data Catalog` in the search box and press **Enter**.
3. Click **Informatica Enterprise Data Catalog**. The **Product Overview** page appears.
4. Click **Continue to Subscribe**. The **Subscribe to this software** page appears.
5. Click **Continue to Configuration**. The **Configure this software** page appears.
6. Based on your requirement, select one of the following options from the **Fulfillment Option** drop-down box to specify where you want to deploy Enterprise Data Catalog:
   - Informatica Enterprise Data Catalog (New VPC). Select this option to deploy Enterprise Data Catalog on a new VPC.
   - Informatica Enterprise Data Catalog (Existing VPC). Select this option to deploy Enterprise Data Catalog on an existing VPC.
7. Click **Continue to Launch**. The **Launch this software** page appears.
8. Select **Launch CloudFormation** from the **Choose Action** drop-down list and click **Launch**. The **Create stack** page appears.
9. Check the region that's displayed in the upper-right corner of the navigation bar, and change it if necessary. This is where the network infrastructure for Enterprise Data Catalog will be built. The template is launched in the US West (Oregon) Region by default.
10. On the **Specify template** section, keep the default setting for the template URL, and then choose **Next**.
11. On the **Specify stack details** page, change the stack name if needed. Review the parameters for the template. Provide values for the parameters that require input. For all other parameters, review the default settings and customize them as necessary. When you finish reviewing and customizing the parameters, choose **Next.**

    In the following tables, parameters are listed by category and described separately for the two deployment options:

    - "Option1:Parameters for deployment into a new VPC" on page 11
    - "Option 2: Parameters for Deployment in an Existing VPC" on page 12

Note: If you're deploying Enterprise Data Catalog into an existing VPC, make sure that your VPC has a minimum of two subnets. If you choose a remote desktop, you can either use one of the two subnets or use a separate public subnet. These subnets require NAT gateways in their route tables, to allow the instances to download packages and software without exposing them to the Internet. You'll also need the domain name option configured in the DHCP options as explained in the Amazon VPC documentation. You'll be prompted for your VPC settings when you launch the deployment.

Option1:Parameters for deployment into a new VPC

The following table lists the parameters in the **Network Configuration** section:

Parameter	Default	Description
Availability Zones	Requires input	The list of Availability Zones to use for the subnets in the VPC. Make sure that you select two zones.
VPC CIDR	10.0.0.0/16	CIDR block for the VPC.
Public Subnet CIDR		CIDR block for the public subnet. Ensure that you include the subnet in the VPC CIDR.
Private Subnet 1 CIDR	10.0.0.0/19	CIDR block for the private subnet 1. Ensure that you include the subnet in the VPC CIDR.
Private Subnet 2 CIDR	10.0.32.0/19	CIDR block for the private subnet 2. Ensure that you include the subnet in the VPC CIDR.
CIDR IP Range	Requires input	The CIDR IP range that is permitted to access Enterprise Data Catalog. To reduce the potential of inbound attacks from unknown IP addresses, you can use a constrained CIDR range to reduce. For example, to specify the range of 10.20.30.40 to 10.20.30.49, enter 10.20.30.40/49.

The following table lists the parameters in the **Bastion Server for Informatica** section:

Parameter	Default	Description
Deploy Bastion Server	Requires input	Select **Yes** to deploy a bastion server to access other resources in the VPC.
Bastion Subnet	Requires input	Select the subnet for the bastion server if you selected **Yes** for the previous option. Ensure that the subnet is included in the VPC that you use and attached to an Internet Gateway.

The following table lists the parameters in the **Amazon EC2 Configuration** section:

Parameter label (name)	Default	Description
Key Pair Name	Requires input	Name of an existing Amazon EC2 key pair to enable SSH and Remote Desktop Protocol (RDP) access to Enterprise Data Catalog

The following table lists the parameters in the **Informatica Enterprise Data Catalog Configuration** section:

Parameter label (name)	Default	Description
License Key S3 Bucket URI	Requires input	The S3 bucket URI for the Enterprise Data Catalog license key in the following format: `s3://<bucket-name>/<directory>/<license key>`
High Availability Deployment	Requires input	Enable or disable high availability for the deployment.
Deployment Type	Requires input	Size of the Informatica embedded cluster that you want to create based on the region you selected. You can select one of the following sizes: - Small - Medium - Large

Option 2: Parameters for Deployment in an Existing VPC

The following table lists the parameters in the **Network Configuration** section:

Parameter	Default	Description
VPC	Requires input	ID of your existing VPC where you want to deploy Enterprise Data Catalog.
Subnet 1	Requires input	Select the subnet ID where you want to deploy Enterprise Data Catalog. Make sure that the subnet is included in the VPC and associated with an Internet Gateway or a NAT Gateway. Ensure that service subnet ID1 and service subnet are not the same.
Subnet 2	Requires input	If you want to deploy Enterprise Data Catalog with high availability enabled, provide the subnet ID that you want to use for high availability. Make sure that subnet 1 and subnet 2 are in different availability zones.
CIDR IP Range	Requires input	The CIDR IP range that is permitted to access Enterprise Data Catalog. To reduce the potential of inbound attacks from unknown IP addresses, you can use a constrained CIDR range. For example, to specify the range of 10.20.30.40 to 10.20.30.49, enter 10.20.30.40/49.
Pre-validation Check	Requires input	Select to validate the VPC and subnets before deploying Enterprise Data Catalog.

The following table lists the parameters in the **Bastion Server for Informatica** section:

Parameter	Default	Description
Deploy Bastion Server	Requires input	Select **Yes** to deploy a bastion server to access other resources in the VPC.
Bastion Subnet	Requires input	Select the subnet for the bastion server if you selected **Yes** for the previous option. Ensure that the subnet is included in the VPC that you use and attached to an Internet Gateway.

The following table lists the parameters in the **Amazon EC2 Configuration** section:

Parameter label (name)	Default	Description
Key Pair Name	Requires input	Name of an existing Amazon EC2 key pair to enable SSH and Remote Desktop Protocol (RDP) access to Enterprise Data Catalog

The following table lists the parameters in the **Informatica Enterprise Data Catalog Configuration** section:

Parameter label (name)	Default	Description
License Key S3 Bucket URI	Requires input	The S3 bucket URI for the Enterprise Data Catalog license key in the following format: `s3://<bucket-name>/<directory>/<license key>`
High Availability Deployment	Requires input	Enable or disable high availability for the deployment.
Deployment Type	Requires input	Size of the Informatica embedded cluster that you want to create based on the region you selected. You can select one of the following sizes: - Small - Medium - Large

1. Click **Next** after you configure the parameters in the **Informatica Enterprise Data Catalog Configuration** section.
2. On the **Configure stack options** page, you can specify tags (key-value pairs) for resources in your stack and set advanced options. When you are done, choose **Next**.
3. On the **Review** page, review and confirm the template settings. Under Capabilities, select the check box to acknowledge that the template will create IAM resources.
4. Choose **Create** to deploy the stack.
5. Monitor the status of the stack. When the status is `CREATE_COMPLETE`, the deployment is complete.

You can use the keys displayed in the **Output** page as shown in the following table to view the details of the deployment:

Key	Description
AdminConsoleUser	Username to access the Informatica domain.
AdminConsoleUserPwd	Password for the Informatica domain username.
BastionPrivateDNS	Private DNS of the Informatica bastion server.
BastionPublicIPDNS	Public DNS of the Informatica bastion server.
DBUserPwd	Password to access the Informatica domain repository database.
DBUsers	Usernames to access the Informatica domain repository database.
HANodeIP	IP address of the Informatica high availability node.

Key	Description
InformaticaAdminConsolePrivateURL	Private URL to access the Informatica Administrator.
InformaticaAdminConsolePublicURL	Public URL to access the Informatica Administrator.
informaticaAdminConsoleServerLogs	Location where the Informatica domain installation log files are stored.
InstanceSetupLogs	Location where the Informatica domain instance setup log files are stored.
NodePrivateEDCURL	Private URL to access Enterprise Data Catalog.
NodePrivateEDCURL1	Private URL to access Catalog Administrator.
NodePublicEDCURL	Public URL to access Enterprise Data Catalog.
NodePublicEDCURL1	Public URL to access Catalog Administrator.

See the Informatica Enterprise Data Catalog User Guide for information about logging in to Enterprise Data Catalog.

A sample completed stack for Enterprise Data Catalog deployed on AWS is shown in the following table:

Stack and Nested Stack Names	Stack Description	Condition for Stack Usage
StackName	Main stack	Not applicable
StackName-LicenseCheckLambdaStack	Template to verify if the license key file is valid.	Default
StackName-BastionSubnetCheckLambdaStack	Template to verify the validity of a subnet.	Default. Creates the AWS Lambda function if the **Pre-validation Check** and the **Bastion Server** parameters are set to **Yes**.
StackName-SubnetsAZCheckLambdaStack	Template to verify if the subnets are in different availability zones.	Default
StackName-PublicIpDnsCheckLambdaStack	Template to verify if an auto-assign public IPv4 address is enabled for subnet.	Default
StackName-Subnet2CheckLambdaStack	Template to verify the validity of a subnet.	Default. Creates the AWS Lambda function if the **Pre-validation Check** parameter is set to **Yes**.
StackName-Subnet1CheckLambdaStack	Template to verify the validity of a subnet.	Default. Creates the AWS Lambda function if the **Pre-validation Check** parameter is set to **Yes**.
StackName-IAMStack	Template to create the IAM roles and profiles used by EC2 instances.	Default
StackName-SecurityGroupStack	Template to create the security groups and network interfaces used by EC2 instances.	Default
StackName-RDSStack	Template to create the RDS resources.	Default

Stack and Nested Stack Names	Stack Description	Condition for Stack Usage
StackName-BastionServerStack	Template to create the bastion server EC2 instances.	Used when the **Deploy Bastion Server** parameter is set to **Yes**.
StackName-IHSStack	Template to create the Informatica embedded Hadoop cluster.	Default
StackName-IHSStack-HadoopGatewayStack	Template to create the Informatica embedded Hadoop cluster gateway node.	Default
StackName-IHSStack-Infa-Hadoop-Node2	Template to create the Informatica embedded Hadoop cluster node.	Used when the **Deployment Type** parameter is set to **Medium**.
StackName-IHSStack-Infa-Hadoop-Node3	Template to create the Informatica embedded Hadoop cluster node.	Used when the **Deployment Type** parameter is set to **Medium**.
StackName-IHSStack-Infa-Hadoop-Node4	Template to create the Informatica embedded Hadoop cluster node.	Used when the **Deployment Type** parameter is set to **Large**.
StackName-IHSStack-Infa-Hadoop-Node5	Template to create the Informatica embedded Hadoop cluster node.	Used when the **Deployment Type** parameter is set to **Large**.
StackName-IHSStack-Infa-Hadoop-Node6	Template to create the Informatica embedded Hadoop cluster node.	Used when the **Deployment Type** parameter is set to **Large**.
StackName-EDCServerStack	Template to create the Informatica Enterprise Data Catalog node.	Default
StackName-EDCHANodeStack	Template to create the Informatica Enterprise Data Catalog high availability node.	Used when the **High Availability Deployment** option is enabled.

Accessing Informatica Administrator and the Apache Ambari Server

Based on the type of deployment, follow the instructions to access Informatica Administrator and the Apache Ambari Server:

- Deployment on a new VPC. Use an instance in the same VPC that you can access from outside the VPC. It is assumed that you have not deployed a remote Windows server to access other resources in the VPC. Use that instance to access other resources. See the following URL formats that you can use:
  - To access the Informatica Administrator, use the URL in the following format: <Private IP Address or private DNS>:6005 and
  - To access the Apache Ambari Server, use the URL in the following format: <Private IP Address or private DNS>:8080
- Deployment on an existing VPC. If the created instances are assigned a public IP address or DNS, you can access the Informatica Administrator and the Apache Ambari Server using the public IP address or the DNS with port numbers as shown in the following sample:
  - To access the Informatica Administrator, use the URL in the following format: *<Public IP Address or DNS>:8443*
  - To access the Apache Ambari Server, use the URL in the following format: *<Public IP Address or DNS>:8080*. You can use the default Hadoop username and password to access the Apache Ambari Server.

- If the public IP address or DNS is not assigned to the created instance, then you can use the private IP address or DNS. Use an instance in the same VPC if you have not deployed a remote Windows server to access other resources in the VPC. Make sure that you can access the instance from outside the VPC. Use that instance to access other resources. For example, to access the Informatica Administrator , use the URL in the following format:*<Private IP Address or private DNS>:6008*. To access the Apache Ambari Server, use the URL in the following format: *<Private IP Address or private DNS>:808*0

Troubleshooting

If you want to troubleshoot issues relating to Informatica Cluster Service, Catalog Service, Installation, and service initialization, you'll have to look at the following log files:

- Installation log file: `/opt/Infa_OneClick_Solution.log`
- Service Initialization log file: `/var/log/cfn-init.log`
- Informatica Cluster Service log file: `/opt/Informatica/Server/logs/InfaNode/services/InfaHadoopService/Informatica_Hadoop_Service/IHS.log`
- Catalog Service log file: `/opt/Informatica/Server/logs/InfaNode/services/CatalogService/Catalog_Service/LDM.log`

If the Informatica EC2 instance is running, but the stack fails or if the Informatica Administrator console does not open, see the following log files to identify the issues:

- /opt/Infa_OneClick_Solution.log
- /var/log/cfn-init-cmd.log
- /var/log/cfn-init.log
- /var/log/cloud-init.log

Verifying Lambda CloudWatch Logs for Failures

Perform the following steps to verify the Lambda CloudWatch logs if any of the Lambda stacks fails:

1. In the Amazon Web Services Console, select **CloudWatch** from the **Services** menu.
2. In the navigation pane, choose **Logs** > **Log Groups**.
3. Enter `/aws/lambda/<stack name>` in the search pane.
4. Select the log group associated with the name of the stack that failed.

Verifying AWS CloudWatch Logs for Failures

CloudWatch is an AWS service that enables you to monitor all of the events in your AWS deployments. CloudWatch provides a central location from which to monitor EC2 instances and other resources. For more information, see the AWSdocumentation.
Perform the following steps to access the Enterprise Data Catalog deployment log files using the AWS console:

1. In the Amazon Web Services Console, select **CloudWatch** from the **Services** menu.
2. In the navigation pane, select **Logs** > **Log Groups**.
3. Search for `Enterprise Data Catalog` in the search pane.

AWS CloudWatch Log Files for Specific Stack Failure Scenarios

See the following AWS CloudWatch log files for the following specific stack failure scenarios:

- Stack fails when the **Pre-validation Check** parameter is set to **Yes**. Search for `/aws/lambda/<Stack-Name>` in the **Log Groups** search box and then click the log group to view the log files for the stack.

- Stack fails when the **Rollback on failure** option is enabled by default. Search for the `InformaticaEDCLogGroup` in the **Log Groups** search box and click any of the following log groups:
  - installation.log
  - cfn-init
  - InformaticaClusterService
  - Catalog Service

FAQ

Q. I encountered a CREATE_FAILED error when I launched the deployment. What should I do?

A. If AWS CloudFormation fails to create the stack, we recommend that you relaunch the template with **Rollback on failure** set to **No**. (This setting is under **Advanced** in the AWS CloudFormation console, **Options** page.) With this setting, the stack's state will be retained, and the instance will be left running, so you can troubleshoot the issue. (You will want to look at the log files in `/var/log/cfn-init-cmd.log`)

Note: When you set **Rollback on failure** to **No**, you will continue to incur AWS charges for this stack. Please make sure to delete the stack when you have finished troubleshooting.

For additional information, see Troubleshooting AWS CloudFormation on the AWS website or contact us on the AWS Discussion Forum.

Q. I encountered errors while installing the Informatica domain and services. Where can I find more information?

A. View the `/opt/Infa_OneClick_Solution.log` log file to see more information about the errors that you encountered.

Q. I encountered a size limitation error when I deployed the AWS CloudFormation templates.

A. We recommend that you launch the deployment templates from the location we've provided or from another S3 bucket. If you deploy the templates from a local copy on your computer or from a non-S3 location, you might encounter template size limitations when you create the stack. For more information about AWS CloudFormation limits, see the AWS documentation.

Q. What should I do before I relaunch the template with **Rollback on failure**?

A. Ensure that you disable rollback when you launch the stack in case of a stack failure. If Informatica Enterprise Data Catalog server stack is created, check the following logs to identify the issues:

- `/opt/Infa_OneClick_Solution.log`
- `/var/log/cfn-init.log`
- `/opt/Informatica/Server/logs/InfaNode/services/InfaHadoopService/Informatica_Hadoop_Service/IHS.log`
- `/opt/Informatica/Server/logs/InfaNode/services/CatalogService/Catalog_ Service/LDM.log`

If the Informatica Enterprise Data Catalog server stack is not created, you can check the CloudWatch logs to identify the issues. For more information, see the Verifying Lambda CloudWatch Logs for Failures section.

Q. I am facing MRS database connectivity issues with the following error messages:

```
MRS_50015 "The Repository Service operation failed. ['[PRSVCSHARED_01707] Connection issues
with the configured database. Request failed with the error message [[PERSISTENCEAPI_0307]
[DBPERSISTER_1005] Failed to process requested operation. This was caused by [informatica]
[Oracle JDBC Driver]No more data available to read.].']"

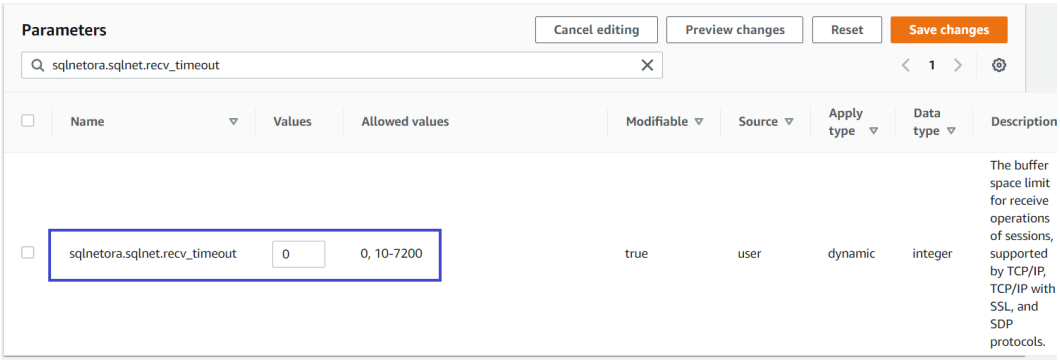
MRS_50015 "The Repository Service operation failed. ['[PRSVCSHARED_01707] Connection issues
with the configured database. Request failed with the error message [[PERSISTENCEAPI_0307]
[PERSISTENCECOMMON_0001] Internal error. The request processing failed. This was caused by
```

```
[informatica][Oracle JDBC Driver]Exception generated during deferred local transaction
handling.  See next exception via SQLException.getNextException for details.]."]"
```

A. Perform the following steps to resolve the MRS database connectivity issue:

1. Launch the AWS RDS page.
2. Select the RDS instance.
3. Click the **Configuration** tab.
4. Scroll down the tab page and click the link under the **Parameter group** section.
5. Search for the following parameters in the list of parameters that appear:
   - sqlnetora.sqlnet.send_timeout
   - sqlnetora.sqlnet.recv_timeout
   - sqlnetora.sqlnet.inbound_connect_timeout
   - sqlnetora.sqlnet.outbound_connect_timeout
   - sqlnetora.tcp.connect_timeout

   Note: If the listed parameters are not present in the list, make sure that add the parameters to the list.
6. Change the value of each of the parameters listed in the previous step to 0 as shown in the following sample image:



7. Click **Save changes**.

Additional Resources

AWS Services

- AWS CloudFormation
- Amazon EBS
- Amazon EC2
- Amazon VPC

Enterprise Data Catalog Documentation

- Product Documentation

Reference Deployments

- AWS Deployment home page

Author

Suraj Jayan