How-To Library

Informatica

Deploying the Informatica® Enterprise Data Preparation 10.4.0 AWS Marketplace Solution

© Copyright Informatica LLC 2019, 2021. Informatica and the Informatica logo are trademarks or registered trademarks of Informatica LLC in the United States and many jurisdictions throughout the world. A current list of Informatica trademarks is available on the web at https://www.informatica.com/trademarks.html.

Abstract

This deployment reference provides step-by-step instructions for deploying the Informatica® Enterprise Data Preparation 10.4.0 AWS Marketplace Solution. Automated reference deployments use AWS CloudFormation templates to launch, configure, and run the AWS compute, network, storage, and other services required to deploy a specific workload on AWS.

Supported Versions

- Enterprise Data Preparation 10.4.0

Table of Contents

Overview. . . . 3
Intended Audience. . . . 3
Specialized Knowledge. . . . 3
Costs and Licenses. . . . 3
Architecture. . . . 4
AWS Resources in the Deployment 6
Informatica Domain. . . . 8
Informatica Clients. . . . 10
Before You Begin 11
Network Prerequisites. . . . 11
Account and Security Prerequisites. . . . 12
Storage Prerequisite. . . . 13
Deploying Enterprise Data Preparation on the AWS Marketplace. . . . 13
Step 1. Get a License for Enterprise Data Preparation. . . . 14
Step 2. Launch the Deployment. . . . 14
Step 3. Configure the Deployment. . . . 15
Step 4. Recycle the Informatica Services. . . . 18
Monitoring Instance Provision and Informatica Domain Creation. . . . 19
Deployed AWS Resources. . . . 19
Output Tab Properties. . . . 21
Deployed Informatica Resources. . . . 22
Solution Deployment Log Files. . . . 22
Troubleshooting. . . . 24
Stack Failure. . . . 24
Stack Launch Failure When the Pre-validation Check Parameter is Enabled. . . . 25
Stack Launch Failure When Rollback On Failure Is Enabled. . . . 25

Additional Resources. 25

Overview

This deployment reference provides step-by-step instructions for deploying the Informatica Enterprise Data Preparation 10.4.0 AWS Marketplace Solution. Automated reference deployments use AWS CloudFormation templates to launch, configure, and run the AWS compute, network, storage, and other services required to deploy a specific workload on AWS.

Enterprise Data Preparation is a collaborative self-service data discovery and preparation solution for data analysts and data scientists. It enables analysts to rapidly discover and turn raw data into insights and allows IT to ensure quality, visibility, and governance. With Enterprise Data Preparation, analysts to spend more time on analysis and less time on finding and preparing data.

Intended Audience

As a user with administrator privileges to deploy applications on AWS, you should be familiar with AWS resources such as CloudFormation, VPC, EC2, S3, RDS, Internet gateway, NAT gateway, route table, security group, and elastic IP. You should also be familiar with concepts such as IP CIDR and public and private IP addresses.

You should also be familiar with Enterprise Data Preparation. To find the product documentation, see the Informatica documentation portal.

Specialized Knowledge

Before you deploy Enterprise Data Preparation, you should be familiar with the following AWS services.

- Amazon VPC
- Amazon EC2
- Amazon RDS
- Amazon S3
- Elastic IP Addresses

To learn more about AWS, see Getting Started with AWS.

Costs and Licenses

You are responsible for the cost of the AWS services used while running this deployment. There is no additional cost for using this deployment.

The AWS CloudFormation template for this deployment includes configuration parameters that you can customize. Some of these settings, such as instance type, will affect the cost of deployment. See the pricing pages for each AWS service you will be using for cost estimates.

The deployment requires a license for Informatica Enterprise Data Preparation. Contact Informatica Global Customer Support to sign up for a demo license.

The following table lists the instance types that you can choose based on sizing requirements:

Virtual Machine	Instance Type	Cluster Size
Database	m5.xlarge	Small, Medium, Large
Informatica Domain	m4.2xlarge / m5.2xlarge	Small, Medium, Large
Bastion Server	m4.xlarge / m5.xlarge	Small, Medium, Large
Informatica Embedded Hadoop Cluster	m4.4xlarge / m5.4xlarge m4.2xlarge / m5.2xlarge	Small Medium, Large
Informatica Compute Cluster on EMR	m4.xlarge / m5.xlarge / m4.2xlarge / m5.2xlarge	Small, Medium, Large

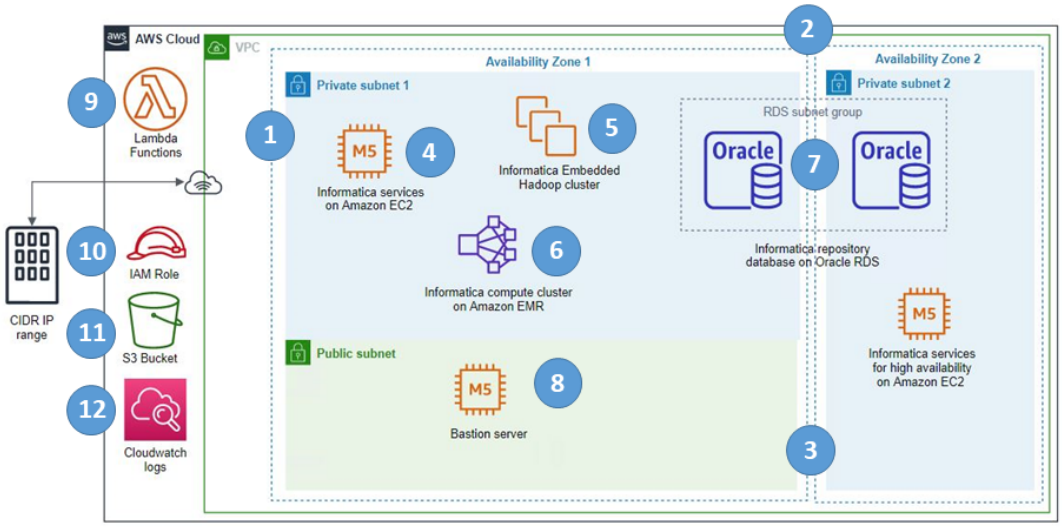
The deployment chooses m4 or m5 instance types based on the instance type availability for the specific AWS region.

Architecture

When you deploy the Enterprise Data Preparation marketplace solution in a new VPC, AWS CloudFormation templates create and connect the following resources in the VPC:

- An Informatica domain server on an EC2 instance, with additional instances to contain nodes in the Data Integration Service grid.
- An embedded Hadoop cluster that Enterprise Data Catalog uses to run metadata processing and profiling jobs.
- Informatica clients on a remote Windows bastion server that runs on a public subnet.
- Amazon S3 storage resources and connections for source and target data in existing Amazon S3 buckets.
- Amazon RDS relational databases for the Informatica repositories.
- AWS security and account management services.
- An Amazon EMR cluster with autoscaling enabled.

The following image shows the architecture of the Informatica Enterprise Data Preparation 10.4.0 AWS Marketplace Solution:



The icons in the architecture diagram correspond to items in the following list:

1. A virtual public cloud (VPC) configured across two Availability Zones to contain the Enterprise Data Preparation deployment.
2. Availability Zones. The deployment provisions two Availability Zones.
3. Subnets to contain specific elements of the deployment. The deployment creates two private subnets, plus one public subnet if you want to use a Windows bastion server for Informatica clients. The deployment creates each of the subnets in a different Availability Zone.
4. The Informatica domain where application services run, including the Catalog Service, the Enterprise Data Preparation Service, the Interactive Data Preparation Service, the Model Repository Service and the Data Integration Service.
5. An internal Hadoop cluster deployed on Amazon EC2 instances that Enterprise Data Catalog uses to run metadata processing and profiling jobs.
6. A compute cluster of Amazon EC2 instances with autoscaling enabled that Enterprise Data Preparation uses to publish prepared data to the data lake.
7. An Oracle database on Amazon RDS for the following Informatica repositories:
   - Data Preparation repository
   - Domain repository
   - Model repository
8. Informatica clients in a separate EC2 bastion server in a public subnet.
9. AWS Lambda functions.
10. IAM roles.
11. Amazon S3 storage, to act as a temporary location to store Enterprise Data Preparation processing data.
12. Amazon CloudWatch.

AWS Resources in the Deployment

The deployment process creates or includes the following components:

- Components created when you specify the solution to create a new VPC.
- Components created, as well as some that already exist, when you specify the solution to use an existing VPC.
- A bastion server, if you choose to create one.
- An internal Hadoop cluster deployed on Amazon EC2 instances that Enterprise Data Catalog uses to run metadata processing and profiling jobs.
- A computer cluster of Amazon EC2 instances with autoscaling enabled that Enterprise Data Preparation uses to publish prepared data to the data lake.

Components in a Deployment on a New VPC

The deployment creates the following components when you deploy the software in a new virtual private cloud:

Component	Number of Components Created
VPC	1
Internet gateway	1
IAM role and instance profile	1
Subnets	4 2
Route tables	2
Security groups	The following security groups: - One security group for RDS. - One security group for the Informatica Enterprise Data Preparation instance. - One security group for the Informatica domain instance. - If you selected the option to create the bastion server, one security group to support remote login to the bastion server. - Two security groups for the embedded Hadoop cluster. - Three security groups for the Informatica compute cluster on EMR. - Based on your choice, two security groups for Informatica EMR workflow connections.
Database subnet group	1
Amazon RDS instance	1 Amazon RDS for Oracle instance with 200 GB of storage capacity that runs on a db.m5.xlarge DB instance type.

Component	Number of Components Created
Amazon EC2 instances	Informatica domain with an EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Informatica domain with an EC2 instance and 200 GB of storage space (or 150 GB of storage space for a Windows node). The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, one to six Informatica node with an EC2 instance and 200 GB(150 GB for windows node) of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice one of the above mentioned node will be used as high availability node. Based on your choice, 1 (one) Informatica high availability domain node on an EC2 instance with 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, 1 (one) Informatica high availability domain node on an EC2 instance with 200 GB of storage space (or 150 GB of storage space on a Windows node). The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, a bastion server with an EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large. Based on your choice, a bastion server with an EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large or m4.large, depending on availability of the instance size in the selected region. Based on the Deployment Type you select, the deployment launches the following: - 1, 3, or 6 EC2 instances with 200 GB of storage space each for the embedded Hadoop cluster. - 2 EC2 instances for the Informatica compute cluster on EMR.
Elastic IP addresses	1 elastic IP address attached to the NAT gateway.
EFS File System	1

Components in a Deployment on an Existing VPC

The deployment creates the following components when you deploy the software in an existing virtual private cloud:

Component	Number of Components Created
Security groups	The following security groups: - One security group for RDS. - One security group for the Informatica Enterprise Data Preparation instance. - One security group for the Informatica domain instance. - If you selected the option to create the bastion server, one security group to support remote login to the bastion server. - Based on your choice, two security groups for Informatica EMR workflow connections.
Database subnet group	1
Amazon RDS instance	1 Amazon RDS for Oracle instance with 200 GB of storage capacity that runs on a db.m5.xlarge DB instance type.

Component	Number of Components Created
Amazon EC2 instances	Informatica domain with an EC2 instance and 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Informatica domain with an EC2 instance and 200 GB of storage space (or 150 GB of storage space for a Windows node). The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, one to six Informatica node with an EC2 instance and 200 GB(150 GB for windows node) of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice one of the above mentioned node will be used as high availability node. Based on your choice, 1 (one) Informatica high availability domain node on an EC2 instance with 200 GB of storage space. The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, 1 (one) Informatica high availability domain node on an EC2 instance with 200 GB of storage space (or 150 GB of storage space on a Windows node). The instance created is of type m5.2xlarge or m4.2xlarge, depending on availability of the instance size in the selected region. Based on your choice, a bastion server with an EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large. Based on your choice, a bastion server with an EC2 instance that has 150 GB of storage space is created. The instance created is of type m5.large or m4.large, depending on availability of the instance size in the selected region. Based on the Deployment Type you select, the deployment launches the following: - 1, 3, or 6 EC2 instances with 200 GB of storage space each for the embedded Hadoop cluster. - 2 EC2 instances for the Informatica compute cluster on EMR.
AWS Lambda functions	5 functions to confirm the validity of the VPC, the subnets, and the Informatica license.
IAM role and instance profile	1
EFS File System	1

Bastion Server

You can optionally deploy a bastion server as an element in the solution. The bastion server acts as a firewall between the internet and the cloud platform network where the solution deploys. It can also act as a remote Windows server which you can log into to run Informatica clients.

Informatica Domain

The Informatica domain is a server component that hosts application services, such as the Model Repository Service and the Data Integration Service. These services, together with domain clients, enable you to create and run mappings and other objects to extract, transform, and write data.

Application Services

Enterprise Data Preparation Service

The Enterprise Data Preparation Service is an application service in the Informatica domain that runs the Enterprise Data Preparation application.

Interactive Data Preparation Service

The Interactive Data Preparation Service is an application service in the Informatica domain that manages data preparation within the Enterprise Data Preparation application.

Catalog Service

The Catalog Service is an application service in the Informatica domain that runs the Enterprise Data Catalog application, and manages connections between service components and external applications.

Model Repository Service

The Model Repository Service is an application service in the Informatica domain that manages the Model repository. The Model repository stores metadata created by Informatica products in a relational database to enable collaboration among the products. Informatica Developer, the Data Integration Service, and the Administrator tool store metadata in the Model repository.

Data Integration Service

The Data Integration Service is an application service in the Informatica domain that performs data integration tasks for the Developer tool and for external clients.

Metadata Access Service

The Metadata Access Service is an application service that allows the Developer tool to access Hadoop connection information to import and preview metadata. The Metadata Access Service contains information about the Service Principal Name (SPN) and keytab information if the Hadoop cluster uses Kerberos authentication.

Content Management Service

The Content Management Service is an application service in the Informatica domain that manages reference data and is responsible for compiling rule specifications into mapplets. The Content Management Service provides reference data information to the Data Integration Service and to the Developer tool and Analyst tool. The Content Management Service stores reference data in a database that you specify.

Analyst Service

The Analyst Service is an application service in the Informatica domain that runs Informatica Analyst. The Analyst Service manages the connection between the service components and the users who log in to Analyst tool. You can perform column and rule profiling, manage scorecards, and manage bad records and duplicate records in the Analyst tool. The Analyst Service stores profiling, scorecarding, and bad and duplicate record data in databases that you specify.

Search Service

The Search Service is a component within the Informatica domain that manages search operations in the Analyst tool and Business Glossary Desktop. By default, the Search Service returns search results from a Model repository, such as data objects, mapping specifications, profiles, reference tables, rules, and scorecards.

Informatica Cluster Service

The Informatica Cluster Service is an application service that runs and manages all the Hadoop services, Apache Ambari server, and Apache Ambari agents on an internal Hadoop cluster.

The Informatica domain can run several other services. For more information about Informatica services, see the Informatica Application Service Guide.

Repositories

Informatica repositories, posted on Oracle SQL databases, store metadata about domain objects. Informatica repositories include the following:

Domain configuration repository

The domain configuration repository stores configuration metadata about the Informatica domain. It also stores user privileges and permissions.

Model repository

The Model repository stores metadata for projects and folders and their contents, including all repository objects such as mappings and workflows. The repository also stores rules you apply during data preparation.

Data Preparation repository

The Data Preparation repository stores worksheet metadata created when you use the Enterprise Data Preparation application to prepare data.

In addition to these domain repositories, the solution also requires a repository for Hive metadata. This repository is hosted on an SQL database. It stores Hive table metadata to enable Hadoop operations.

For more information about domain repositories, see the Informatica Application Service Guide.

Clusters

Informatica uses the following EC2 node clusters.

Embedded Hadoop cluster

An embedded Hadoop cluster deployed on Amazon EC2 instances that the Catalog Service uses to run metadata processing and profiling jobs.

Informatica compute cluster on EMR

A compute cluster of Amazon EC2 instances with autoscaling enabled that Enterprise Data Preparation uses to publish prepared data to the data lake.

Informatica Clients

You can use different clients with Data Engineering Streaming:

Administrator tool

The Administrator tool enables you to create and administer services, connections, and other domain objects.

Analyst tool

The Analyst tool enables you to analyze, cleanse, standardize, profile, and score data in your enterprise.

Developer tool

The Developer tool enables you to create and run mappings and other objects that enable you to access, transform, and write data to targets.

Command line interface

The command line interface offers hundreds of commands to assist in administering the Informatica domain, creating and running repository objects, administering security features, and maintaining domain repositories.

Bastion server

You can optionally deploy a bastion server as an element in the solution. The bastion server is an AWS Windows instance installed with the Developer tool and the command line interface clients. The bastion server acts as a firewall for access to the VPC. It can also act as a remote Windows server which you can log into to run Informatica clients.

Administrator tool

The Administrator tool enables you to create and administer services.

PowerCenter client

The PowerCenter Client consists of applications that you can use to design mappings, mapplets, create sessions and workflows to load the data, and monitor workflow progress.

Bastion server

You can optionally deploy a bastion server as an element in the solution. The bastion server is an AWS Windows instance installed with the Developer tool and the command line interface clients. The bastion server acts as a firewall for access to the VPC. It can also act as a remote Windows server that you can log into to run Informatica clients.

Before You Begin

Before you launch the automated deployment on the AWS Marketplace, verify the prerequisites and make the choices described in this section.

Network Prerequisites

Before you deploy Data Engineering Streaming, verify that you have completed the following network prerequisites for the Informatica application services.

Network prerequisites depend on whether you choose to deploy the solution to a new subnet or an existing subnet.

When you deploy to a new VPC, the automated deployment builds a new AWS environment consisting of the VPC, subnets, NAT gateways, security groups, and other infrastructure components. The deployment then deploys Data Engineering Streaming to the new VPC.

When you deploy to an existing VPC, the automated deployment provisions Data Engineering Streaming in your existing AWS infrastructure.

New VPC Prerequisites

Complete the following prerequisites when you choose to deploy into a new VPC:

- Ensure the selected Availability Zone has sufficient capacity to create a new subnet, route table, internet gateway, and NAT gateway.
- Ensure that the VPC creation has not exceeded the supported limit.

Existing VPC Prerequisites

Complete the following prerequisites when you choose to deploy into an existing VPC:

- Ensure that the VPC has DNS resolution enabled.
- You must also attach the subnet to a route table with a local route to the VPC CIDR.
- Verify that each subnet in the VPC is associated with a route table.
- Verify that the subnet associated with an internet gateway has the Auto-assign Public IPv4 property enabled.
- Verify that each route listed in the route table is defined as one of the following gateway types:
  - Internet gateway
  - NAT gateway
  - Virtual private gateway
- Verify that the subnet is associated with one of the following gateways:
  - Internet gateway
  - NAT gateway

  - Virtual private gateway
- Verify that the service subnets you choose are attached to a route table under the selected VPC.
- Verify that your VPC has a minimum of two subnets. If you choose to create a bastion server, you can deploy it to one of the two subnets or use a separate public subnet.

Subnets require an Internet gateway, an NAT gateway, or a virtual private gateway in their route tables, to allow the instances to download packages and software without exposing them to the Internet. Record the domain name configured in the DHCP options as explained in the Amazon VPC documentation. You provide the domain name in the VPC settings when you configure the deployment.

Account and Security Prerequisites

Complete the prerequisites in this section to verify that you obtained the required AWS account and private keys.

AWS Account and Region

If you do not already have an AWS account, create an account at aws.amazon.com.

Use the region selector in the navigation bar to choose the AWS region where you want to deploy Enterprise Data Preparation on AWS.

Note: The region you choose determines the availability of certain sizes of EC2 instances.

If necessary, request a service limit increase for the EC2 instance type you've decided to use for the Informatica domain. You might need to do this if you already have an existing deployment that uses this instance type, and you think you might exceed the default limit with this reference deployment.

Private Keys

Create a key pair in your preferred region.

When you log in to an EC2 instance, you use a password file for authentication. The file is called a private key file and has a file name extension of .pem.

If you do not have an existing .pem key to use, follow the instructions in the AWS documentation to create a key pair.

Note: Your administrator might ask you to use an existing key pair.

When you create a key pair, you save the .pem file to your desktop system. Simultaneously, AWS saves the key pair to your account. Make a note of the key pair that you want to use for the deployment, so that you can provide the key pair name during network configuration.

Account User Policies

Choose to grant the AWS user one of the following privileges:

- Administrator privileges
- The following minimum policies attached to the user account:
  - arn:aws:iam::aws:policy/AWSCloudFormationFullAccess
  - arn:aws:iam::aws:policy/AmazonEC2FullAccess
  - arn:aws:iam::aws:policy/CloudWatchFullAccess
  - arn:aws:iam::aws:policy/IAMFullAccess
  - arn:aws:iam::aws:policy/AmazonS3FullAccess
  - arn:aws:iam::aws:policy/AmazonRDSFullAccess
  - arn:aws:iam::aws:policy/AmazonVPCFullAccess
  - arn:aws:iam::aws:policy/AWSLambdaFullAccess

- arn:aws:iam::aws:policy/AmazonElasticMapReduceFullAccess
- arn:aws:iam::aws:policy/AmazonElasticFileSystemFullAccess

Storage Prerequisite

Create an Amazon S3 bucket for the PowerCenter license.

Amazon S3 stores objects, consisting of one or more files and accompanying file metadata, in "buckets." For information about Amazon S3, see the AWS documentation.

To use the automated solution on the cloud marketplace, you must create an S3 bucket to store the license. You can store the license on the bucket level or in a folder within the S3 bucket.

Note the location of the license key and supply it when you configure the solution in the **License Key S3 Bucket URI** parameter.

Deploying Enterprise Data Preparation on the AWS Marketplace

The automated deployment of Enterprise Data Preparation on the AWS marketplace uses an AWS CloudFormation template to guide you through the process and launch the solution deployment.

Use the AWS CloudFormation console to specify the stack name and to enter parameter values in the template that you use to deploy the application. A stack is a logical representation of an AWS deployment.

Specify the identifiers for the following components:

- The VPC.
- The subnet that contains the Amazon EC2 instances where the deployment process deploys the application components.
- The subnet where the deployment process creates the Model repository.

The deployment process provisions the application components and creates the Model repository and the Data Preparation repository based on the values that you specify.

Availability Zones and Subnets

The deployment process deploys the components and services for Enterprise Data Preparation into an Amazon VPC configured across two Availability Zones. Each Availability Zone is provisioned with two subnets.

In the first subnet, the deployment process provisions the following application services on an Amazon EC2 instance within the VPC:

- Data Integration Service
- Model Repository Service
- Metadata Access Service
- Catalog Service
- Enterprise Data Preparation Service
- Interactive Data Preparation Service

For more information about domain components, see the Informatica Administrator Guide and the Informatica Application Service Guide.

In the second subnet, the process deploys an Oracle database on Amazon RDS for use as the Model repository and the Data Preparation repository.

Step 1. Get a License for Enterprise Data Preparation

1. This deployment requires a license for Informatica Enterprise Data Preparation. To sign up for a demo license, please contact Informatica.
2. Place the license key file for the software in the S3 bucket or directory that you created. You will be prompted for the S3 bucket location during the deployment process.

Step 2. Launch the Deployment

You are responsible for the cost of the AWS services used while running this deployment. There is no additional cost for using this deployment. See the pricing pages for each AWS service you will be using for full details. Each deployment takes approximately one to two hours to complete.

Perform the following steps to launch the deployment:

1. Login to the AWS Marketplace using your AWS account.
2. Type `Informatica Enterprise Data Preparation` in the search box and press **Enter**.
3. Click **Informatica Enterprise Data Preparation 10.4.0**.

   The **Product Overview** page appears.
4. Click **Continue to Subscribe**.

   The **Subscribe to this software** page appears.
5. Click **Continue to Configuration**.

   The **Configure this software** page appears.
6. Choose whether to deploy the solution on a new VPC or an existing VPC.
7. Click **Continue to Launch**.

   The **Launch this software** page appears.
8. Select **Launch CloudFormation** from the **Choose Action** drop-down list and click **Launch**.

   The **Create stack** page appears.
9. Verify the region for the deployment. The console displays the region in the upper-right corner of the navigation bar. Change the region if necessary.
10. In the **Stack** pane, ensure that the default option `Amazon S3 URL` is selected in the **Specify Template** section.

    Caution: Do not change any settings on this page or navigate away from the page. If you leave the page, the Amazon S3 URL might be erased.

    Click **Next**. The **Specify stack details** page opens.

Step 3. Configure the Deployment

On the Specify stack details page, review the parameters for the template. Provide values for the parameters that require input. For all other parameters, review the default settings and customize them as necessary.

Parameters to Deploy to an Existing VPC

You can deploy Enterprise Data Preparation with AWS CloudFormation to a VPC that already exists within your Amazon infrastructure.

Specify the identifiers for the VPC, for the subnet containing the Amazon EC2 instances where the deployment process deploys the application components, and for the subnets used by the Informatica services. The deployment process provisions the application components and creates the required repository databases based an the values you specify.

1. Enter the name of the stack to create.
2. Enter the parameters required to deploy the application.

   The following table describes the parameters you enter to deploy to an existing VPC:

Parameter	Default	Description
VPC	Requires input	Identifier for the Amazon VPC on which to deploy the application. For example: `vpc-0343606e`
Subnet 1	Requires input	Identifier for the subnet within the VPC where the installer deploys the application.
Subnet 2	Requires input when High Availability is enabled	Identifier for the subnet where the installer deploys the application in high availability mode.
CIDR IP Range	Requires input	Classless Inter-Domain Routing IP address range associated with the Amazon VPC. The CIDR range must have permissions to access the public subnet where the application runs. Format the string as `<IP address>/<IP network prefix>`. For example, to specify a range from `10.20.30.40` to `10.20.30.47`, enter `10.0.0.0/16`.
Pre-validation Check	No	Set to **Yes** to validate the VPC and subnet configurations before you deploy the application.
Deploy Bastion Server	No	Select **Yes** to deploy a bastion server.
Bastion Server Subnet	Requires input if you choose to deploy a bastion server	Subnets for the bastion server.
Key Pair Name	Requires input	Name for key pair required to access the EC2 nodes where the Informatica domain runs.

Parameter	Default	Description

License Key Bucket URI	Requires input	The URI for the S3 bucket that contains the Enterprise Data Preparation license key. Use one of the following formats:
		`s3://<bucket name>/license.key`
		`s3://<bucket name>/<folder name>/license.key`

High Availability Deployment	Disable	Choose to enable or disable high availability for the deployment. For information about high availability for the Informatica domain, see the Informatica documentation.
Deployment Type	Requires selection	Select the Amazon EC2 instance type where you deploy the embedded Hadoop cluster and the Informatica compute cluster on EMR. Choose from the following options for the embedded Hadoop cluster: - Small: one m*.4xlarge node - Medium: three m*.2xlarge nodes - Large: six m*.2xlarge nodes Choose from the following options for the Informatica compute cluster on EMR: - Small: one m*.4xlarge node - Medium: one to three m*.2xlarge nodes - Large: one to six m*.2xlarge nodes **Note:** The instance types you can select might depend on your AWS region.
S3 Bucket Name for Enterprise Data Preparation and Amazon EMR	N/A	Specify the name of an existing S3 bucket. If you do not specify a value, the deployment process creates an S3 bucket.
Cluster ID for Amazon EMR	N/A	Specify the identifier for an existing Amazon EMR cluster to use for the Informatica compute cluster. If you do not specify a value, the deployment process creates a cluster with autoscaling enabled. The number of EC2 nodes in the cluster depends on the Deployment Type value you set.

3. Click **Next**.
4. On the Options page, expand the **Advanced** tab.
5. Set the **Rollback on Failure** property to **No**.
6. Click **Next**.
7. Click **Create**.
8. When the status for the stack reads CREATE_COMPLETE, click the stack name.
9. Click the **Outputs** tab to find the URLs for the Informatica Administrator, Catalog, and Enterprise Data Preparation applications.

Parameters to Deploy to a New VPC

You can deploy Enterprise Data Preparation with AWS CloudFormation to a new VPC.

Specify the identifiers for the VPC, for the subnet containing the Amazon EC2 instances where the deployment process deploys the application components, and for the subnets used by the Informatica services. The deployment process provisions the application components and creates the required repository databases based on the values you specify.

1. Enter the name of the stack to create.

2. Enter the parameters required to deploy the application.

   The following table describes the parameters you enter to deploy to a new VPC:

Parameter	Default	Description
Availability Zone	Requires selection	Amazon Availability Zone where you want to create the VPC.
VPC CIDR	Requires input	Classless Inter-Domain Routing IP address range associated with the Amazon VPC. The CIDR range must have permissions to access the public subnet where the application runs. Format the string as `<IP address>/<IP network prefix>`. For example, to specify a range from `10.20.30.40` to `10.20.30.47`, enter `10.0.0.0/16`.
Public Subnet CIDR	Requires input	Classless Inter-Domain Routing IP address range associated with the public subnet.
Private Subnet 1 CIDR	Requires input when High Availability is enabled	Classless Inter-Domain Routing IP address range associated with the first private subnet.
Private Subnet 2 CIDR	Requires input	Classless Inter-Domain Routing IP address range associated with the second private subnet.
CIDR IP Range	Requires input	Classless Inter-Domain Routing IP address range that has permissions to access the application.
Deploy Bastion Server	No	Select **Yes** to deploy a bastion server.
Key Pair Name	Requires input	Name for key pair required to access the EC2 nodes where the Informatica domain runs.

License Key Bucket URI	Requires input	The URI for the S3 bucket that contains the Enterprise Data Preparation license key. Use one of the following formats:
		`s3://<bucket name>/license.key`
		`s3://<bucket name>/<folder name>/license.key`

High Availability Deployment	Disable	Choose to enable or disable high availability for the deployment. For information about high availability for the Informatica domain, see the Informatica documentation.

Parameter	Default	Description
Deployment Type	Requires selection	Select the Amazon EC2 instance type where you deploy the embedded Hadoop cluster and the Informatica compute cluster on EMR. Choose from the following options for the embedded Hadoop cluster: - Small: one m*.4xlarge node - Medium: three m*.2xlarge nodes - Large: six m*.2xlarge nodes Choose from the following options for the Informatica compute cluster on EMR: - Small: one m*.4xlarge node - Medium: one to three m*.2xlarge nodes - Large: one to six m*.2xlarge nodes **Note:** The instance types you can select might depend on your AWS region.
S3 Bucket Name for Enterprise Data Preparation and Amazon EMR	N/A	Specify the name of an existing S3 bucket. If you do not specify a value, the deployment process creates an S3 bucket.
Cluster ID for Amazon EMR	N/A	Specify the identifier for an existing Amazon EMR cluster to use for the Informatica compute cluster. If you do not specify a value, the deployment process creates a cluster with autoscaling enabled. The number of EC2 nodes in the cluster depends on the Deployment Type value you set.

3. Click **Next**.
4. On the Options page, expand the **Advanced** tab.
5. Set the **Rollback on Failure** property to **No**.
6. Click **Next**.
7. Click **Create**.
8. When the status for the stack reads CREATE_COMPLETE, click the stack name.
9. Click the **Outputs** tab to find the URLs for the Informatica Administrator, Catalog, and Enterprise Data Preparation applications.

Step 4. Recycle the Informatica Services

After the deployment completes, use the Administrator tool to recycle the Informatica application services.

1. Log in to the Administrator tool.
2. Click the **Services and Nodes** tab.
3. In the Domain Navigator, select the Model Repository Service.
4. From the **Actions** menu, select **Search Index > Re-index**.
5. From the **Search Index > Re-index** menu, select **Recycle Service**.
6. Select the Data Integration Service, and then select **Actions > Recycle Service**.
7. Select the Content Management Service, and then select **Actions > Recycle Service**.
8. Select the Interactive Data Preparation Service, and then select **Actions > Recycle Service**.
9. Select the Enterprise Data Preparation Service, and then select **Actions > Recycle Service**.
10. Log in to the Enterprise Data Preparation application.

11. Log out of the Enterprise Data Preparation application, and then log in again.

Monitoring Instance Provision and Informatica Domain Creation

You can use cloud platform dashboards, logs, or other artifacts to see whether cluster creation succeeded and how to locate and identify the Informatica domain and Data Engineering Integration on the cloud platform.

Deployed AWS Resources

You can monitor the progress of the solution deployment in the Create Stack section of the AWS portal. Each row of the display shows stack components and creation status. The stack is complete when the status reads CREATE_COMPLETE.

When you deploy the solution to an existing VPC, the console displays the following stack components:

Component	Description
StackName/PCStack	The main stack. Created by default.
StackName-LicenseCheckLambdaStack	Checks whether the license key file is valid. Created by default.
StackName-BastionSubnetCheckLambdaStack	Checks whether the bastion server subnet is valid. Created by default. Lambda function will be created only when: - Pre-validation check parameter set to Yes - Deploy Bastion Server set to Yes
StackName-SubnetsAZCheckLambdaStack	Checks whether the subnets are in different Availability Zones. Created by default.
StackName-PublicIpDnsCheckLambdaStack	Checks whether an auto-assign public IPv4 address is enabled for a subnet. Created by default.
StackName-Subnet2ChecklambdaStack	Checks whether the second subnet is valid. Created by default. Creates the Lambda function if the Pre-validation Check parameter is set to Yes.
StackName-Subnet1CheckLambdaStack	Checks whether the first subnet is valid. Created by default. Creates the Lambda function if the Pre-validation Check parameter is set to Yes.
StackName-IAMStack	Creates IAM roles and profiles used by EC2 instances. Created by default.
StackName-EFSStack	Creates Elastic File system and its endpoint. Created only when the **Create EFS** parameter set to `Yes` and the region supports it. **Note:** Applicable for Linux deployment.
StackName-SecurityGroupStack	Creates security groups and network interfaces used by EC2 instances. Created by default.

Component	Description
StackName-RDSStack	Creates RDS resources. Created by default.
StackName-BastionServerStack	Creates bastion server on an EC2 instance. Created when you choose to create the bastion server.
StackName-PCServerStack	Creates the Informatica PowerCenter gateway node EC2 instance. Created by default.
StackName-PCNode2Stack	Creates the Informatica PowerCenter node2 or HA node EC2 instance. Created only when there are two or more PowerCenter nodes or when **High Availability** is enabled.
StackName-PCNode3Stack	Creates the Informatica PowerCenter node3 EC2 instance. Created only when there are three or more PowerCenter nodes.
StackName-PCNode4Stack	Creates the Informatica PowerCenter node4 EC2 instance. Created only when there are four or more PowerCenter nodes.
StackName-PCNode5Stack	Creates the Informatica PowerCenter node5 EC2 instance. Created only when there are five or more PowerCenter nodes.
StackName-PCNode6Stack	Creates the Informatica PowerCenter node6 EC2 instance. Created only when there are six or more PowerCenter nodes.
StackName-S3BucketStack	Creates an Amazon S3 bucket. Created when you leave the S3 Bucket Name value for Enterprise Data Preparation and Amazon EMR empty.
StackName-EMRStack	Creates an Amazon EMR cluster. Created when you leave the Cluster ID for Amazon EMR value empty.
StackName-IHSStack	Creates an Informatica Embedded Hadoop Cluster. Created by default.
StackName-IHSStack-HadoopGatewayStack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Small. Medium, or Large.
StackName-IHSStack-MultiNodeHadoopNode2Stack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Medium or Large.
StackName-IHSStack-MultiNodeHadoopNode3Stack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Medium or Large. .
StackName-IHSStack-MultiNodeHadoopNode4Stack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Large.
StackName-IHSStack-MultiNodeHadoopNode5Stack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Large.

Component	Description
StackName-IHSStack-MultiNodeHadoopNode6Stack	Creates an Informatica Embedded Hadoop Cluster. Created when you set Deployment Type to Large.
StackName-PCServerStack	Creates an Informatica gateway node on which PowerCenter runs on an EC2 instance. Created by default.
StackName-PCHANodeStack	Creates a high availability node on which PowerCenter runs on an EC2 instance. Created when you set High Availability Deployment to Enabled.

When you deploy to a new VPC, the console displays the following stack components:

Component	Condition for Creation
StackName-VPCStack	Creates network-related resources including the VPC, subnets, internet gateway, and the NAT gateway. Created by default.
PCStack	The main stack. Created by default.

Output Tab Properties

Click the **Output** tab to view the list of resources that the deployment process created in the VPC. The following table lists the outputs:

Output	Description
AdminConsoleUser	User name of the Administrator tool user.
AdminConsoleUserPwd	Password for the Administrator tool user.
BastionPrivateDNS	DNS of the private bastion server. For example: ip-<IPv4 address>-<region designation>.compute.internal
BastionPublicIPDNS	DNS of the public bastion server. For example: ip-<IPv4 address>-<region designation>.compute.amazonaws.com
DBUserPwd	Password for the database users of the domain databases created in the deployment.
DBUsers	User IDs for the database users of the domain databases created in the deployment.
HANodeIP	IPv4 address of the high availability or failover node for the domain.
InformaticaAdminConsolePrivateURL	URL:port of the Administrator tool hosted on a private server.

Output	Description
InformaticaAdminConsolePublicURL	Public URL:port of the Administrator tool.
InformaticaAdminConsoleServerLog	Path to the domain installation log.
InformaticaEDPPrivateURL	URL:port of the Enterprise Data Preparation application user interface hosted on a private server.
InformaticaEDPPublicURL	Public URL:port of the of the Enterprise Data Preparation application user interface.
InstanceSetupLogs	Path to the domain EC2 host setup log.
InformaticaAnalystPublicURL	Informatica Analyst tool URL (Public DNS)
InformaticaAnalystPrivateURL	Informatica Analyst tool URL (Private DNS)

The table below lists the outputs for domain expansion:

Key	Description
NodePrivateIP	Informatica Joined Node Private IP.

Use the URL, user ID, and password to access the Administrator tool to perform administrator tasks.

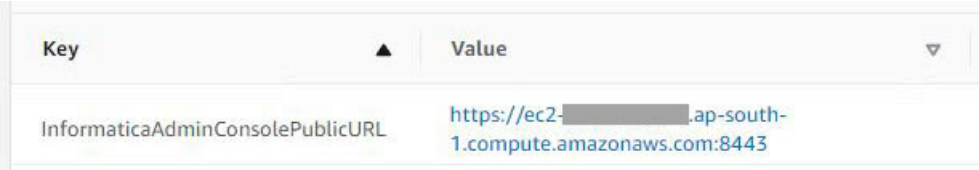
For a description of domain and PowerCenter administrator tasks, see the Informatica Administrator Guide.

For a description of domain and PowerCenter administrator tasks, see the Informatica Administrator Guide and the Informatica Enterprise Data Preparation Administrator GuideData Engineering Integration Administrator GuideApplication Services Guide.

Deployed Informatica Resources

The list of deployed resources includes the **InformaticaAdminConsolePublicURL**. Click the linked value of this property.

The following image shows an example of this property, with the IP address portion redacted:



The Administrator tool opens in a browser window.

Use the Informatica Administrator user name and password listed in the **Outputs** list to log in to the Administrator tool.

You can use the Administrator tool to view and administer the Informatica application services and resources that the automated deployment created.

Solution Deployment Log Files

You can review the following log files to see the solution deployment as it happened.

Review the following log files in the `/var/log` directory:

- cfn-init.log
- cfn-init-cmd.log
- cfn-wire.log
- cloud-init.log
- messages

Review the following log file in the `/opt` directory:

- Infa_OneClick_Solution.log

Review the following logs in the `/var/log` directory in Linux deployment and `C:\cfn\log\` directory in Windows deployment:

- *cfn-init.log*
- *cfn-init-cmd.log*
- Available only in Linux deployment.*cloud-init.log*

Review the following log in the `/opt` directory in Linux deployment and `C:\` directory in Windows deployment :

- Infa_OneClick_Solution.log

Informatica Domain Logs

You can access the following logs on the Informatica domain VM:

Command execution log

This log records the following events:

- Creation of Informatica connections, cluster configurations, and services.
- Population of the data warehouse and SQL databases.
- PowerCenter Integration Service recycling to register all changes.

At the top of the log file is a summary section that lists automated tasks and their status. Beneath the summary section are detailed sections about each task. If any of the tasks failed to complete successfully, you can look at the detailed section for the task to troubleshoot the task.

File name: *Infa_OneClick_Solution.log*

Linux Location: `/opt`

Windows Location: `C:\`

Linux Location: `/opt`

File name: *cfn-init-cmd.log*

Windows Location: `C:\cfn\log\`

Linux Location: `/var/log`

File name: *cfn-init.log*

Windows Location: `C:\cfn\log\`

File name: *cloud-init.log*

Linux Location: `/var/log`

Command execution log

This log records the following events:

- Creation of Informatica connections, cluster configurations, and services.
- Population of the data warehouse and SQL databases.
- Data Integration Service recycling to register all changes.

At the top of the log file is a summary section that lists automated tasks and their status. Beneath the summary section are detailed sections about each task. If any of the tasks failed to complete successfully, you can look at the detailed section for the task to troubleshoot the task.

File name: `Infa_OneClick_Solution.log`

Location: `/opt`

CloudWatch logs for Informatica

CloudWatch is an AWS service that enables you to monitor all of the events in your AWS deployments. CloudWatch provides a central location from which to monitor EC2 instances and other resources. for more information, see the AWS documentation.

To monitor PowerCenter events in CloudWatch, go to the CloudWatch page and search for "Informatica."

If no log files are returned by the search, log into the VM where the Informatica domain is installed and check for logs with the filename "`cfn*.log`"

Access the Informatica deployment logs through the AWS console:

1. In the Amazon Web Services Console, select CloudWatch from the Services menu.
2. In the navigation pane, choose **Logs** > **Log groups**.
3. Enter InformaticaPCLogGroup in the search pane.

CloudWatch logs for Lambda functions

If the automated deployment fails during execution, CloudWatch logs display information about process errors.

Access the CloudWatch logs through the AWS console:

1. In the Amazon Web Services Console, select CloudWatch from the Services menu.
2. In the navigation pane, choose **Logs** > **Log groups**.
3. Enter `/aws/lambda/<stack name>` in the search pane.
4. Select the log group associated with the name of the stack that failed.

Troubleshooting

Review the AWS log files to troubleshoot issues with your Data Privacy Management deployment.

Stack Failure

Review the following log files if an EC2 instance where Informatica services run is running but the stack fails, or if the Administrator tool fails to launch:

- Installation log file: `/home/Infa_OnceClick_Solution.log`
- Command Initialization log file: `/var/log/cfn-init-cmd.log`
- Service Initialization log file: `/var/log/cfn-init.log`
- Cloud log file: `/var/log/cloud-init.log`

Stack Launch Failure When the Pre-validation Check Parameter is Enabled

Review the Lambda CloudWatch log files if a stack failure occurs when the Pre-validation Check option is enabled.

1. In the AWS Management Console, select **CloudWatch** from the **Services** menu.
2. Select **Logs**, then and then select **Log Groups**.
3. Enter the following string in the search bar:

   `/aws/lambda/<stack name>`
4. Click the log group associated with the stack.

Stack Launch Failure When Rollback On Failure Is Enabled

Review the Lambda CloudWatch log files if a stack failure occurs when the **Rollback on Failure** option is enabled.

1. In the AWS Management Console, select **CloudWatch** from the **Services** menu.
2. Select **Logs**, then and then select **Log Groups**.
3. Enter the following string in the search bar: `InformaticaDPMLogGroup`
4. Click a log file in the log group associated with the stack.

   You might want to review the following log files:

   - installation.log
   - cfn-init
   - InformaticaClusterService
   - EDCService
   - DPMService
   - EDPService

Additional Resources

AWS services

- AWS CloudFormation
- Amazon EBS
- Amazon EC2
- Amazon VPC

Enterprise Data Preparation documentation

- Informatica Documentation

Reference deployments

- AWS Deployment home page

Author

Dan Hynes